

mcp-emacs — Source Reference

A guided tour of the code, its structure, and the Emacs Lisp it leans on

2026-09-07

Gunnar Bastkowski

v1.11.0+4

Contents

About This Document	7
1 What This Project Is	8
1.1 The one-sentence version	8
1.2 What that actually means	8
1.3 The longer-term intent	8
1.4 Scale and shape	9
2 Repository Layout	10
2.1 Top-level map	10
2.2 The <code>elisp/</code> directory in dependency order	10
2.3 What is deliberately <i>not</i> here	12
3 Architecture and the Decisions Behind It	14
3.1 The request path	14
3.2 Decision: the server lives in Emacs	14
3.3 Decision: idempotent lifecycle	15
3.4 Decision: async tools for human-answered requests	15
3.5 Decision: two JSON gotchas, encoded as helpers	16
3.6 Decision: registries, not cond chains	17
3.7 Decision: return status strings, not errors, for empty results	17
4 <code>mcp-emacs.el</code> — The Helper Layer	19
4.1 The foundational abstraction: “the current buffer”	19
4.2 Reading state	20
4.3 Position arithmetic	20
4.4 Editing	21
4.5 Diagnostics: two backends, one interface	21
4.6 Navigation via <code>imenu</code>	22
4.7 Org helpers	23
4.8 Environment and self-diagnosis	24
4.9 Project tools with a projectile fallback	24
5 <code>mcp-emacs-server.el</code> — The MCP Server	26
5.1 Schema construction	26
5.2 The tool registry	26
5.3 Resources	27
5.4 JSON-RPC dispatch	27
5.5 The HTTP handler and the deferral	28
5.6 Releasing a kept-alive connection	30
5.7 The permission gate route	30
5.8 Body parsing and the JSON decode configuration	31
6 The Interactive Diff Review	33
6.1 The problem	33

6.2	The shared review function	33
6.3	Saving and restoring the window layout	34
6.4	The startup hook: bindings and the quit hook	35
6.5	Single-path resolution	35
6.6	Failure during setup	36
6.7	The synchronous caller	36
6.8	The asynchronous caller	37
7	The Org Task-Session Protocol	39
7.1	The data format	39
7.2	The change token	39
7.3	Identifying items without owning them	40
7.4	Appending without disturbing	41
7.5	Validating keywords against the user's own workflow	41
7.6	Waiting for the human	42
7.7	Why this shape	42
8	The Shared Agent Layer	43
8.1	Why a base class at all	43
8.2	The optional verbs, and what their defaults say	44
8.3	The event vocabulary	44
8.4	Finding a conversation without a registry	45
8.5	Sharing what the human is looking at	45
8.6	agent-prompt.el — a buffer, not a minibuffer	46
8.6.1	Windows, and the timer	46
8.7	agent-permission.el — a tool call as a question	47
8.7.1	The measured glob behaviour	48
8.7.2	Writing to a file someone else owns	48
8.7.3	The deadline	49
9	mcp-emacs-ide.el — The Claude Code IDE Surface	50
9.1	Why this exists at all	50
9.2	What the spike found	50
9.3	Session state as a struct	51
9.4	The discovery lockfile	51
9.5	Binding a port	52
9.6	Message dispatch and the deferred sentinel	52
9.7	The deferred openDiff	53
10	mcp-emacs-run.el — The Terminal Runner	55
10.1	What it does	55
10.2	Sessions without a registry	55
10.3	Resolving which session gets input	56
10.4	Window placement: an ordinary window, deliberately	56
10.5	Beating the framework at display-buffer	57
10.6	Passing the IDE port through the environment	58
10.7	The eat-term-send-string trap	58
10.8	Quitting with a deadline	59
10.9	Selection references	59
10.10	Headless one-shot queries	60

11	claude-client.el — The Terminal-Free Runner	61
11.1	The shape	61
11.2	Two differences from the opencode backend, stated up front	61
11.3	Constructing the command line	62
11.3.1	The settings file layers, it does not replace	63
11.3.2	Installing the PreToolUse gate	63
11.4	The event log and the subscriber hook	64
11.5	Framing NDJSON	65
11.6	Turning messages into events	66
11.7	Full re-render on every event	67
11.8	Turns over one process	67
11.9	Reopening a past session	68
11.10	Notes: the human as a second producer	68
11.11	Interruption: the design question, answered by measurement	68
11.12	Backpressure: three defects, found by probing	69
11.13	Guarding against a second run	69
11.14	The agent-backend methods	70
11.15	The major mode	70
11.15.1	Two keymaps, and why	70
12	opencode-client.el — The Second Backend	72
12.1	The shape	72
12.2	One server per project, on parallel ports	72
12.3	What its agent-backend methods reveal	73
12.4	The gap the overview has to work around	73
13	agent-session-overview.el — One Live View	74
13.1	Enumeration without a registry, again	74
13.2	Reporting only what a backend can actually observe	74
13.3	The subscriber swallows its own errors	75
13.4	Two window-management traps	75
13.5	One list drives three surfaces	76
14	mcp-emacs-remote.el and mcp-emacs-run-resume.el	77
14.1	Remote prompt input	77
14.2	The Org transcript via advice	77
14.3	The human’s side of the same record	78
14.4	Quiet tools	79
14.5	The resume picker: reading someone else’s data format	79
15	orgspec — Concept and Data Model	81
15.1	The workflow	81
15.2	The marker table	81
15.3	The model: three structs	82
15.4	Deliberate omissions	83
16	orgspec — Parsing with org-element	84
16.1	The org-element AST	84
16.2	The section gotcha	84
16.3	Extracting body, drawers, scenarios	85
16.4	The two entry points	86

17 orgspec — The Fold	87
17.1 The four design commitments	87
17.2 Why the order	87
17.3 Subtree surgery	87
17.4 The four operations	88
17.5 The MODIFIED drop guard	89
17.6 The pure core	89
17.7 Grouping and the transaction boundary	90
18 Building a Domain Model on Org	92
18.1 The question	92
18.2 Why Org is not a DOM	92
18.3 Measuring the round-trip	93
18.4 The pattern: four rules	94
18.4.1 Rule 1 — Parse into your own structs, not the AST	94
18.4.2 Rule 2 — Carry the raw source	94
18.4.3 Rule 3 — Mutate with Org’s buffer API, in a temp buffer	95
18.4.4 Rule 4 — Keep the transformation pure	95
18.5 Where the model actually lives	96
18.6 What you give up, and when to reconsider	96
18.7 Applying this elsewhere	97
19 orgspec — Validator, Lifecycle, Review, Tools	98
19.1 The hard-gate validator	98
19.2 The lifecycle and the agenda payoff	99
19.3 The review	99
19.4 Typed MCP tools	100
20 The Test Suite	101
20.1 Still not ERT, but no longer bare	101
20.1.1 test/test-helper.el: one check, and a description	101
20.1.2 bin/test-emacs.sh: an Emacs of its own	101
20.1.3 The report	102
20.2 What batch Emacs does and does not have	102
20.3 Technique 1: cl-letf function stubbing	103
20.4 Technique 2: advice-add :override for capture	104
20.5 Technique 3: stub the missing package	104
20.6 Testing the pure core is easy	105
20.7 Testing time and timers	106
20.8 A reusable test macro	106
20.9 Testing the dispatch layer without a socket	107
20.10 What is and is not covered	108
20.11 Running the tests yourself	109
21 Modifying the Code	110
21.1 Setting up a development loop	110
21.2 Adding a tool	110
21.3 Adding a tool that needs a human answer	111
21.4 Adding an orgspec op or rule	112
21.5 Adding a subscriber to the event log	112
21.6 Adding a backend	113
21.7 Things that will bite you	113

21.8	Style conventions to match	114
21.9	Before you push	115
22	Packaging, CI, and the Plugin Surface	116
22.1	Distribution as a Claude Code plugin	116
22.2	MCP client configuration	116
22.3	Skills	116
22.4	CI	117
22.5	Cutting a release	118
23	Cross-Cutting Idioms	119
23.1	Degrade, do not fail	119
23.2	Guard cells for single delivery	119
23.3	Build in memory, write once	119
23.4	Derive state, do not store it	120
23.5	Publish events, do not advise	120
23.6	Comment the failure, not the code	120
24	The Claude Code CLI Surface	121
24.1	Three ways the project invokes the CLI	122
24.2	The flags, and what each one buys	122
24.3	The generated configuration files	123
24.4	The stream-json wire format	123
24.5	The negative finding: <code>can_use_tool</code>	124
24.5.1	What does work: the <code>PreToolUse</code> hook	124
24.6	Three gates, and the hole between them	126
24.7	The IDE protocol handshake	126
24.8	Version fragility, and how the project handles it	129
24.9	If you wanted to talk to the API directly	129
25	The Event Model	130
25.1	The problem with turns	130
25.2	The aggregates already exist	130
25.3	What exists today, and precisely what is missing	131
25.4	The target shape	132
25.5	The cooperative loop, as an event stream	133
25.6	The three hard questions, and their answers	133
25.7	The smallest first step, and what it decided	135
25.8	What this means if you are modifying the code today	135

About This Document

This is a reference for the source code of `mcp-emacs`: what each file does, how the pieces fit together, and — where the code reaches for something beyond everyday Emacs Lisp — an explanation of the feature it is using and why.

It assumes you can read basic Emacs Lisp: you know what `defun`, `let`, `setq`, and a buffer are; you have written a small `.emacs` customisation; you are comfortable with lists and `alist`s. It does *not* assume you know `cl-lib` structs, `org-element`, process filters, advice, dynamic binding subtleties, or the MCP protocol. Those are explained as they come up, in place, in shaded “Elisp Feature” notes.

The document is organised as follows.

- Chapters 1–3 set the stage: what the project is, how the repository is laid out, and the architectural decisions that everything else follows from.
- Chapters 4–7 cover the core: the helper layer, the MCP server, the interactive diff review, and the Org task-session protocol.
- Chapter 8 covers the shared agent layer — the backend base class, the prompt composition buffer, and the permission gate — which both chat clients are built on and where a third backend would plug in.
- Chapters 9–14 cover the agent-facing surfaces: the Claude Code IDE integration, the terminal runner, the terminal-free runner, the opencode client, the cross-backend session overview, and the remote/resume modules.
- Chapters 15–19 cover `orgspec`, the org-native spec workflow, which is effectively a second project sharing the same host.
- Chapters 20–23 cover the test suite and its runner, how to modify the code safely, the packaging and release process, and a set of cross-cutting idioms worth learning from here.
- Chapters 24–25 cover the Claude Code CLI surface the project depends on, and the event model the human/ AI workflow is built on.

Diagrams are rendered from the PlantUML sources in `docs/`; each is named in its caption so you can regenerate it with `plantuml -tpng docs/<name>.puml`.

Line-number references are given as `file:line` and were accurate at the time of writing; treat them as signposts rather than as guarantees.

1 What This Project Is

1.1 The one-sentence version

mcp-emacs runs a Model Context Protocol server *inside a live Emacs session*, so that an AI coding agent can see and manipulate the same buffers, windows, Org state, and diagnostics that the human is looking at.

1.2 What that actually means

Most editor integrations for AI agents work by having the agent read and write files on disk, and then telling the editor to reload them. That is a reasonable design, and it is also a lossy one. A file on disk does not know:

- which buffer the human is currently looking at,
- whether that buffer has unsaved changes,
- where point (the cursor) is,
- what the active region is,
- what Flycheck thinks is wrong with the code right now,
- which Org task the human has clocked in,
- what xref would say about a symbol given the currently-running LSP server.

All of that lives in the Emacs process. An external tool can only approximate it. So mcp-emacs inverts the arrangement: rather than the agent driving Emacs from outside, the *server the agent talks to is itself an Emacs Lisp program running in the Emacs process*. A tool call is an HTTP request that lands in an Emacs process filter, is dispatched to an ordinary Elisp function, and that function calls `buffer-substring-no-properties` or `org-get-heading` or `flycheck-current-errors` directly on live state.

There is no `emacsclient` round-trip on the request path, no subprocess, and no serialisation of editor state into a file. The trade-off is that the server inherits Emacs's single-threaded, cooperative execution model, which turns out to matter a great deal — see Chapter 3 and Chapter 6.

1.3 The longer-term intent

The README and docs/VISION.md frame this as a step away from the chatbot request/response loop and toward an environment in which the human and the AI work the same live artifacts concurrently. Several design choices only make sense in that light:

- The Org task-session tools (Chapter 7) let the human and the agent edit the same checklist in the same buffer, with the human as the driver.
- The interactive `ediff` review (Chapter 6) makes every agent-proposed file change a thing the human physically accepts or rejects in the editor.

- The permission gate (Chapter 8) does the same for tool calls: a shell command the agent is not allowed to run becomes a question in a buffer rather than a refusal it reports afterwards.
- The event-log design in `claude-client.el` (Chapter 11) publishes an append-only stream of what the agent is doing, so multiple consumers can subscribe — the beginning of “one log, many readers”. The session overview (Chapter 13) is the second reader, and exists to prove the first one was not a special case.

You do not need to buy the vision to read the code, but it explains why the codebase keeps choosing “human stays in the loop, in the editor” over “fastest path to the edit”.

1.4 Scale and shape

At the time of writing the project is roughly 9,600 lines of Emacs Lisp across 23 files in `elisp/`, plus about 6,600 lines of tests across 24 files in `test/`, and some 630 lines of shell in `bin/`. Every file is `lexical-binding: t`, every public function has a docstring, and CI byte-compiles everything and runs the full suite on Emacs 29.4 — 22 suites and around 800 assertions.

The test suites are not `ert` suites, despite what the file names suggest; they are batch scripts with a shared check vocabulary. Chapter 20 explains why, and what runs them.

It is not a large codebase. It *is* a fairly dense one: several files spend more lines on commentary explaining a protocol quirk or an Emacs gotcha than on the code implementing the workaround. That is a deliberate style, and this document leans on those comments heavily.

2 Repository Layout

2.1 Top-level map

mcp-emacs/	
├─ elisp/	the entire implementation (23 .el files)
├─ test/	the suites (23 files) + test/init/ (isolated init)
├─ bin/	test runner, release script, permission-gate hook
├─ skills/	Claude Code "skill" definitions (Markdown)
├─ commands/	a slash-command definition
├─ .claude/commands/	project-local slash commands (orgspec, opsx)
├─ .claude-plugin/	Claude Code plugin + marketplace manifests
├─ docs/	vision, diagrams, comparison write-ups
├─ orgspec/	this project's own orgspec specs and changes
├─ .github/workflows/ci.yml	byte-compile + test, all via bin/test-emacs.sh
├─ .mcp.json	MCP client config pointing at localhost:8765
├─ opencode.json	the same, for the opencode agent
├─ AGENTS.md	conventions and gotchas for contributors
├─ README.md	user-facing overview and tool table
└─ PLUGIN.md	plugin install notes

bin/ is newer than the rest and worth naming early, because two of the three scripts are load-bearing rather than convenience:

Script	Lines	Role
bin/test-emacs.sh	434	Runs the suites in an Emacs of their own. The single entry point, shared with CI.
bin/release.sh	139	Stamps the version in every place it is declared, tags, publishes.
bin/permission-gate.sh	60	The PreToolUse hook: asks Emacs to approve a tool call, fails closed.

2.2 The elisp/ directory in dependency order

The files fall into four groups. Within each group they are listed roughly in the order one file requires another.

Group 1 — the MCP core.

File	Lines	Role
<code>mcp-emacs.el</code>	1365	Every helper that touches editor state. No HTTP, no protocol.
<code>mcp-emacs-report.el</code>	144	Files GitHub issues about the tooling itself, via gh.
<code>mcp-emacs-server.el</code>	810	The HTTP server, tool/resource registries, JSON-RPC dispatch, the permission gate route.

Group 2 — the shared agent layer.

These three know nothing about any particular agent. They were factored out of the clients once there were two of them (issues #41, #56), and they are where a third backend would plug in.

File	Lines	Role
<code>agent-backend.el</code>	540	An EIEIO base class naming what every backend can do, the shared event vocabulary, and the shared keymap.
<code>agent-prompt.el</code>	297	A composition buffer for prompts — the git-commit shape, not the minibuffer.
<code>agent-permission.el</code>	451	A buffer that asks the human to allow or deny one tool call.

Group 3 — agent-facing surfaces.

File	Lines	Role
<code>mcp-emacs-ide.el</code>	467	A WebSocket server speaking Claude Code's unofficial IDE protocol.
<code>mcp-emacs-run.el</code>	680	Launches the Claude CLI in an eat terminal buffer; window management.
<code>mcp-emacs-run-resume.el</code>	172	A native picker over past Claude sessions on disk.
<code>mcp-emacs-remote.el</code>	371	Sends prompts to a running session; records activity as an Org transcript.
<code>claude-client.el</code>	1648	A terminal-free runner: subprocess + NDJSON stream + rendered buffer.

File	Lines	Role
opencode-client.el	915	An HTTP + SSE client for the opencode agent server.
agent-session-overview.el	441	One live tabulated-list view of every session, across all backends.

Group 4 — orgspec, the org-native spec workflow.

File	Lines	Role
orgspec.el	95	The marker table: op tags, property names, regexes, TODO roles.
orgspec-model.el	61	Three cl-defstructs: scenario, requirement, change.
orgspec-parse.el	129	org-element extraction of the model from buffers.
orgspec-fold.el	173	The delta fold — the load-bearing algorithm.
orgspec-validate.el	112	The hard-gate validator.
orgspec-commands.el	273	The verbs: new, status, archive.
orgspec-lifecycle.el	81	Moving a delta requirement through TODO states.
orgspec-agenda.el	60	One org-agenda custom command as an in-flight dashboard.
orgspec-review.el	93	Ediff the fold before it writes.
orgspec-mcp.el	230	Typed MCP tools over the orgspec verbs.

2.3 What is deliberately *not* here

There is no package.json, no build system, and no bundled dependency vendoring. The hard dependency list is one item long: web-server, for the HTTP listener. Everything else is either built into Emacs (json, cl-lib, org, org-element, ediff, project, xref, imenu, treesit) or a *soft* dependency loaded with (require 'foo nil t) and guarded at the point of use: websocket, eat, markdown-mode, plz, flycheck, projectile, lsp-mode.

Elisp Feature: soft requires

(require 'websocket) signals an error if the package is missing. (require 'websocket nil t) returns nil instead — the second argument is a filename override (nil = use the default search), and the third, NOERROR, suppresses the error.

The codebase uses this everywhere an optional package is involved, and then pairs it with one of three guards at the call site:

- `(featurep 'eat)` — was the feature actually provided?
- `(fboundp 'gfm-view-mode)` — does the specific function exist?
- `(bound-and-true-p flycheck-mode)` — is the variable both bound *and* true?

The last one matters because a plain `(when flycheck-mode ...)` signals `void-variable` if flycheck was never loaded. `bound-and-true-p` is a macro that expands to a `boundp` check first.

You will also see `declare-function` near the top of files:

```
(declare-function websocket-server "websocket" (port &rest plist))
```

This declares nothing at runtime. It exists purely to tell the byte-compiler “this function comes from that file, stop warning me that it is undefined”. Since CI byte-compiler with `-werror`-like scrutiny, every soft-dependency call site needs one.

3 Architecture and the Decisions Behind It

3.1 The request path

A single MCP tool call travels this route:

1. The agent (Claude Code, opencode, anything MCP-speaking) issues an HTTP POST to `http://localhost:8765/mcp` with a JSON-RPC body.
2. Emacs's network process for the listening socket receives bytes. The `web-server` package frames them into an HTTP request object and calls the registered handler — `mcp-emacs-server--handler` — **from inside a process filter**.
3. The handler parses the JSON body, looks up the named tool in a registry, and calls its `:handler` closure with the parsed arguments alist.
4. The handler closure calls an ordinary `mcp-emacs-*` function, which touches live buffer/window/Org state and returns a string.
5. The string is wrapped in MCP's `{"content": [{"type": "text", ...}]}` shape, encoded as JSON, and written back to the socket.

Steps 2–5 happen synchronously inside the filter, with one important exception covered in §3.4.




Figure 3.1: The HTTP transport: emacsclient starts the server once, then every request goes direct (docs/architecture.puml)

3.2 Decision: the server lives in Emacs

AGENTS.md states the intent plainly: “tool calls are dispatched directly to the `mcp-emacs-*`

helpers, so they observe the real buffers, windows, and Org state of the running session.”

The alternative — a Node.js server shelling out to `emacsclient --eval` per call — was the earlier design and was abandoned. The costs it imposed were a process spawn per tool call, a serialisation boundary that made rich values awkward, and the impossibility of anything long-lived (a review that waits for a human, a wait-for-change subscription).

`emacsclient` still has a role, but only to *start* the server:

```
emacsclient --eval '(mcp-emacs-server-ensure)'
```

which returns the endpoint URL and is safe to call repeatedly.

3.3 Decision: idempotent lifecycle

`mcp-emacs-server-ensure` (`mcp-emacs-server.el:666`) is the public entry point, and it is deliberately idempotent:

```
(defun mcp-emacs-server-ensure ()
  "Start the MCP server if it is not already running; return its URL."
  (interactive)
  (unless (mcp-emacs-server-running-p)
    (mcp-emacs-server-start))
  (format "http://localhost:%d/mcp" mcp-emacs-server-port))
```

This lets the same call appear in a Doom `config.el` startup hook, in a launcher script, and in a manual M-x invocation without anyone having to track whether the server is up.

Note the small detail in `mcp-emacs-server-start`:

```
(let ((proc (ws-process mcp-emacs-server--process)))
  (when (processp proc)
    (set-process-query-on-exit-flag proc nil)))
```

Without this, the listening socket counts as an active process, and Emacs asks “Active processes exist; kill them?” every time you quit. A one-line fix for a daily annoyance.

3.4 Decision: async tools for human-answered requests

This is the subtlest architectural decision in the project, and it is worth understanding before reading any of the diff-review code.

Emacs runs process filters with `inhibit-quit` bound to `t` and, critically, **without a command loop**. A process filter is not a normal interactive context: it is a callback that runs to completion while the rest of Emacs waits. If a filter blocks — say, in a loop waiting for the user to press a key — then the keys the user presses are never processed, because the code that would process them is the very thing being blocked.

That is fatal for a tool like `apply_diff`, whose entire purpose is to open an `ediff` session and wait for the human to press C-c C-c or q. Done synchronously in the filter, the diff appears on screen and its keys are dead.

The fix is a two-handler tool descriptor. A tool may declare both:

- `:handler` — the synchronous version, used by direct callers, and

- `:async-handler` — a function of `(args done)` that starts the work, returns immediately, and later calls `done` with the result string.

The HTTP handler checks for `:async-handler` first (`mcp-emacs-server.el:601`) and, when present, holds the TCP connection open, returns control to the event loop, and writes the response from the completion callback. The mechanism for holding the connection is a web-server-specific throw:

```
(throw 'close-connection :keep-alive)
```

Chapter 5 walks through the code; Chapter 6 walks through the review itself.

Elisp Feature: process filters and inhibit-quit

A *process filter* is a function attached to an Emacs subprocess or network connection, called whenever output arrives:

```
(make-process :name "x" :command '("cat") :filter (lambda (proc chunk) ...))
```

Two properties of filters bite people:

1. **They are not re-entrant with the command loop.** While a filter runs, Emacs is not reading keyboard input. Blocking in a filter freezes the editor's interactivity even though `accept-process-output` may still let other process output through.
2. **`inhibit-quit` is t.** Normally C-g sets the variable `quit-flag`, and the next time Emacs checks it, execution unwinds. Inside a filter that check is suppressed, so C-g cannot rescue a stuck filter.

The idiomatic escape is the one used here: do not block. Start the work, return, and complete the request from a callback (a timer, a hook, a sentinel) once the event loop is running again.

3.5 Decision: two JSON gotchas, encoded as helpers

Emacs's `json-encode` maps an *empty alist* to `null`, because an empty list and `nil` are the same object. MCP schemas need `{}`. The project's answer is a tiny constructor (`mcp-emacs-server.el:63`):

```
(defun mcp-emacs-server--obj (&rest pairs)
  "Build a JSON object value from PAIRS preserving key order.
With no PAIRS, return an empty hash table so it encodes as `{}`
rather than `null` (an empty alist would)."
  (if (null pairs)
      (make-hash-table :test 'equal)
      (let (acc)
        (while pairs
          (push (cons (pop pairs) (pop pairs)) acc))
        (nreverse acc))))
```

Two things are going on. First, the empty case returns a hash table, which `json-encode` renders as `{}`. Second, the non-empty case builds an alist by pushing pairs and then reversing, which preserves the caller's key order — JSON objects are notionally unordered, but readable schemas and stable test fixtures both want deterministic output.

mcp-emacs-ide.el has its own near-identical pair (`--obj` and `--empty-object`) because that module is intentionally decoupled from the HTTP server; the duplication is the price of the isolation.

The second gotcha is in `web-server` itself, and `AGENTS.md` flags it: the HTTP verb is the *key* of the headers alist, not a `:method` entry. Hence:

```
(let* ((is-post (assoc :POST headers)) ...)
```

rather than the `(alist-get :method headers)` one would expect.

3.6 Decision: registries, not cond chains

Tools and resources are lists of plists. Dispatch is a `seq-find` over the list comparing `:name`. Adding a tool means appending a plist; it never means editing a dispatcher.

Optional modules extend the set without touching the core list by pushing onto a separate variable:

```
(defvar mcp-emacs-server-extra-tools nil
  "Tool descriptors registered by optional modules (e.g. orgspec).")
```

```
(defun mcp-emacs-server--all-tools ()
  (append mcp-emacs-server--tools mcp-emacs-server-extra-tools))
```

`orgspec-mcp.el` calls `(orgspec-mcp-register)` at load time, and that is the entire integration: requiring the file exposes eight new tools.

Elisp Feature: plists vs alists

Both are association structures, and this codebase uses both, on purpose.

A **plist** (“property list”) is a flat list of alternating keys and values:

```
(list :name "open_file" :description "..." :handler (lambda (args) ...))
```

Read with `(plist-get tool :name)`. Keys are compared with `eq`, so they are almost always keyword symbols. Plists are used here for *code-authored* records: tool descriptors, event structures, diff-session state.

An **alist** (“association list”) is a list of cons cells:

```
'((path . "/tmp/x.el") (save . t))
```

Read with `(alist-get 'path args)`. Keys are compared with `eq` by default (hence symbol keys) or with a supplied test. Alists are used here for *data-derived* records: anything that came out of `json-read-from-string`, which produces symbol-keyed alists when `json-object-type` is `'alist`.

The rule of thumb the codebase follows: plists for things written by hand in source, alists for things parsed from the wire.

3.7 Decision: return status strings, not errors, for empty results

`AGENTS.md` is explicit: “return a plain status string for empty results (e.g. ‘No active selection’), not an error.”

The reason is that MCP surfaces an error as a protocol-level failure, which an agent tends to treat as “the tool is broken” rather than “the answer is nothing”. So `get_selection` with no region returns the *string* “No active selection”, and only genuine failures — unknown tool, unparseable arguments, a signal from the helper — become JSON-RPC errors, wrapped as `-32603` (internal error) or `-32601` (method not found).

You can see this convention shaping helper signatures throughout `mcp-emacs.el`: functions return either useful text or a human-readable explanation of why there is none, and the wrapping is done at the tool descriptor with `or`:

```
:handler (lambda (_args)
            (or (mcp-emacs-get-selection) "No active selection"))
```

4 mcp-emacs.el — The Helper Layer

This is the largest file in the project (1318 lines) and the least surprising one. It contains the functions that actually touch Emacs state. No HTTP, no JSON, no protocol — just Emacs that reads and writes buffers, and returns strings.

4.1 The foundational abstraction: “the current buffer”

Every buffer-reading tool starts from one three-line function (`mcp-emacs.el:40`):

```
(defun mcp-emacs--current-buffer ()  
  "Return the buffer associated with the currently selected window."  
  (window-buffer (frame-selected-window (selected-frame))))
```

This is more careful than it looks. The obvious implementation would be `(current-buffer)`, but `current-buffer` returns whatever buffer the *code* is currently operating in — and when the code is running inside a process filter, that may be an arbitrary internal buffer, not the one the human is looking at.

Going `frame` → `selected window` → `its buffer` asks a different question: “what is displayed where the user’s attention is?” That is the question the agent means when it says “the current buffer”.

Every helper then opens with:

```
(with-current-buffer (mcp-emacs--current-buffer)  
  ...)
```

Elisp Feature: `with-current-buffer` and the `buffer-local` world

Emacs has a notion of “the current buffer” as ambient state. Most buffer functions (`point`, `buffer-substring`, `insert`, `save-buffer`) operate on it implicitly. `with-current-buffer` is a macro that binds it for the duration of a body and restores it afterwards, even on a non-local exit:

```
(with-current-buffer some-buffer  
  (buffer-substring-no-properties (point-min) (point-max)))
```

The related macro `save-excursion` saves and restores *point and mark* within the current buffer — used when a function needs to move around to gather information but must not disturb where the user’s cursor is. You will see both stacked constantly in this file:

```
(with-current-buffer (mcp-emacs--current-buffer)  
  (save-excursion  
    (org-back-to-heading t)  
    ...))
```

There is also `save-restriction`, which preserves narrowing (see `mcp-emacs-org-task--find-item` at line 706, which narrows to a subtree to scope a search and then

unwinds).

4.2 Reading state

The read-only helpers are the simplest code in the project:

```
(defun mcp-emacs-get-buffer-content ()
  "Return the full text of the current buffer without properties."
  (with-current-buffer (mcp-emacs--current-buffer)
    (buffer-substring-no-properties (point-min) (point-max))))

(defun mcp-emacs-get-selection ()
  "Return the active region text for the current buffer, or nil."
  (with-current-buffer (mcp-emacs--current-buffer)
    (when (use-region-p)
      (buffer-substring-no-properties (region-beginning) (region-end)))))
```

Note `buffer-substring-no-properties` rather than `buffer-substring`. Emacs strings can carry *text properties* — font-lock faces, keymap overlays, invisible markers. Sending those over JSON is both useless and, for large fontified buffers, expensive. The `-no-properties` variant strips them.

`get_selection` returning `nil` rather than `""` for “no region” is what lets the tool descriptor’s or produce the status string.

4.3 Position arithmetic

Agents speak in 1-based line and column numbers; Emacs speaks in buffer positions (1-based character offsets). The converter (`mcp-emacs.el:65`) is defensive about both:

```
(defun mcp-emacs--line-column-position (line column)
  "Return buffer position for 1-based LINE and COLUMN."
  (let ((line (if (and (integerp line) (> line 0)) line 1))
        (column (if (and (integerp column) (> column 0)) column 1)))
    (save-excursion
      (goto-char (point-min))
      (forward-line (1- line))
      (let* ((line-start (line-beginning-position))
             (line-end (line-end-position))
             (line-length (- line-end line-start))
             (offset (min (max 0 (1- column)) line-length)))
        (+ line-start offset)))))
```

Three defensive moves worth naming:

1. Non-integer or non-positive inputs silently become 1 rather than signalling. An agent sending null for a column should get a sane edit, not an error.
2. `(min ... line-length)` clamps a too-large column to the end of the line rather than spilling onto the next line.
3. `save-excursion` means computing a position does not move the user’s point.

4.4 Editing

`mcp-emacs-edit-file-region` (line 183) is the coordinate-based editor. Its shape is representative of the whole file:

```
(let* ((buffer (or (get-file-buffer path) (find-file-noselect path)))
      ...)
  (with-current-buffer buffer
    (let ((start (mcp-emacs--line-column-position s-line s-col))
          (end   (mcp-emacs--line-column-position e-line e-col)))
      (when (> start end)
        (error "Start position %d:%d must be before end position %d:%d" ...))
      (let ((inhibit-read-only t))
        (save-excursion
          (goto-char start)
          (delete-region start end)
          (insert text)))
      (when save (save-buffer))
      (format "Edited %s at %d:%d-%d:%d%s" ...))))
```

Points of interest:

- `(or (get-file-buffer path) (find-file-noselect path))` — prefer an already-open buffer (which may hold unsaved changes the agent should be editing) over reading the file fresh. `find-file-noselect` opens without displaying.
- `inhibit-read-only` is bound to `t` so the edit works even in a buffer the user has marked read-only. This is a deliberate choice: the agent was asked to edit, so it edits.
- Saving is opt-in. The default leaves the buffer modified, so the human sees a dirty buffer and can review or undo.
- The return value is a human-readable confirmation, not a status code.

4.5 Diagnostics: two backends, one interface

Emacs has two competing on-the-fly checkers: Flycheck (a package) and Flymake (built in, and the channel `eglot` uses for LSP diagnostics). The helper layer detects which is live and normalises both to lines of text (`mcp-emacs.el`:298):

```
(defun mcp-emacs--buffer-diagnostics (buffer)
  (with-current-buffer buffer
    (cond
      ((bound-and-true-p flycheck-mode)
       (let ((errors (flycheck-current-errors)))
         (when errors (mapconcat (lambda (err) (format "%s: %s [%s]" ...))
                                errors "\n"))))
      ((bound-and-true-p flymake-mode)
       (let ((diags (flymake-diagnostics)))
         (when diags (mapconcat (lambda (d) (format "%s: %s: %s" ...))
                                diags "\n"))))
      (t nil))))
```

Returning `nil` for “no diagnostics” and for “no checker” lets the caller distinguish the two cases and produce different messages:

```
(or (mcp-emacs--buffer-diagnostics buffer)
    (with-current-buffer buffer
      (if (or (bound-and-true-p flycheck-mode) (bound-and-true-p flymake-mode))
          "No diagnostics in buffer"
          "Neither Flycheck nor Flymake is active"))))
```

Project-wide aggregation (`mcp-emacs-get-project-diagnostics`, line 343) walks (`buffer-list`), keeps buffers whose file is under the project root, and concatenates their sections. The docstring is honest about the limitation: “Files that are not open in a buffer are not covered.” This is a real constraint of the approach — diagnostics only exist for buffers a checker has run on — and documenting it beats pretending otherwise.

Elisp Feature: `when-let`, `if-let`, and friends

```
(when-let ((lines (mcp-emacs--buffer-diagnostics buf)))
  (format "%s:\n%s" buffer-file-name lines))
```

`when-let` binds a set of variables like `let*`, but aborts and returns `nil` the moment any binding evaluates to `nil`. It is the Elisp answer to nested null checks. `if-let` adds an `else` branch.

A note on style: newer Emacs deprecates the bare `when-let` in favour of `when-let*` (the plain form’s semantics were subtly different in old versions). `mcp-emacs-run.el`:157 carries a TODO acknowledging exactly this.

4.6 Navigation via imenu

`goto_line` accepts either coordinates or a function name. The function-name path goes through `imenu`, Emacs’s language-agnostic symbol index (`mcp-emacs.el`:87):

```
(defun mcp-emacs--find-imenu-position (target entries)
  "Find TARGET in imenu ENTRIES and return its position."
  (catch 'mcp-emacs--found
    (dolist (entry entries nil)
      (cond
        ((or (null entry)
              (and (stringp (car entry)) (string-prefix-p "*" (car entry))))
         nil)
        ((and (consp entry) (imenu--subalist-p entry))
         (let ((child (mcp-emacs--find-imenu-position target (cdr entry))))
           (when child (throw 'mcp-emacs--found child))))
        ((and (consp entry) (stringp (car entry)) (string= target (car entry)))
         (let ((pos (mcp-emacs--normalize-position (cdr entry))))
           (when pos (throw 'mcp-emacs--found pos))))))))))
```

Two Elisp features carry this function.

Elisp Feature: `catch` / `throw` for early exit

Elisp has no `return`. `catch` establishes a labelled exit point; `throw` jumps to it with a value:

```
(catch 'found
  (dolist (x list)
    (when (good-p x) (throw 'found x)))))
```

The tag is compared with `eq`, so it is conventionally a symbol — and here a deliberately namespaced one (`mcp-emacs--found`) to avoid colliding with a `catch` in surrounding code.

This is a *non-local exit*: it unwinds through intervening frames, which is exactly what makes the recursive `imenu` search work. The recursive call throws from arbitrary depth straight to the top-level `catch`.

Related: `unwind-protect` guarantees cleanup on any exit path, including a `throw` or a signal. `mcp-emacs-apply-diff` uses it to guarantee the temporary proposal buffer is killed.

The other subtlety is that `imenu` positions are not uniform. Depending on the major mode's `imenu` backend, an entry's `cdr` may be an integer, a marker, an overlay, or a cons whose `car` is a position. Hence `mcp-emacs--normalize-position` (line 78):

```
(cond
  ((markerp pos) (marker-position pos))
  ((overlayp pos) (overlay-start pos))
  ((numberp pos) pos)
  ((and (consp pos) (numberp (car pos))) (car pos))
  (t nil))
```

Elisp Feature: markers and overlays

A **marker** is a position that moves with the text. Insert 10 characters before a marker and its `marker-position` increases by 10. Buffer positions (plain integers) do not track edits, which is why anything stored across an edit should be a marker.

An **overlay** is a region (start, end) with properties attached, used for transient decoration — the highlight on a Flycheck error, an `ediff` difference region. Like markers, overlays move with edits.

Both need explicit conversion to an integer before arithmetic, hence the normaliser.

4.7 Org helpers

The Org functions are thin wrappers that mostly exist to handle the “not in an Org buffer” and “no task here” cases gracefully:

```
(defun mcp-emacs-get-current-clocked-task ()
  (when (org-clocking-p)
    (let ((marker (and (boundp 'org-clock-marker) org-clock-marker)))
      (when (and marker (marker-buffer marker))
        (with-current-buffer (marker-buffer marker)
          (save-excursion
            (goto-char marker)
            (org-back-to-heading t)
            (org-get-heading t t t t)))))))
```

`org-get-heading` takes four booleans controlling whether to strip TODO keyword, priority, tags, and comment marker respectively; `t t t t` means “just the text”. `org-clock-marker` is a marker into whichever buffer holds the clocked task — possibly a different file entirely — hence the marker-buffer dance.

4.8 Environment and self-diagnosis

`mcp-emacs-diagnose` (line 500) assembles a report: Emacs version, system type, Doom version if present, the full `exec-path`, whether a handful of language servers are on `PATH`, and any active `lsp-mode` workspaces.

This exists because the single most common class of “the AI can’t help me” problem is an Emacs environment problem — a GUI Emacs launched from Finder with a `PATH` that lacks `~/local/bin`, so no LSP starts, so no diagnostics exist, so every tool returns “nothing found”. The `diagnose-emacs` skill (Chapter 19) exists to route that class of problem here rather than into a fruitless code hunt.

The LSP workspace collector is worth a look for its de-duplication idiom:

```
(let ((seen (make-hash-table :test 'eq))
      (entries '()))
  (dolist (buf (buffer-list))
    (with-current-buffer buf
      (when (and (boundp 'lsp-mode) lsp-mode (fboundp 'lsp-workspaces))
        (dolist (ws (lsp-workspaces))
          (unless (gethash ws seen)
            (puthash ws t seen)
            (push (mcp-emacs--describe-workspace ws) entries))))))
    (nreverse entries))
```

Many buffers share one workspace object, so identity-keyed (`:test 'eq`) hashing is both correct and cheap. The `push` + `nreverse` pair is the standard Emacs way to build a list in order without repeated `append`.

4.9 Project tools with a projectile fallback

`project.el` is built in; `projectile` is a popular package with a different API. The project helpers prefer `projectile` *when it is already loaded* and fall back to `project.el` otherwise (line 1254):

```
(if (featurep 'projectile)
    (let ((root (projectile-project-root))) ...)
    (if (not (require 'project nil t))
        "project.el is not available"
        ...))
```

Note the asymmetry: `projectile` is checked with `featurep` (already loaded — do not load it on the agent’s behalf), while `project` is required with `NOERROR`. Combined with the `declare-function` declarations at line 1236, this introduces no hard dependency in either direction.

`mcp-emacs-switch-project` (line 1273) deserves a note. To “switch project” in Emacs is not a well-defined operation — there is no global current project, only whatever `project-current` computes from the current buffer’s `default-directory`. So the implementation opens the root in `dired`:

```
(when (fboundp 'project-remember-project)
  (ignore-errors (project-remember-project (project-current nil dir))))
(dired dir)
(format "Switched project to %s" dir)
```


Opening a dirent buffer at the root makes it the selected window's buffer, so every subsequent tool that resolves through `mcp-emacs--current-buffer` computes the new project. It is a pragmatic answer that makes the *observable* behaviour match the tool's name.

5 mcp-emacs-server.el — The MCP Server

686 lines, of which roughly 350 are the tool registry — a flat list of descriptors. The interesting machinery is the other half.

5.1 Schema construction

MCP tools advertise a JSON Schema for their inputs. Writing those by hand as nested alists is unreadable, so the file defines three combinators (mcp-emacs-server.el:63–90):

```
(defun mcp-emacs-server--no-args ()
  (mcp-emacs-server--obj "type" "object" "properties" (mcp-emacs-server--obj)))

(defun mcp-emacs-server--prop (type description)
  (mcp-emacs-server--obj "type" type "description" description))

(defun mcp-emacs-server--position-prop (which)
  (mcp-emacs-server--obj
   "type" "object"
   "description" (format "%s of the replacement region" (capitalize which))
   "properties" (mcp-emacs-server--obj
                  "line" (mcp-emacs-server--prop "integer" "1-based line number")
                  "column" (mcp-emacs-server--prop "integer" "1-based column number"))
   "required" (vector "line" "column")))
```

Note (vector "line" "column") rather than ("line" "column"). This is the second JSON-encoding trap in Emacs: a *list* encodes as a JSON object if it looks like an alist and as null if empty, whereas a *vector* always encodes as a JSON array. Anything that must be a JSON array in the output is built with vector or vconcat.

5.2 The tool registry

A descriptor is a plist:

```
(list :name "open_file"
      :description "Open a file in the current Emacs window"
      :schema (mcp-emacs-server--obj
                "type" "object"
                "properties" (mcp-emacs-server--obj
                              "path" (mcp-emacs-server--prop
                                       "string" "Absolute path to the file to open"))
                "required" (vector "path"))
      :handler (lambda (args)
                  (let ((path (alist-get 'path args)))
```

```
(mcp-emacs-open-file path)
(format "Opened file: %s" path)))
```

The handler is a closure of one argument, the parsed arguments alist. It does argument extraction and result formatting; the actual work is delegated to a `mcp-emacs-*` helper. `AGENTS.md` names this as the rule for adding tools, and the registry follows it consistently — the handlers are one to five lines each.

Boolean arguments are extracted with `(eq (alist-get 'save args) t)` rather than a truthiness test. That is because `json-read-from-string` decodes JSON `false` as the symbol `:json-false` by default, which is *truthy* in Emacs. Comparing to `t` explicitly avoids the trap.

The `eval` tool is the one place where a handler is doing real work:

```
:handler (lambda (args)
  (let* ((expr (alist-get 'expression args))
        (form (car (read-from-string
                     (format "(with-current-buffer (mcp-emacs--current-buffer) (progn %s))"
                             expr)))))
    (format "%s" (eval form t))))
```

The expression string is wrapped so it evaluates in the visible buffer's context, `read-from-string` parses it into a form, and `(eval form t)` evaluates it with lexical binding enabled (the `t`). This is, obviously, an arbitrary-code-execution tool; it is also the escape hatch that made the `orgspec` MVP drivable before typed tools existed.

5.3 Resources

MCP has a second concept alongside tools: *resources*, addressable read-only content identified by URI. The registry is smaller (`mcp-emacs-server.el`:445):

```
(list :uri "org-tasks://all"
      :name "org-tasks"
      :description "All TODO items from org-mode agenda files"
      :mime "text/plain"
      :reader (lambda () (mcp-emacs-get-org-tasks)))
```

Three are defined: all Org agenda TODOs, the `*Messages*` buffer, and the `*Warnings*` buffer. The dispatch is identical in shape to tools — `seq-find` on `:uri`, call `:reader`, wrap the string.

5.4 JSON-RPC dispatch

`mcp-emacs-server--dispatch` (line 547) is a `cond` over the method name:

```
(cond
  ((null method) nil)
  ((string= method "initialize") ...)
  ((string= method "tools/list") ...)
  ((string= method "tools/call")
   (condition-case err
     (mcp-emacs-server--result id (mcp-emacs-server--tools-call params))
     (error (mcp-emacs-server--error id -32603 (error-message-string err)))))
```

```
((string= method "resources/list") ...)
((string= method "resources/read") ...)
((null id) nil)
(t (mcp-emacs-server--error id -32601 (format "Method not found: %s" method))))
```

Three details:

- The `(null id)` clause near the end handles **notifications**. In JSON-RPC, a request without an id expects no reply. Returning `nil` from dispatch signals the HTTP layer to answer 202 Accepted with an empty body.
- Errors from tool and resource execution are caught and converted to -32603 (internal error). An unknown method is -32601.
- `initialize` echoes a fixed protocol version constant, "2025-06-18".

Elisp Feature: `condition-case`

Elisp’s `try/catch` for *signals* (as opposed to `catch/throw`, which is for control flow):

```
(condition-case err
  (risky-thing)
  (user-error (message "user problem: %s" (error-message-string err)))
  (error (message "anything else: %s" (error-message-string err))))
```

The first form is the protected body. Each subsequent clause names an error symbol (or a list of them) and a handler body; `err` is bound to the error object, which is a cons of `(SYMBOL . DATA)`.

Error symbols form a hierarchy: `user-error` is a child of `error`, so an `error` clause catches both — clause order matters. `error-message-string` renders an error object into the message the user would have seen in the echo area.

`condition-case` appears roughly 20 times in this codebase, almost always in the same shape: catch the error, return its message as the tool’s text result rather than propagating it. That is the “status string, not error” convention from §3.7 applied at the boundary.

5.5 The HTTP handler and the deferral

`mcp-emacs-server--handler` is the web-server callback. Two request shapes reach it: `/mcp`, the JSON-RPC surface, and `/permission-gate`, which is neither JSON-RPC nor a tool — see “The permission gate route” below.

```
(defun mcp-emacs-server--handler (request)
  (with-slots (process headers) request
    (let* ((is-post (assoc :POST headers))
          (path (cdr is-post))
          (rpc (ignore-errors (mcp-emacs-server--parse-body request))))
      (cond
        ;; Held open like an async tool, because a human answers it.
        ((and is-post (equal path "/permission-gate"))
         (if (and rpc (mcp-emacs-server--permission-gate rpc process))
             (throw 'close-connection :keep-alive)
             ;; An unusable payload still gets a well-formed deny, not a 400:
             ;; the hook parses this reply, and a gate must fail closed.
```

```

(ws-response-header process 200 ...)
(process-send-string process (json-encode ...)))
((and is-post rpc)
 (if (and (equal (alist-get 'method rpc) "tools/call")
              (alist-get 'id rpc)
              (mcp-emacs-server--tools-call-async
               (alist-get 'params rpc)
               (alist-get 'id rpc)
               (lambda (response)
                 ;; Cleanup sits outside the liveness check, so a client
                 ;; that gave up waiting still frees the socket.
                 (unwind-protect
                  (when (process-live-p process)
                    (let ((json (json-encode response)))
                      (ws-response-header process 200
                                           ("Content-Type" . "application/json"))
                      (process-send-string process json)))
                  (mcp-emacs-server--release-connection process))))))
      (throw 'close-connection :keep-alive)
      (let* ((response (mcp-emacs-server--dispatch rpc)))
        (if response
            (progn (ws-response-header process 200 ...)
                    (process-send-string process (json-encode response)))
            (ws-response-header process 202 ...))))))
(t
 (ws-response-header process 200 ("Content-Type" . "text/plain"))
 (process-send-string process "mcp-emacs server\n"))))

```

Read the async branch carefully, because the control flow is unusual.

`mcp-emacs-server--tools-call-async` returns non-`nil` **only if** the named tool has an `:async-` handler, in which case it has already started the work and stashed the completion callback. That callback closes over `process` — the live TCP connection — and writes the response directly when it eventually fires.

If the tool was `async`, the handler does not return a response. It throws `'close-connection` with the value `:keep-alive`, which is web-server's protocol for “do not close this socket; I will finish it myself”. Emacs returns to its event loop, the `ediff` panel becomes interactive, the human answers, the callback fires, and only then is the HTTP response written and the process deleted.

Elisp Feature: `with-slots` and `EIEIO`

```
(with-slots (process headers) request ...)
```

`web-server` represents a request as an `EIEIO` object — Emacs's CLOS-inspired object system, from the `eieio` library. `with-slots` binds slot names as *symbol macros*, so `process` inside the body reads (and could write) the object's `process` slot without an explicit accessor call.

The single-slot accessor is `oref`:

```
(oref request body)
```

and `slot-boundp` checks whether a slot has been assigned at all — which the body parser needs, because a GET request has no body slot value:

```
(let ((body (and (slot-boundp request 'body) (oref request body))))
  ...)
```

EIEIO appears in two places: here, at the web-server boundary, where the library chose it; and deliberately, as the backend class hierarchy in `agent-backend.el` (Chapter 8). Everywhere else the codebase uses `cl-defstruct` or plain plists.

5.6 Releasing a kept-alive connection

Holding a socket open makes the *caller* responsible for closing it, and the first version of that cleanup leaked. `mcp-emacs-server--release-connection` exists because deleting the process is only half the job:

```
(defun mcp-emacs-server--release-connection (process)
  (when-let* ((server (plist-get (process-plist process) :server)))
    (setf (ws-requests server)
          (cl-remove-if (lambda (r) (eq process (ws-process r)))
                        (ws-requests server))))
  (when (process-live-p process)
    (delete-process process)))
```

The comment records the mechanism precisely: `ws-filter` only prunes its requests list on the path that deletes the process itself, and it leaves a kept-alive request with `ws-active` set. So a request whose callback never deleted the process leaked the entry forever — the list only grew, every later filter call scanned the dead entries, and `ws-stop` still held their processes.

Doing both here is what makes *abandoning* a review cost nothing. That is the case worth noticing: the human closing an ediff without answering, or an MCP client timing out, is normal operation rather than an error path.

Hence the `unwind-protect` in the handler’s callback above. Answering a dead client is pointless, but releasing the connection is not, so the cleanup sits outside the liveness check.

5.7 The permission gate route

`/permission-gate` is the one HTTP surface that is not JSON-RPC and not a tool, and the commentary at the top of the file explains the second part:

A *tool* would be worse than useless here: the model being gated could call it and approve itself. So it lives off the registry, `tools/list` never mentions it, and the model has no name for it.

The Claude CLI’s `PreToolUse` hook posts a pending tool call here and blocks on the reply. Emacs raises the decision buffer from Chapter 8 and answers when the human does — so the route uses exactly the same keep-alive deferral as `apply_diff`, for the same reason: a human is in the loop.

The reply is the hook’s own output shape rather than anything MCP-flavoured:

```
(defun mcp-emacs-server--gate-decision (decision reason)
  (mcp-emacs-server--obj
   "hookSpecificOutput"
   (mcp-emacs-server--obj
```

```
"hookEventName" "PreToolUse"
"permissionDecision" (if (eq decision 'allow) "allow" "deny")
"permissionDecisionReason" (or reason "")))
```

Note the asymmetry in that `if`: only `allow` allows. An unrecognised decision value denies rather than falling open — the same rule the shell script and the timeout follow.

Failing to *raise* the gate is itself handled, because a gate that cannot ask must still answer:

```
;; A gate that cannot even be raised must still answer, or the CLI
;; waits out its own timeout with no explanation.
(error
  (funcall reply 'deny
    (format "Permission gate could not be raised in Emacs: %s"
      (error-message-string err)))
  t)
```

The `permissionDecisionReason` string reaches the model verbatim, which is why every one of them is written for the model rather than for a log: they say what happened and whether retrying will help. Chapter 24 covers the hook side, including the fail-closed shell script.

5.8 Body parsing and the JSON decode configuration

```
(defun mcp-emacs-server--parse-body (request)
  (let ((body (and (slot-boundp request 'body) (oref request body))))
    (when (and body (not (string-empty-p body)))
      (let ((json-object-type 'alist)
            (json-array-type 'vector)
            (json-key-type 'symbol))
        (json-read-from-string body)))))
```

The three `let`-bound variables configure the older `json.el` reader: objects become symbol-keyed alists, arrays become vectors. This is why every handler uses `(alist-get 'path args)` with a quoted symbol.

Elisp Feature: dynamic binding as configuration

`json-object-type` is a *special* (dynamically scoped) variable. `let`-binding it does not create a new lexical variable shadowing the old one; it temporarily changes the global value for the dynamic extent of the `let`, including inside any function called from within it.

That is how these libraries take configuration: you bind the knob around the call. The same idiom appears with `inhibit-read-only`, `default-directory`, `process-environment`, `display-buffer-alist`, `ediff-window-setup-function`, and `org-inhibit-startup` throughout this codebase.

Under lexical-binding: `t` (which every file here declares), ordinary `let` bindings are lexical, but variables declared with `defvar`/`defcustom` remain special and continue to bind dynamically. That dual regime is why files sometimes contain a bare `(defvar ediff-control-buffer)` with no value — it declares the symbol special for the byte-compiler without defining it.

Note also `(ignore-errors (mcp-emacs-server--parse-body request))` at the call site: a malformed body yields `nil` rather than an error, and `nil` falls through to the health-response

branch.

6 The Interactive Diff Review

This is the most intricate piece of the codebase, and the one that best rewards close reading. Everything in this chapter lives in `mcp-emacs.el`, lines 946–1229.

6.1 The problem

An agent wants to change a file. The project’s position is that a human should see and approve that change *in the editor*, with the ability to hand-edit the proposal before accepting it. Emacs already has the perfect UI for this: `ediff`.

The complications are all in the plumbing:

1. `ediff` is an interactive, stateful UI driven by keys in a control buffer. The tool call arrives in a process filter with no command loop (§3.4).
2. `ediff` rearranges windows aggressively, and will *fail outright* if the selected window is a side window or a dedicated window — which it very often is in a Doom setup with Treemacs open.
3. The review must be answerable exactly once, whether the human accepts, rejects, presses plain `q`, or lets it time out.
4. Two callers need it: the synchronous `apply_diff` tool and the deferred IDE `openDiff` flow (Chapter 9), with different completion semantics.

6.2 The shared review function

`mcp-emacs--ediff-review` (line 990) takes five required arguments plus two optional ones:

```
(defun mcp-emacs--ediff-review (buffer-a buffer-b entry-content result
                                &optional on-resolve tab-name)
```

- `buffer-a` — the buffer visiting the real file.
- `buffer-b` — a temporary buffer holding the proposed content.
- `entry-content` — a snapshot of `buffer-a`’s text taken *before* the review.
- `result` — a one-element list used as a mutable cell.
- `on-resolve` — an optional callback for the async caller.
- `tab-name` — an optional name for the control buffer.

Elisp Feature: a cons cell as a mutable box

```
(let ((result (list nil)))
  ...
  (setcar result 'applied)
  ...
  (car result))
```

Elisp closures capture variables by reference under lexical binding, so a closure *can* setq an outer variable. But when the value must be visible to a *different* closure cre-

ated at a different time — here, the accept key binding, the reject key binding, the quit hook, and the polling loop all need to agree — passing a shared mutable container is clearer and avoids questions about which closure captured which binding.

(list nil) creates a one-element list; setcar mutates its first cell in place; every holder of the same cons sees the change. This is the Elisp equivalent of a Ref or a one-element array.

The same idiom appears a second time in mcp-emacs-apply-diff-async as (done (list nil)), guarding single delivery against a race between the quit hook and the timeout timer.

6.3 Saving and restoring the window layout

The first thing the function does, before touching any window:

```
(let ((winconf (current-window-configuration))
      control
      (ediff-window-setup-function 'ediff-setup-windows-plain)
      (ediff-split-window-function 'split-window-horizontally))
```

current-window-configuration captures the complete arrangement of windows in the frame — sizes, buffers, point positions — as an opaque object. set-window-configuration restores it. Capturing *first* is essential, because the very next thing the function does is destroy windows:

```
(dolist (window (window-list))
  (when (window-parameter window 'window-side)
    (ignore-errors (delete-window window))))
```

Side windows (Treemacs, an imenu-list sidebar, a Doom popup) cannot be split. ediff-buffers tries to split the selected window; if that fails, the setup aborts *before the control buffer exists*, leaving no UI and no way to resolve the review. Deleting side windows up front prevents that.

The optional directional placement follows:

```
(unless (eq mcp-emacs-ediff-window-direction 'plain)
  (ignore-errors
   (select-window
    (display-buffer-in-direction
     buffer-a `((direction . ,mcp-emacs-ediff-window-direction))))))
```

Elisp Feature: window parameters, side windows, and display-buffer

Emacs windows carry arbitrary key/value *parameters*. The window-side parameter marks a window as belonging to the frame’s side-window group — these are the docked panels that packages like Treemacs create. Side windows have restrictions: they cannot be split, and delete-other-windows skips them.

display-buffer is the central “show me this buffer” entry point, and it is policy-driven rather than imperative. It consults display-buffer-alist (user rules), then a list of *action functions* to try in order. Passing an action explicitly:

```
(display-buffer buf `((display-buffer-in-direction)
                      (direction . right)
                      (window-width . 120)))
```

means “use this specific placement strategy with these parameters”.

This matters for a package that must coexist with a user’s framework. See `mcp-emacs-run--display-popup` (Chapter 10) for the technique of *prepending* a rule to a locally-bound `display-buffer-alist` so it wins over Doom’s popup system.

6.4 The startup hook: bindings and the quit hook

`ediff-buffers` accepts a list of *startup hook* functions, run once setup completes and the control buffer exists. Everything interesting is installed there:

```
(ediff-buffers
  buffer-a buffer-b
  (list (lambda ()
    (setq control ediff-control-buffer)
    (when tab-name
      (let ((wanted (format "*mcp diff: %s*" tab-name)))
        (unless (get-buffer wanted)
          (with-current-buffer ediff-control-buffer
            (rename-buffer wanted))))))
    (with-current-buffer ediff-control-buffer
      (local-set-key (kbd "C-c C-c") (lambda () (interactive) ...accept...))
      (local-set-key (kbd "C-c C-k") (lambda () (interactive) ...reject...))
      (setq-local ediff-quit-hook (list (lambda () ...)))
      (message "mcp diff: C-c C-c to accept, C-c C-k (or q) to reject")))))
```

The comment explaining the rename is instructive: `ediff` derives its own control-buffer name and clobbers `ediff-control-buffer-suffix` during setup, so the startup hook is the only reliable place to inject a custom name. And if the name is taken (two concurrent reviews of the same tab), `ediff`’s own unique name is kept rather than erroring.

The accept binding is the interesting one, because it does not blindly copy:

```
(defun mcp-emacs--apply-diff-accept (buffer-a buffer-b entry-content result)
  (when (buffer-live-p buffer-a)
    (with-current-buffer buffer-a
      (when (string= (buffer-substring-no-properties (point-min) (point-max))
                    entry-content)
        (erase-buffer)
        (insert-buffer-substring buffer-b))))
  (setcar result 'applied))
```

If `buffer-a` still matches `entry-content` — the human just looked and approved — the proposal is copied in wholesale. But if the human *edited* buffer A during the review (`ediff` allows this; it is the point of reviewing in a real editor), their edits are preserved and the proposal is discarded. The outcome is still applied, and the caller returns buffer A’s *actual* final content. “Accept” means “accept what is now in buffer A”, not “accept what the agent proposed”.

6.5 Single-path resolution

The quit hook is the linchpin:

```
(setq-local
 ediff-quit-hook
 (list (lambda ()
        (unless (car result)
          (setcar result 'rejected))
        (ignore-errors (set-window-configuration winconf))
        (when on-resolve (funcall on-resolve))))))
```

Three responsibilities, in order:

1. **Default to rejection.** If the human pressed plain q — never touching either binding — result is still nil, and this sets it to rejected. The safe default is “do not change the file”.
2. **Restore the layout.** Exactly once, on the single quit path, bringing back the side windows deleted during setup.
3. **Fire the resolution callback,** if the async caller supplied one.

Because ediff-quit-hook runs on *every* exit — accept, reject, or bare q — this is the one place that guarantees “exactly once” semantics. The accept and reject bindings record a decision and then call ediff-really-quit, which runs the hook. Nothing else calls on-resolve.

Elisp Feature: setq-local and buffer-local hooks

setq-local creates a *buffer-local binding* for a variable: the new value is visible only in the current buffer, while other buffers keep the global value.

Setting ediff-quit-hook buffer-locally in the control buffer means this review’s quit behaviour is scoped to this review. Two concurrent ediff sessions each get their own hook, closing over their own result cell and winconf. Setting it globally would make the second review clobber the first.

Note that this replaces rather than adds to the hook (it uses setq-local with an explicit list rather than add-hook). That is deliberate here: the review owns the quit path.

6.6 Failure during setup

The whole ediff-buffers call is wrapped:

```
(condition-case err
  (ediff-buffers ...)
  (error (ignore-errors (set-window-configuration winconf))
        (signal (car err) (cdr err))))
```

If setup fails before the quit hook could ever be installed, the layout is restored here and the error is re-signalled unchanged. signal takes the error symbol and its data, so (signal (car err) (cdr err)) re-raises the original error rather than wrapping it in a new one.

6.7 The synchronous caller

mcp-emacs-apply-diff (line 1168) is the version used when the tool is dispatched directly rather than over HTTP. It polls:

```
(let ((deadline (+ (float-time) secs)))
  (while (and (null (car result))
```

```
(< (float-time) deadline))
(accept-process-output nil mcp-emacs-org-task-wait-poll-interval)))
```

Elisp Feature: accept-process-output as a cooperative yield

(accept-process-output PROCESS SECONDS) waits for output from a process, or until the timeout, **running the event loop while it waits**. With nil as the process argument it waits on any process.

That makes it Elisp's cooperative yield. A busy while loop with sit-for or nothing at all would starve everything; a loop with accept-process-output lets timers fire, other process filters run, and — crucially for the org-task wait in Chapter 7 — other MCP tool calls be served while this one waits.

What it does *not* do is give you a command loop. Keyboard input is still not processed. That is precisely why the synchronous variant cannot work from a process filter, and why the async variant exists.

On timeout, the session is force-quit:

```
(when (and control (buffer-live-p control))
  (with-current-buffer control
    (if (fboundp 'ediff-really-quit)
        (ignore-errors (ediff-really-quit nil))
        (kill-buffer control))))
"Status: timeout"
```

And an unwind-protect guarantees the temporary proposal buffer is killed on every exit path.

6.8 The asynchronous caller

mcp-emacs-apply-diff-async (line 1095) has the same setup but no polling loop. Its docstring is the clearest statement of the problem in the whole codebase:

This exists because the synchronous variant cannot be resolved by a human when it is served from a process filter: Emacs runs filters with inhibit-quit bound to t and no command loop for the ediff control panel, so the review is displayed but its keys are dead.

The structure is a finish closure guarded by a done cell:

```
(let ((finish
      (lambda (outcome)
        (unless (car done)
          (setcar done t)
          (when timer (cancel-timer timer))
          (when (buffer-live-p buffer-b) (kill-buffer buffer-b))
          (funcall on-done outcome)))))
  (setq control
    (mcp-emacs--ediff-review
     buffer-a buffer-b entry-content result
     (lambda ()
       (funcall finish
        (if (eq (car result) 'applied)
```

```

      (format "Status: applied\n%s" ...buffer-a content...)
      "Status: rejected")))))
(setq timer
  (run-at-time secs nil
    (lambda ()
      ...force-quit ediff...
      (funcall finish "Status: timeout")))))

```

Two racers — the quit hook and the timer — both call `finish`; the `done` cell ensures only the first one is delivered. Note also the ordering comment: the timer is armed *after* `mcp-emacs--ediff-review` returns successfully, so a failed setup cannot leave a timer pointing at a session that never existed.

Elisp Feature: timers

```

(run-at-time SECS REPEAT FUNCTION &rest ARGS) ; one-shot when REPEAT is nil
(run-with-timer SECS REPEAT FUNCTION &rest ARGS) ; same thing, older name
(cancel-timer TIMER)

```

Timers fire from the event loop, meaning they get a proper execution context — unlike process filters, code in a timer *can* interact with the UI. That is why `claude-client.el`'s notes (recorded in the project's memory) recommend driving `ediff` via `run-at-time` rather than from an `emacsclient --eval` callback.

`run-at-time` returns a timer object; hold onto it if you may need to cancel.

7 The Org Task-Session Protocol

`mcp-emacs.el:649–911`. This is the cooperative human / AI workspace: an Org file both parties edit, live, in the same buffer.

7.1 The data format

The commentary block at line 649 defines it precisely:

- A “session task file” is an Org file whose **first top-level heading** is the task.
- That heading’s **TODO keyword** is the session status.
- Its **SESSION property** holds a session id (a plain label).
- Its **child headings** are the TODO checklist items.
- Progress notes are appended to the task’s body; new items as children.

And then the constraint that shapes every function below it:

The AI mutates only items it can identify (by ID / CUSTOM_ID property, else heading text) and never reorders, deletes, or rewrites human-authored items. All edits go through the live buffer; nothing saves to disk unless a tool is defined to save.

No tool is defined to save. The human owns the file.

7.2 The change token

```
(defun mcp-emacs-org-task--token ()  
  "Return the change token for the current buffer."  
  (buffer-chars-modified-tick))
```

Elisp Feature: modification ticks

Every buffer maintains two monotonically increasing counters:

- `buffer-modified-tick` — bumped by any change, including text-property changes and some non-textual modifications.
- `buffer-chars-modified-tick` — bumped only when the actual *characters* change.

The latter is the right choice here: font-lock re-fontifying the buffer must not look like the human edited it.

The tick is an opaque baseline. A caller reads it, later compares, and if it differs, something changed. It says nothing about *what* changed — that is what re-reading the session view is for.

7.3 Identifying items without owning them

Two helpers do the addressing:

```
(defun mcp-emacs-org-task--item-id ()
  "Return the ID or CUSTOM_ID property of the heading at point, or nil."
  (or (org-entry-get (point) "ID")
      (org-entry-get (point) "CUSTOM_ID")))

(defun mcp-emacs-org-task--find-item (ref)
  (when (and (stringp ref) (not (string-empty-p ref))
            (mcp-emacs-org-task--goto-task))
    (let ((task-level (org-current-level))
          (found nil))
      (save-restriction
        (org-narrow-to-subtree)
        (goto-char (point-min))
        (org-map-entries
         (lambda ()
           (when (and (not found)
                     (= (org-current-level) (1+ task-level))
                     (or (equal ref (mcp-emacs-org-task--item-id))
                         (equal ref (org-get-heading t t t t))))
             (setq found (point))))
         nil 'tree))
      (when found (goto-char found) t))))
```

REF matches a stable ID property if present, else exact heading text. The `(= (org-current-level) (1+ task-level))` guard restricts matching to *direct children* — a nested sub-item with the same text is not a match.

Elisp Feature: org-map-entries and narrowing

`org-map-entries` walks headings and calls a function at each one with point positioned on the heading:

```
(org-map-entries FUNC MATCH SCOPE)
```

MATCH is a tags/property/TODO query (nil for all); SCOPE limits the walk — 'tree means the current subtree, 'agenda means all agenda files, nil means the accessible portion of the buffer.

“Accessible portion” is the key phrase, and it interacts with **narrowing**. `org-narrow-to-subtree` restricts the buffer’s accessible region to the current subtree — `point-min` and `point-max` now refer to the subtree’s bounds, and code outside cannot be seen or modified.

`save-restriction` saves and restores the narrowing state, so the narrowing is a scoped tool rather than a lasting change to the user’s buffer. The combination `save-restriction + org-narrow-to-subtree + org-map-entries` is the standard way to walk one subtree safely.

7.4 Appending without disturbing

mcp-emacs-org-task-append-note (line 808) has to insert into the task's body *after* whatever the human wrote and *before* the first child heading:

```
(let ((task-level (org-current-level))
      (limit (save-excursion (org-end-of-subtree t t) (point))))
  (org-end-of-meta-data t)
  (if (re-search-forward org-heading-regexp limit t)
      (goto-char (match-beginning 0))
      (goto-char limit))
  (let ((inhibit-read-only t))
    (unless (bolp) (insert "\n"))
    (insert note "\n")))
```

org-end-of-meta-data with a non-nil argument skips past the planning line (SCHEDULED/DEADLINE), the property drawer, *and* the logbook — landing at the start of the actual body text. Then a bounded search for the next heading finds where the body ends. (unless (bolp) (insert "\n")) — “beginning of line predicate” — avoids gluing the note onto the end of an existing line.

mcp-emacs-org-task-append-item (line 835) is simpler: go to the end of the subtree and insert a heading one level deeper.

```
(org-end-of-subtree t t)
(unless (bolp) (insert "\n"))
(insert (make-string (1+ task-level) ?*) " " kw " " text "\n"))
```

(make-string N ?*) builds the stars. ?* is the character literal syntax — ?a is the character a, ?\n is newline.

7.5 Validating keywords against the user's own workflow

Every status-setting function checks:

```
(unless (member status org-todo-keywords-1)
  (user-error "Unrecognized status keyword: %s" status))
```

org-todo-keywords-1 is the flattened list of every TODO keyword configured in this buffer (including per-file #+TODO: lines). The agent cannot invent IN_PROGRESS if the human's workflow does not have it. Similarly, org-task-append-item defaults to (car org-todo-keywords-1) — whatever the user's first keyword is, not a hardcoded TODO.

Elisp Feature: user-error vs error

(error "...") signals a generic error — a bug, a broken invariant. (user-error "...") signals user-error, which Emacs treats as “the user did something wrong”: no backtrace is shown even with debug-on-error enabled.

The convention here is that anything caused by bad input from the agent or a missing file is a user-error, and it is then caught by a condition-case ... (user-error (error-message-string err)) wrapper and returned as a friendly text result. Genuine internal failures propagate to the dispatcher and become JSON-RPC errors.

7.6 Waiting for the human

`mcp-emacs-org-task-wait-for-change` (line 878) is what turns this from a set of file operations into a *protocol*:

```
(let* ((baseline (cond ((integerp token) token)
                       ((and (stringp token) (not (string-empty-p token)))
                        (string-to-number token))
                       (t nil)))
      (secs (min mcp-emacs-org-task-wait-max-timeout
                  (if (and (numberp timeout) (> timeout 0))
                      timeout
                      mcp-emacs-org-task-wait-default-timeout)))
      (deadline (+ (float-time) secs)))
  (while (and baseline
              (= (mcp-emacs-org-task--token) baseline)
              (< (float-time) deadline))
    (accept-process-output nil mcp-emacs-org-task-wait-poll-interval))
  (let ((changed (or (null baseline)
                    (/= (mcp-emacs-org-task--token) baseline))))
    (format "Changed: %s\n%s" (if changed "yes" "no")
            (mcp-emacs-org-task-read path))))
```

The token is accepted as an integer or a numeric string, because JSON clients are inconsistent about large-integer handling. A `nil` baseline, or one that is already stale, returns immediately — no hanging on a caller that forgot to pass a token.

The poll interval is 200ms, and every iteration yields to the event loop, so the human’s edits are processed and other tool calls are served. The timeout is capped at 300 seconds.

The result includes both the change flag *and* a fresh session view, so a single call is a complete “wait then read” round-trip.

7.7 Why this shape

Together these tools implement a loop the `org-task-loop` skill drives:

1. `org_task_session` — read the checklist, note the token.
2. Work the first actionable item.
3. `org_task_set_item_status` / `org_task_append_note` — report progress.
4. `org_task_wait_for_change` with the token — block.
5. The human edits the file: reprioritises, adds an item, marks something done, writes a note.
6. The wait returns with the new state. Go to 2.

The human never has to type into a chat box. They edit an Org file — the thing they were already using — and the agent notices. That is the “shared live workspace” idea in its most concrete form.

8 The Shared Agent Layer

Three files — `agent-backend.el`, `agent-prompt.el`, `agent-permission.el` — know nothing about Claude or opencode in particular. They exist because there came to be two chat clients, and the second one made it obvious which parts of the first were about *agents* and which were about *that* agent.

Read this chapter before Chapters 11 and 12: both clients are subclasses of what is defined here, and their chapters assume it.

8.1 Why a base class at all

The first client, `claude-client.el`, was written standalone. When `opencode-client.el` arrived it needed the same verbs — send a prompt, interrupt a turn, add a note, quit — over a completely different transport: a subprocess with an NDJSON stream on one side, HTTP with an SSE event stream on the other. Nothing about “send a prompt” is transport-specific, but every implementation of it was.

`agent-backend.el` names the verbs once, as generic functions over an EIEIO class:

```
(defclass agent-backend ()
  ((note-policy :initarg :note-policy :initform :steer ...)
   (session-id  :initarg :session-id  :initform nil ...)
   (buffer      :initarg :buffer      :initform nil ...)))
```

Only state *every* backend has lives in the base class. Everything else — `claude-client`’s process handle and stdout accumulator, `opencode`’s port and server registry — stays in the subclass.

Elisp Feature: EIEIO and `cl-defgeneric`

EIEIO is Emacs’s built-in CLOS-style object system: `defclass` defines a class with typed, documented slots, `oref` reads one, `make-instance` constructs.

Methods are not stored in the class. `cl-defgeneric` declares a function, and `cl-defmethod` adds an implementation *specialised on an argument’s type*:

```
(cl-defgeneric agent-backend-send (backend prompt))
(cl-defmethod agent-backend-send ((backend claude-client-backend) prompt) ...)
(cl-defmethod agent-backend-send ((backend opencode-client-backend) prompt) ...)
```

A call to `agent-backend-send` dispatches on the runtime class of its first argument. This is single dispatch used deliberately as a plugin seam: the caller — `agent-backend-explain-selection`, say — never learns which client it is talking to.

The interesting choice here is *which* generics get a default method. The five required verbs (`connect`, `quit`, `send`, `interrupt`, `add-note`) deliberately have none, so a subclass that forgets one fails with `cl-no-primary-method` naming the gap. Every optional capability has a default on the base class — a no-op, `nil`, or a `user-error` — so a minimal backend compiles without writing stubs.

That split is the whole design, stated in the commentary: *required verbs fail loudly, optional verbs degrade quietly*.

8.2 The optional verbs, and what their defaults say

The defaults are worth reading as documentation of intent (`elisp/agent-backend.el:148`):

Generic	Default	Why that default
<code>reply-permission</code>	<code>no-op</code>	A backend without a permission gate has nothing to reply to.
<code>reply-question</code>	<code>no-op</code>	Likewise.
<code>list-sessions</code>	<code>nil</code>	“I know of no sessions” is a true answer.
<code>resume</code>	<code>user-error</code>	Refusing is honest; silently doing nothing looks like success.
<code>seed-history</code>	<code>no-op</code>	Nothing stored means nothing to seed.
<code>project-root</code>	<code>nil</code>	
<code>render</code>	<code>no-op</code>	A backend may render by other means.
<code>mention</code>	delegates to <code>add-note</code>	See below.
<code>query</code>	<code>user-error</code> naming the backend	

Two of those are more than bookkeeping.

`agent-backend-mention` is “offer this text to the conversation *without* sending a turn” — how a selection reference reaches the agent so the human can keep typing around it. The default routes it through `add-note`, because notes already mean “the human said this, deliver it with the next turn”, which is exactly what a mention needs. So every backend gets a working mention for free, and one with a real editable input line can override to insert there instead. The commentary is explicit that this is deliberately *not* `agent-backend-send`: mentioning is not a turn.

`agent-backend-query` is “answer this once, outside any session”. The two backends satisfy it in incompatible ways — Claude shells out to `claude -p`, while `opencode`, which is session-only over HTTP, opens a session, asks, and deletes it. The default refuses *and names the backend*, so a backend that can do neither says so rather than pretending.

8.3 The event vocabulary

`agent-backend-event-functions` is an abnormal hook run with `(BUFFER EVENT)` for every backend event. `EVENT` is a plist whose `:kind` is one of `started`, `prompt`, `text`, `tool-use`, `tool-result`, `finished`, `interrupted`, `note`, `notes-delivered`, `note-dropped`, `resumed`, `permission-request`, `question-request`, or `error`.

The contract in the docstring is the load-bearing part:

Backends may emit kinds beyond this list; subscribers **MUST** ignore unknown kinds (guard their own pcase with a catch-all), so a publisher never has to know who is listening.

That is what makes the log extensible without coordinating releases. Chapter 25 picks this up as the foundation of the event model; Chapter 13 is the subscriber that proves it works.

Publishing goes through one function, and it is careful to do nothing else:

```
(defun agent-backend--publish (buffer event)
  (run-hook-with-args 'agent-backend-event-functions buffer event))
```

Rendering is *not* triggered here. A backend renders when it decides to; publishing only notifies. Chapter 23 lists this under “publish events, do not advise”.

8.4 Finding a conversation without a registry

Several commands need “the conversation the human means” — from a code buffer, with no argument. There is no session registry; there is a buffer-list scan for a buffer-local variable (elisp/agent-backend.el:286):

```
(defun agent-backend--conversation-buffers ()
  (seq-filter (lambda (buffer)
                (buffer-local-value 'agent-backend--instance buffer))
              (buffer-list)))
```

Any client that sets agent-backend--instance is found for free — the same no-registry choice agent-session-overview--sessions makes, and for the same reason: a registry is a second copy of the truth that has to be kept in step.

agent-backend--resolve-conversation then ranks candidates in tiers and takes the first non-empty one:

1. same project as the current buffer, and visible in a window
2. same project, hidden
3. any project, visible
4. any project, hidden

Visible beats hidden because a conversation on screen is the one being worked with; same-project beats other projects because that is what the selection is about. Several candidates in the winning tier means asking.

There is a deliberate second function, agent-backend--visible-conversation, which never signals and never prompts. The distinction matters: a caller with a fallback needs to *ask* whether a conversation is on screen, not to be stopped when none is.

8.5 Sharing what the human is looking at

agent-backend-selection-reference returns a reference to the selection for embedding in a prompt. In a file-visiting buffer that is an at-mention with a line span — @elisp/foo.el:12-40 — rather than the text itself. The reasoning is worth quoting: it is *a pointer the agent can read for itself, which beats pasting the text and stays right if the file moves on*.

There is one off-by-one guard worth learning once, because it recurs:

```
;; A region ending at column 0 covers up to the previous line.
(end-line (line-number-at-pos (if (and (use-region-p) (> end beg)
                                     (save-excursion
                                       (goto-char end) (bolp)))
                                (1- end)
                                end)))
```

Select three whole lines and point lands at the beginning of the fourth. Reporting the line point sits on would claim a line that was never selected.

agent-backend-explain-selection shows how the routing customs are meant to be read. agent-backend-explain-route has three values, and the docstring argues for each rather than just listing them: session-first (the default, because “explain this” is usually part of the work already in context), one-shot (for when asides pollute the conversation and push real work out of the model’s context), and session-only (for never spending a separate call).

8.6 agent-prompt.el — a buffer, not a minibuffer

Prompts used to be read with read-string. The commentary is blunt about why that was wrong: no multiline, no editing commands worth the name, no pasting a code block, *and a minibuffer that loses its place the moment an apply_diff ediff rearranges the windows.*

agent-prompt-read replaces it with the git-commit / org-capture shape: an ordinary buffer, C-c C-c to send, C-c C-k to abort, M-p/M-n for history. Under evil, ZZ and ZQ do the same and the buffer opens in insert state — chosen because *the buffer exists because someone is about to type into it.*

It is a **minor** mode, not a derived major mode, so the buffer keeps whatever editing mode suits the text (markdown, in practice, for the font-lock on code fences) and only gains the send and abort bindings.

The cost is stated up front: commands that used to read a string inline become asynchronous. Their interactive spec drops the read-string and the send moves into a callback. The non-interactive cores stay directly callable, so MCP tools and tests are unaffected.

Prompt history is session-only, and the docstring says why rather than leaving it to be discovered: *a prompt is a working note, not something worth carrying across restarts, and prompts routinely contain pasted source.*

8.6.1 Windows, and the timer

Placement follows the conversation, not the frame: the prompt window is a split of the agent’s own output window, so several conversations can each have their own prompt. Split below by default; split right only when the output window is too short to give up lines (agent-prompt-split-height-threshold, 20 lines — below that *a horizontal split leaves both halves useless*).

Closing deletes the window rather than restoring a saved configuration, and the comment says why in one line that will sound familiar from Chapter 6:

The split was taken from the conversation window, so removing it hands those lines straight back to the conversation. A saved configuration would also be stale by now if an ediff ran in between.

And the callback runs from a timer, not inline:

```
(agent-prompt--close)
(when callback
  (run-at-time 0 nil callback text))
```

The callback typically spawns a process and displays buffers. Doing that while `agent-prompt-send` is still unwinding the window it was invoked from is how conversation windows get lost. This is one of three places in the codebase using `run-at-time 0` for exactly this reason — see also `agent-permission--resolve` below, and Chapter 6.

8.7 `agent-permission.el` — a tool call as a question

The opening line of the commentary is the design:

A tool call the agent is not allowed to make should be a question, not a dead end.

The surface is a small read-only buffer naming the tool and the exact call, answered with one key: a/y allow once, d/n deny, A always allow. It is `agent-prompt`'s shape for the same reasons, with one difference — a prompt collects text and a decision collects a choice, so this derives from `special-mode` and single keys are commands.

Two invariants the callers depend on, both stated in the commentary:

- **Every pending decision is answered exactly once.** Whoever is waiting is blocked until then, so a decision that resolves twice would answer a call that already happened.
- **Not answering is a denial.** These requests arrive while nobody may be looking, so the timeout falls closed.

The first invariant is enforced structurally rather than with a flag, which is the trick worth stealing:

```
(defun agent-permission--take (id)
  "Remove and return the pending decision for ID, or nil when it is gone."
  (when-let* ((entry (assoc id agent-permission--pending)))
    (setq agent-permission--pending
          (assoc-delete-all (car entry) agent-permission--pending))
    (cdr entry)))
```

Lookup and removal are one step, so the registry *is* the answer to “is this still open”. A second answer finds nothing and does nothing. `agent-permission--resolve` returns non-nil only when its call was the one that answered, which is how the UI can say “that decision was already answered” instead of silently double-replying.

Note `assoc`, not `assq`: ids are strings, and an `assq` here would never match. The comment says so, because it is the kind of bug that presents as “the gate never fires”.

A long tool input is elided *in the middle*, on the grounds that the head says what the call is and the tail says what it ends with — and the omission is stated rather than implied. A decision is only meaningful if the human can see what they are approving, but a whole file passed to a write tool would bury the keys below the fold.

Elisp Feature: `special-mode` for single-key UIs

`special-mode` is the base major mode for read-only, non-editing buffers (`*Help*`, `direx-like` listings). Deriving from it gets you `buffer-read-only`, `q` to quit, and — crucially — a keymap where plain letters are free to be commands rather than self-inserting.

This mode sets `cursor-type` to `nil` as well: there is nowhere to type, so a blinking cursor would be a lie about what the buffer accepts.

Under `evil`, a read-only buffer still needs its keys re-registered in normal state, because `evil`'s state maps outrank a major mode's own keymap. Both this file and `agent-session-overview.el` do that in a `with-eval-after-load 'evil` block. Chapter 13 explains what breaks without it.

Note what `q` does here: it **denies**. The docstring is explicit that *there is no such thing as closing this buffer without answering it, because something is blocked waiting*.

8.7.1 The measured glob behaviour

`A` (always allow) writes a durable rule into `.claude/settings.local.json`. The rule syntax is Claude Code's, and the commentary records four *measured* behaviours rather than assumed ones:

```
Bash(echo hi)  matches `echo hi'          -- exact
Bash(echo)     does NOT match `echo hi'   -- a bare prefix
                                   silently misses
Bash(echo *)   matches `echo hi'          -- as expected
Bash(echo *)   ALSO matches `echo hi && rm -rf ...' -- the `*' runs
                                   past `&&' into
                                   another command
```

The last line is why the *exact* command is the default proposal. “Allow all echo commands” reads harmless and is not: it authorises `echo x && anything`. The prefix form stays available, because it is what makes an allowlist worth keeping, but the human has to choose it knowing that — hence the confirmation, which the docstring insists *is not ceremony*: the settings file is shared with the terminal CLI, so this changes what other sessions may do.

`A` is a capital letter on purpose, so it does not sit under the same finger as `a`.

8.7.2 Writing to a file someone else owns

`agent-permission-add-rule` is deliberately additive: read the existing settings, append one entry to `permissions.allow`, write the rest back untouched. Two JSON details in that path are guarded with comments, and both are the sort of thing that corrupts a shared file quietly:

```
;; A vector, not a list: an empty JSON array parses to nil,
;; and `json-encode' renders nil as `null'. Writing
;; "deny": null back into a file the CLI also reads is how a
;; well-meaning rewrite corrupts someone else's settings.
(allow* (vconcat allow (list rule)))
```

and, in the reader, `json-array-type` 'vector for the same reason: arrays this code never touches must survive the round-trip as arrays.

A malformed settings file **signals** rather than starting from scratch — silently beginning fresh would drop every rule the human already had.

And when the rule cannot be written, the call is still allowed:

```
;; The call was approved either way; only the durable part failed.
;; Allowing anyway is what the human just asked for.
```


8.7.3 The deadline

`agent-permission-request` takes a `:timeout` and reports expiry as the decision timeout, rather than leaving each caller to invent a default. The resolution runs from a timer for the now-familiar reason — the continuation writes to a process and may rearrange windows, and doing that while the command is still unwinding its own window is how windows get lost.

Raising a decision for an id that is already pending signals: two callers waiting on one id could not both be told the answer.

The file is backend-agnostic throughout. A pending decision is a plist and a continuation; *who* asked and how the answer travels back is the caller's business. Claude reaches this through a `PreToolUse` hook, `opencode` through its own permission API and `agent-backend-reply-permission` — and neither is mentioned in the file. Chapter 24 covers the hook wiring, including why the gate is emphatically not an MCP tool.

9 mcp-emacs-ide.el — The Claude Code IDE Surface

467 lines implementing an *unofficial, reverse-engineered* protocol. The commentary block is unusually specific about provenance, and for good reason.

9.1 Why this exists at all

The HTTP MCP server gives Claude Code a set of tools, including `apply_diff`. But Claude Code's *own* built-in Edit and Write tools do not go through MCP — they write files directly. They only route through an in-editor diff review when the CLI believes it is connected to an "IDE".

Becoming an IDE means: run a WebSocket server, write a discovery lockfile, and launch the CLI with `CLAUDE_CODE_SSE_PORT` and `ENABLE_IDE_INTEGRATION` in its environment. This module is that server.

9.2 What the spike found

The commentary records the protocol facts, verified live against Claude Code 2.1.212:

- `protocolVersion` is "2025-11-25"; `initialize` echoes it back.
- Before every native edit the CLI calls `closeAllDiffTabs`, then `getDiagnostics`, then `openDiff`, and **blocks until each is answered**. So all four tools must be served, not just `openDiff`.
- `openDiff` is deferred: no immediate response. When the human resolves the review, the stored request id is completed with `FILE_SAVED` plus final content (accept) or `DIFF_REJECTED` plus the tab name (reject).
- Claude Code writes the file itself on `FILE_SAVED`.

That last point is important: Emacs does not write the file. It returns the approved content and the CLI performs the write.

The whole surface is opt-in:

```
(defcustom mcp-emacs-ide-enabled nil
  "When non-nil, allow the IDE integration surface to start.
Off by default: the Claude Code IDE protocol is unofficial and
version-fragile, so it is opt-in and isolated from the HTTP MCP server."
  :type 'boolean :group 'mcp-emacs-ide)
```

9.3 Session state as a struct

```
(cl-defstruct (mcp-emacs-ide-session (:constructor mcp-emacs-ide--make-session))
  "State for one IDE WebSocket server and its connected client."
  server          ; websocket server process
  client          ; connected websocket client (or nil)
  port            ; server port
  project-dir     ; workspace root
  lockfile        ; path to the discovery lockfile
  (diffs (make-hash-table :test 'equal)) ; tab-name -> diff plist
  (deferred (make-hash-table :test 'equal)) ; tab-name -> pending request id
```

Elisp Feature: `cl-defstruct`

From `cl-lib`, this defines a record type and generates:

- a constructor, `make-NAME` by default — renamed here via `(:constructor mcp-emacs-ide--make-session)` to keep it in the private namespace;
- an accessor per slot, `NAME-SLOT`, e.g. `mcp-emacs-ide-session-client`;
- a predicate, `NAME-p`;
- setf-ability: `(setf (mcp-emacs-ide-session-client s) ws)` works.

Slots may have defaults, written as `(slot default)`. Note that the default expression is evaluated *per construction*, which is why `(diffs (make-hash-table :test 'equal))` gives each session its own table rather than sharing one.

Structs are used for the two places in this codebase with genuinely multi-field, mutable state: the IDE session here, and the orgspec model (Chapter 15). Everything smaller uses plists.

Two hash tables, both keyed by tab name:

- `diffs` — the live buffers and control buffer for a review in progress, so it can be torn down.
- `deferred` — the JSON-RPC request id awaiting completion.

They are separate because their lifetimes differ: the deferred id is removed the moment the response is sent, while the diff entry survives until cleanup.

9.4 The discovery lockfile

```
(defun mcp-emacs-ide--write-lockfile (port project-dir)
  (let ((path (mcp-emacs-ide--lockfile-path port)))
    (make-directory (file-name-directory path) t)
    (with-temp-file path
      (insert (json-encode
                (mcp-emacs-ide--obj
                 "pid" (emacs-pid)
                 "workspaceFolders" (vector (expand-file-name project-dir))
                 "ideName" "Emacs"
                 "transport" "ws")))))
    path))
```

Written to `~/.claude/ide/<port>.lock`. It is removed on stop and, best-effort, on `kill-emacs-hook`.

Elisp Feature: with-temp-file and with-temp-buffer

`with-temp-buffer` creates a scratch buffer, makes it current for the body, and kills it afterwards. Used constantly here for “build a string by manipulating it as a buffer” — see `orgspec-fold-area`, which does whole Org subtree surgery in a temp buffer and returns `(buffer-string)`.

`with-temp-file` `FILE` is the same thing plus a write: the body fills the temp buffer, and its contents are written to `FILE` on exit. It replaces the file atomically enough for these purposes and saves a `write-region` call.

9.5 Binding a port

```
(defun mcp-emacs-ide--start-server ()
  (let ((min (car mcp-emacs-ide--port-range))
        (max (cdr mcp-emacs-ide--port-range))
        (attempts 0)
        found)
    (while (and (null found) (< attempts mcp-emacs-ide--max-port-attempts))
      (let ((port (+ min (random (- max min)))))
        (condition-case nil
          (let ((server (websocket-server port :host "127.0.0.1" ...)))
            (setq found (cons server port)))
          (error (setq attempts (1+ attempts))))))
      (or found (error "mcp-emacs-ide: no free port in range %d-%d" min max))))
```

Random port in [10000, 65535], retry up to 100 times on bind failure, bind to loopback only. There is no port registry to consult, so try-and-see is the available strategy.

The server is created with three callbacks:

```
:on-open      (lambda (ws) (setf (mcp-emacs-ide-session-client ...) ws))
:on-message   (lambda (ws frame)
                (mcp-emacs-ide--handle-message
                 mcp-emacs-ide--session ws (websocket-frame-text frame)))
:on-close     (lambda (_ws) (setf (mcp-emacs-ide-session-client ...) nil))
```

Each guards on `mcp-emacs-ide--session` being non-nil, so a late frame after a stop is ignored rather than erroring.

9.6 Message dispatch and the deferred sentinel

`mcp-emacs-ide--handle-message` handles `initialize`, `tools/list`, `tools/call`, `notifications`, and `unknown` methods. The `tools/call` branch has the interesting shape:

```
(let* ((res (condition-case err
              (mcp-emacs-ide--call-tool session name args id)
              (error (mcp-emacs-ide--result
                     id (mcp-emacs-ide--text-content
                        (format "Error: %s" (error-message-string err)))))))
      (unless (eq res 'deferred)
        (mcp-emacs-ide--send client res)))
```

`mcp-emacs-ide--call-tool` returns either a response object *or* the symbol `deferred`, which means “I have stored the id and will answer later”. The sentinel-value protocol is simple and works because a response object is always a cons, never a symbol.

`getDiagnostics` is honestly stubbed:

```
((string= name "getDiagnostics")
;; Stubbed: Claude Code blocks until this is answered, but an empty
;; result is acceptable. A future version can wire Flycheck/Flymake.
(mcp-emacs-ide--result id (mcp-emacs-ide--text-content "[]")))
```

It must be answered — the CLI blocks — but an empty array satisfies it. The comment marks the shortcut rather than hiding it.

9.7 The deferred *openDiff*

```
(puthash tab-name id (mcp-emacs-ide-session-deferred session))
(let ((control
      (mcp-emacs--ediff-review
       buffer-a buffer-b entry-content result
       (lambda ()
         (if (eq (car result) 'applied)
             (mcp-emacs-ide--complete-open-diff
              session tab-name "FILE_SAVED"
              (with-current-buffer buffer-a
                (buffer-substring-no-properties (point-min) (point-max))))
             (mcp-emacs-ide--complete-open-diff
              session tab-name "DIFF_REJECTED" tab-name))
         (mcp-emacs-ide--cleanup-diff session tab-name))
       tab-name)))
      (puthash tab-name (list :buffer-a buffer-a :buffer-b buffer-b
                             :control control :result result)
               (mcp-emacs-ide-session-diffs session)))
'deferred
```

The comment on the ordering is the kind of detail that only shows up after a bug: “Register the deferred id before opening ediff so a very fast resolve still finds it.” If the human somehow resolved before `puthash` ran, the completion would find no id and silently drop the response.

Completion itself is defensive:

```
(defun mcp-emacs-ide--complete-open-diff (session tab-name status &rest extra)
  (let* ((deferred (mcp-emacs-ide-session-deferred session))
         (id (gethash tab-name deferred)))
    (when id
      (remhash tab-name deferred)
      (mcp-emacs-ide--send
       (mcp-emacs-ide-session-client session)
       (mcp-emacs-ide--result
        id (apply #'mcp-emacs-ide--text-content status extra))))))
```

(when id ...) plus immediate `remhash` makes double-completion a no-op.

This is the same `mcp-emacs--ediff-review` the `apply_diff` tool uses, with a different

on-resolve. One review implementation, two protocols.

10 mcp-emacs-run.el — The Terminal Runner

678 lines, and unlike the previous chapters, almost none of it is protocol. This is window management, session bookkeeping, and terminal I/O.

10.1 What it does

Launches the Claude Code CLI — a full-screen ANSI TUI — inside an eat terminal buffer, one or more sessions per project, with commands to show, hide, switch, kill, and send input to them.

eat is a soft dependency; markdown-mode is a second soft dependency used only by the popup output window.

10.2 Sessions without a registry

Buffers are named `*claude:<project>:<n>*`, and that name *is* the registry:

```
(defconst mcp-emacs-run--buffer-name-regexp
  "\\`\\*claude:\\`\\(\\.+\\):\\`\\([0-9]+\\)\\`\\*\\`"
  "Regex matching a runner buffer name `*claude:<project>:<n>*'.")

(defun mcp-emacs-run--sessions-list ()
  (seq-filter (lambda (buf)
                (string-match-p mcp-emacs-run--buffer-name-regexp
                                (buffer-name buf)))
              (buffer-list)))
```

A session's project is not stored either — it is recomputed from the buffer's own default-directory:

```
(defun mcp-emacs-run--buffer-project (buf)
  "Return the project root that runner buffer BUF belongs to.
Derived from BUF's own `default-directory' via `mcp-emacs-run--project-root',
so it survives without any registry."
  (with-current-buffer buf
    (mcp-emacs-run--project-root)))
```

The payoff: no state to get out of sync. Kill a buffer and the session is gone; no cleanup hook, no stale entry. The cost is a buffer-list scan per query, which at this scale is free.

Elisp Feature: regexp syntax in Elisp strings

`"\\`*claude:\\`\\(\\.+\\):\\`\\([0-9]+\\)\\`\\*\\`"` looks alarming until you separate the two levels of escaping.

Elisp strings use `\\` for a literal backslash, so every regexp backslash is doubled. Strip that and the regexp is:

```
\\`*claude:\\(\\.\\):\\([0-9]+\\)\\`*'
```

- `\\`` — buffer/string start (Elisp's `\\A`; plain `^` matches line starts).
- `\\`` — buffer/string end.
- `\\(\\.\\.\\)` — a *group*. In Emacs regexps, bare parentheses are literal and backslashed ones group; this is the opposite of PCRE.
- `\\`*` — a literal asterisk.

Group 1 is the project name, group 2 the session number, retrieved with `(match-string 1 name)` after a successful string-match.

Session numbers refill gaps:

```
(defun mcp-emacs-run--next-number (root)
  "Return the lowest positive session number not in use for project ROOT."
  (let ((used ...) (n 1))
    (while (memq n used) (setq n (1+ n)))
    n))
```

10.3 Resolving which session gets input

When the user runs a send command, which of possibly several sessions should receive it? `mcp-emacs-run--resolve-session` (line 235) ranks candidates into four tiers and takes the first non-empty one:

```
(let* ((visible (lambda (buf) (get-buffer-window buf t)))
      (same (seq-filter (lambda (b) (string= (mcp-emacs-run--buffer-project b) root))
                        all))
      (tiers (list (seq-filter visible same)      ; 1. same project, visible
                   (seq-remove visible same)     ; 2. same project, hidden
                   (seq-filter visible all)       ; 3. any project, visible
                   (seq-remove visible all)))     ; 4. any project, hidden
      (tier (seq-find #'identity tiers)))
  (mcp-emacs-run--pick-session tier))
```

`(seq-find #'identity tiers)` finds the first non-nil element — here, the first non-empty list. Only if the winning tier has more than one candidate does it prompt.

10.4 Window placement: an ordinary window, deliberately

```
(defun mcp-emacs-run--display (buffer)
  (let* ((horizontal (memq mcp-emacs-run-window-direction '(left right)))
        (size (if horizontal
                    `(window-width . ,(mcp-emacs-run--resolved-width))
                    `(window-height . ,mcp-emacs-run-window-height)))
        (window (display-buffer
                  buffer
                  `((display-buffer-in-direction)
                    (direction . ,mcp-emacs-run-window-direction))
```



```

        ,size))))
(when window (set-window-dedicated-p window t))
(when (and window mcp-emacs-run-focus-on-show) (select-window window))
window))

```

The docstring explains the choice: “The runner uses an ordinary window in this direction, so it can be split, navigated, and closed like any other window.” That is a reaction against side windows, which cannot be split and behave irregularly — including breaking ediff (§6.3).

(`set-window-dedicated-p window t`) marks it *weakly* dedicated: `display-buffer` prefers not to reuse it for another buffer, but may if there is no alternative. Passing 'strong instead would make it refuse.

Elisp Feature: backquote and unquote

```

`((display-buffer-in-direction)
  (direction . ,mcp-emacs-run-window-direction)
  ,size)

```

The backquote ``` starts a template: everything is quoted literally *except* forms preceded by `,` (unquote), which are evaluated and spliced in. `,@` (unquote-splicing) splices a list’s elements rather than the list itself.

This is how Elisp builds data structures with holes. It appears throughout this codebase for `display-buffer` action alists, JSON bodies in `claude-client.el`, and the agenda command entry in `orgspec-agenda.el`.

The width computation is a small piece of care:

```

(defun mcp-emacs-run--resolved-width ()
  (let ((cap (truncate (* mcp-emacs-run-window-max-width-fraction (frame-width))))
        (cols (or mcp-emacs-run-window-width-columns
                   (truncate (* mcp-emacs-run-window-width (frame-width))))))
    (min cols cap)))

```

Prefer an absolute column count (120 by default — a sane width for the CLI’s output), but never let it exceed half the frame, so the code pane is not crushed on a laptop screen.

10.5 Beating the framework at `display-buffer`

The popup output window has a harder problem: Doom Emacs’s `+popup` module installs rules in `display-buffer-alist` that would turn this window into a transient, auto-hiding thing. The fix (line 414):

```

(defun mcp-emacs-run--display-popup (buffer)
  (let* ((rule `((, (regexp-quote (buffer-name buffer))
                  (display-buffer-reuse-window display-buffer-in-direction)
                  (direction . ,mcp-emacs-run-popup-direction)
                  ,size))
         (display-buffer-alist (cons rule display-buffer-alist)))
    (display-buffer buffer)))

```

A rule matching this exact buffer name is *prepended* to a let-bound copy of `display-buffer-alist`. `display-buffer` consults the alist in order and takes the first match, so this rule wins.

And because the binding is dynamic and scoped to the `let`, the user's global configuration is untouched.

This is a genuinely useful pattern for any package that must place a window predictably inside someone else's framework.

10.6 Passing the IDE port through the environment

```
(let* ((ide-port (mcp-emacs-run--ide-port))
      (process-environment
       (if ide-port
           (append (list (format "CLAUDE_CODE_SSE_PORT=%d" ide-port)
                         "ENABLE_IDE_INTEGRATION=true"
                         "TERM_PROGRAM=emacs")
                   process-environment)
           process-environment))
      (buffer (apply #'eat-make name mcp-emacs-run-executable nil switches)))
  ...)
```

`process-environment` is a list of "NAME=VALUE" strings that any process started within the binding inherits. Prepending wins, since the list is searched front to back.

The port resolution degrades gracefully:

```
(defun mcp-emacs-run--ide-port ()
  (when mcp-emacs-run-ide-integration
    (condition-case err
      (when (require 'mcp-emacs-ide nil t)
        (or (mcp-emacs-ide-port) (mcp-emacs-ide-start)))
      (error (message "mcp-emacs-run: IDE integration unavailable: %s"
                     (error-message-string err))
             nil))))
```

Missing websocket package, disabled surface, port exhaustion — all produce `nil` and a message, and the session still launches without IDE integration. An optional feature failing should not stop the main thing from working.

10.7 The eat-term-send-string trap

This one cost someone real debugging time, and the docstring says so:

```
(defun mcp-emacs-run--send-to-buffer (buf string)
  "Send STRING to runner buffer BUF's terminal.
`eat-term-send-string' resolves the input process from the current
buffer's eat state, not from its TERM argument, so this must run with
BUF current or it silently sends nothing when invoked from another
buffer (e.g. via \\[execute-extended-command] from a file buffer)."
  (with-current-buffer buf
    (let ((term (buffer-local-value 'eat-terminal buf)))
      (unless term (user-error "Runner session is not a live terminal"))
      (eat-term-send-string term string))))
```

The API *looks* like it takes the terminal as an argument. It does, but it resolves the actual input process from buffer-local state, so calling it with the wrong current buffer sends nothing — silently. The wrapper makes the buffer current unconditionally.

Elisp Feature: `buffer-local-value`

```
(buffer-local-value 'eat-terminal buf)
```

Reads a variable's value *as seen in another buffer* without making that buffer current. It is the read-only counterpart to `(with-current-buffer buf var)`, and it is cheaper and clearer when all you need is one value.

Note the interesting detail in the code above: it uses `buffer-local-value` *inside* `with-current-buffer`, which is redundant for the read — but the `with-current-buffer` is there for `eat-term-send-string`'s benefit, not the read's.

10.8 Quitting with a deadline

```
(defconst mcp-emacs-run--quit-sequence "\003\003"
  "Sequence that makes the Claude CLI exit: two Ctrl-C characters.")

(defun mcp-emacs-run-quit ()
  (interactive)
  (let ((buf (mcp-emacs-run--resolve-session)))
    (mcp-emacs-run--send-to-buffer buf mcp-emacs-run--quit-sequence)
    (run-with-timer mcp-emacs-run-quit-timeout nil
      #'mcp-emacs-run--force-kill-buffer buf)
    (message "Quitting Claude runner session %s" (buffer-name buf))))
```

Ask nicely, then set a 10-second timer to force-kill. The command returns immediately; Emacs is never blocked waiting for a CLI to notice a signal. `"\003"` is the string containing character 3, i.e. C-c.

10.9 Selection references

```
(defun mcp-emacs-run--selection-reference ()
  (let* ((beg (if (use-region-p) (region-beginning) (point)))
        (end (if (use-region-p) (region-end) (point)))
        (file (buffer-file-name)))
    (if file
      (let* ((root (mcp-emacs-run--project-root))
            (rel (file-relative-name file root))
            (start-line (line-number-at-pos beg))
            (end-line (line-number-at-pos
                      (if (and (use-region-p) (> end beg)
                            (save-excursion (goto-char end) (bolp)))
                        (1- end)
                        end))))
        (if (and (use-region-p) (/= start-line end-line))
          (format "@%s:%d-%d" rel start-line end-line)
          (format "@%s:%d" rel start-line end-line)))
      (format "%s" file)))
```

```
(format "@%s:%d" rel start-line)))
...verbatim text...)))
```

For a file buffer it produces `@src/foo.el:10-25`, the at-mention syntax the CLI understands, so the agent resolves the reference itself rather than being handed a wall of text.

The `bolp` check handles an off-by-one that bites everyone: selecting three whole lines with a line-oriented command leaves the region ending at column 0 of the *fourth* line. Without the adjustment the reference would claim one line too many.

10.10 Headless one-shot queries

```
(defun mcp-emacs-run--query-headless (prompt callback)
  (let* ((default-directory (file-name-as-directory (mcp-emacs-run--project-root)))
        (out (generate-new-buffer " *mcp-emacs-query-out*"))
        (err (generate-new-buffer " *mcp-emacs-query-err*")))
    (make-process
     :name "mcp-emacs-query"
     :buffer out :stderr err :noquery t
     :command (list mcp-emacs-run-executable "-p" prompt "--output-format" "text")
     :sentinel (lambda (proc _event) ...))))
```

Details: buffer names starting with a space are hidden from the buffer list; `:noquery t` suppresses the kill-on-exit prompt; `default-directory` is bound so the CLI runs in the project root.

Elisp Feature: process sentinels

A *sentinel* is the state-change counterpart to a filter — called when a process exits, is signalled, or otherwise changes status:

```
:sentinel (lambda (proc event)
  (when (memq (process-status proc) '(exit signal))
    (let ((code (process-exit-status proc)))
      ...)))
```

`event` is a human-readable string like `"finished\n"`; the reliable check is `process-status` plus `process-exit-status`. Sentinels run from the event loop, so unlike filters they have a normal execution context.

The sentinel here uses `unwind-protect` around the callback so both scratch buffers are killed even if the callback signals.

`mcp-emacs-explain-selection-in-current-session` (line 658) ties it together with a routing decision: if a runner window for this project is *visible*, send the prompt into the live session; otherwise fire a headless query and render the answer in a markdown popup. The user gets an answer either way, without a TUI appearing unbidden.

11 `claude-client.el` — The Terminal-Free Runner

1648 lines, and the largest file in the project. This is the runner that removes the terminal emulator entirely, and it has grown the most since it was written: the event model, the note channel, the permission gate wiring, and the `agent-backend` methods all landed here first and were generalised outwards afterwards.

11.1 The shape

Spawn the CLI as a plain subprocess with `--output-format stream-json`, frame its stdout as NDJSON, turn each message into an *event*, append the event to a per-buffer log, notify subscribers, and re-render.

```
subprocess stdout → filter (frame lines) → parse JSON
  → handle-message (build events) → push-event
    → append to log
    → run-hook-with-args (subscribers)
    → render
```

Compare with `mcp-emacs-run.el`, which pipes the CLI's ANSI output into a terminal emulator and lets the human read pixels. Here the output is *structured*, so Emacs can do something with it.

11.2 Two differences from the `opcode` backend, stated up front

The commentary names them:

- **Transport is a subprocess, not HTTP+SSE.** `--output-format stream-json` is NDJSON: one complete JSON object per line. Frame on newlines, not on SSE's blank-line separator.
- **Edits are gated at the MCP tool boundary.** Claude's own mutating tools are disabled with `--disallowedTools`, and the `mcp-emacs` server is passed via `--mcp-config`, so a write can only happen by calling `mcp__emacs__apply_diff` — which opens an ediff the human answers.

And a negative finding, which is the most valuable line in the file:

The CLI's `can_use_tool` stdio control request is NOT used: it is never emitted in headless mode (verified against 2.1.220), and the `permission_denials` in the result are a post-hoc audit record rather than a gate.

Someone tried the obvious approach, found it does not work, and wrote down why. That saves the next person a day.

11.3 Constructing the command line

```
(defun claude-client--command (&optional resume-id)
  (append
    (list claude-client-executable
          "--print"
          "--output-format" "stream-json"
          "--input-format" "stream-json"
          "--verbose"
          "--mcp-config" (claude-client--mcp-config-file)
          "--strict-mcp-config"
          "--settings" (claude-client--settings-file)
          "--append-system-prompt-file" (claude-client--system-prompt-file))
    (when resume-id (list "--resume" resume-id))
    (when claude-client-model (list "--model" claude-client-model))
    (when claude-client-disallowed-tools
      (cons "--disallowedTools" claude-client-disallowed-tools))))
```

`--resume` carries a flag-naming trap the docstring records: it is *not* `--session-id`. `--resume ID` reopens that past session with its history, keeping the id; `--session-id ID` means “create a new session with this id” and fails outright if a transcript already exists.

Three temp files are generated per run:

```
(defun claude-client--mcp-config-file ()
  (or claude-client-mcp-config
    (let ((file (make-temp-file "claude-client-mcp-" nil ".json")))
      (url (format "http://localhost:%s/mcp"
                  (if (boundp 'mcp-emacs-server-port)
                      (symbol-value 'mcp-emacs-server-port)
                      8765))))
    (with-temp-file file
      (insert (json-encode
                `((mcpServers . ((emacs . ((type . "http")
                                             (url . ,url))))))))
      file)))
```

Note `(if (boundp 'mcp-emacs-server-port) (symbol-value ...) 8765)` — the client works whether or not the server module happens to be loaded. `symbol-value` reads a variable whose name is computed at runtime, which is required here because the compiler cannot know the symbol is bound.

The security model is expressed in two defcustoms:

```
(defcustom claude-client-disallowed-tools
  '("Write" "Edit" "MultiEdit" "NotebookEdit")
  "Claude built-in tools to disable so edits route through mcp-emacs.
`MultiEdit' must be disabled alongside `Edit': otherwise Claude can
batch file changes through it and bypass the ediff gate entirely."
  ...)

(defcustom claude-client-allowed-mcp-tools
  '("mcp__emacs__apply_diff" "mcp__emacs__open_file" ...)
  "MCP tools Claude may call without an interactive permission prompt."
```

Proxied MCP tools are denied by default, so the write path `\('apply_diff')` has to be listed here or every edit is refused before the human ever sees the ediff."

```
...)
```

Both docstrings record a specific way the gate can be defeated or accidentally disabled. If you extend this, read them before changing either list.

11.3.1 The settings file layers, it does not replace

```
(defun claude-client--settings-file ()
  (let ((file (make-temp-file "claude-client-settings-" nil ".json")))
    (hooks (claude-client--gate-hooks)))
    (with-temp-file file
      (insert (json-encode
                (append
                  `((permissions
                     . ((allow . ,(vconcat claude-client-allowed-mcp-tools))
                       (deny . []))))
                  (when hooks `((hooks . ,hooks))))))
              file))
```

The docstring makes a promise that is easy to get backwards: `--settings` *adds to* `.claude/settings.json` and `settings.local.json` rather than overriding them, so the project's existing permission rules still apply to the session. This is also why agent-permission's "always allow" (Chapter 8) has any lasting effect at all — it writes to the project file, which this layers on top of.

Note `(deny . [])`: an empty vector, for the reason Chapter 8 gives about `json-encode` rendering `nil` as `null`.

11.3.2 Installing the PreToolUse gate

`claude-client--gate-hooks` builds the hook entry, one matcher per entry in `claude-client-gate-tools`:

```
`((matcher . ,matcher)
  (hooks . [((type . "command")
              (command . ,(format "MCP Emacs port=%d %s"
                                   port
                                   (shell-quote-argument
                                    claude-client-gate-script))))))
```

Three decisions are recorded here, and each is a small lesson:

The default is `("Bash")`, deliberately not `"*"`. The mutating built-ins are already disabled outright, so gating them would raise a menu for a call that cannot happen; gating the read and search tools would raise one on nearly every call. The docstring states the real criterion — *a shell command is the case where the tool name alone does not tell you whether to allow it* — and then the operational one: **a gate that is tedious gets switched off.**

The port travels in the hook's environment, not by discovery. This Emacs knows which port it is serving on; a script left guessing would find the wrong instance whenever a test daemon or a worktree is also up. (Chapter 20 explains why a second Emacs on port 8775 is routine here.)

A missing script switches the gate off rather than breaking every call. `claude-client--gate-hooks` returns `nil` when the script does not exist, so a misconfigured path degrades to the previous behaviour instead of wiring every tool call to a hook that cannot run. Note the asymmetry with the *script's* own posture: once installed, the gate fails closed. Failing open is only acceptable at the point where the gate was never installed at all.

The script path itself is resolved from the file's own location, so a checkout or a straight build finds its own copy rather than whatever is on `PATH`.

11.4 The event log and the subscriber hook

```
(defvar-local claude-client--events nil
  "Parsed conversation events, oldest first.
  Each is a plist: :kind, plus kind-specific keys. This is the render
  model, and the append-only log issue #39 wants to subscribe to.")

(defun claude-client--push-event (buffer event)
  "Append EVENT to BUFFER's log, notify subscribers, and re-render."
  (when (buffer-live-p buffer)
    (with-current-buffer buffer
      (setq claude-client--events (append claude-client--events (list event)))
      (agent-backend--publish buffer event)
      (claude-client--render)))))
```

Three lines, three responsibilities, in a deliberate order: the log is the source of truth, subscribers see every event, and rendering is *just another consumer* that happens to be built in.

The log stays buffer-local to this client, but the *hook* does not: publishing goes through `agent-backend--publish`, and the event vocabulary is defined once on the shared layer (Chapter 8). This file is where the pattern was invented; it was moved outward when the second backend needed the same thing.

Elisp Feature: hooks, normal and abnormal

A **normal hook** is a variable holding a list of functions called with no arguments, by convention named `*-hook`:

```
(add-hook 'after-save-hook #'my-function)
(run-hooks 'after-save-hook)
```

An **abnormal hook** takes arguments and/or interprets return values, and by convention is named `*-functions`:

```
(add-hook 'agent-backend-event-functions #'my-subscriber)
(run-hook-with-args 'agent-backend-event-functions buffer event)
```

The naming convention matters: `-functions` warns the reader “check the docstring for the calling convention”.

There is also `defvar-local`, used for the log itself, which is shorthand for `defvar` plus `make-variable-buffer-local` — every buffer gets its own value automatically.

`mcp-emacs-remote.el` subscribes to exactly this hook (§14.3), and its comment explains why the design matters:

The IDE taps above are advice: the transcript only learns what the IDE surface hap-

pens to be asked. `claude-client` instead *publishes* an append-only event log and announces each event on a hook, so this is a plain subscriber — no advice, and rendering stays one consumer among several.

Publishing beats advising. Adding a reader costs the writer nothing.

11.5 Framing NDJSON

```
(defun claude-client--filter (buffer proc chunk)
  "Frame CHUNK from PROC as NDJSON lines and dispatch them to BUFFER."
  (ignore proc)
  (when (buffer-live-p buffer)
    (with-current-buffer buffer
      (setq claude-client--stdout (concat claude-client--stdout chunk))
      (let (pos)
        (while (setq pos (string-search "\n" claude-client--stdout))
          (let ((line (substring claude-client--stdout 0 pos)))
            (setq claude-client--stdout (substring claude-client--stdout (1+ pos)))
            (unless (string-empty-p (string-trim line))
              (let ((msg (ignore-errors
                           (json-parse-string line
                             :object-type 'alist
                             :array-type 'list
                             :null-object nil
                             :false-object nil))))
                (when msg (claude-client--handle-message buffer msg))))))))))
```

This is the canonical Emacs stream-framing pattern, and it is worth internalising:

1. Append the chunk to a buffer-local accumulator.
2. Loop while a complete delimiter exists in the accumulator.
3. Split off the complete unit; keep the remainder.
4. Parse and dispatch the unit.

A process filter receives *arbitrary* chunks. A single JSON object may arrive split across three calls; three objects may arrive in one call. Only consuming up to the last delimiter and keeping the tail handles both.

Elisp Feature: `json-parse-string` vs `json-read-from-string`

Emacs 27 added `json-parse-string`, a C implementation, much faster than the Lisp `json.el` reader. It takes keyword arguments instead of dynamic variables:

```
(json-parse-string str
  :object-type 'alist ; or 'hash-table (default), 'plist
  :array-type 'list ; or 'array (default)
  :null-object nil ; default is :null
  :false-object nil) ; default is :false
```

The last two matter enormously in practice. By default JSON `null` becomes the symbol `:null` and `false` becomes `:false`, both of which are **truthy** in Emacs. Mapping both to `nil` makes the parsed data behave the way Emacs code naturally expects.

Note the codebase uses both readers: `json-parse-string` in the newer files

(`claude-client`, `opencode-client`, `mcp-emacs-ide`) and the dynamic-variable `json-read-from-string` in `mcp-emacs-server` and `mcp-emacs-run-resume`. If you touch a parse site, match the local style rather than mixing.

11.6 Turning messages into events

```
(defun claude-client--handle-message (buffer msg)
  (let ((type (alist-get 'type msg)))
    (cond
      ((and (equal type "system") (equal (alist-get 'subtype msg) "init"))
        ...push (:kind 'started :session ... :model ...))
      ((equal type "assistant")
        (dolist (event (claude-client--events-for-message (alist-get 'message msg)))
          (claude-client--push-event buffer event)))
      ((equal type "user")
        ;; Tool results come back as a synthetic user message.
        (dolist (block (alist-get 'content (alist-get 'message msg)))
          (when (equal (alist-get 'type block) "tool_result")
            ...push (:kind 'tool-result :text ...))))
      ((equal type "result")
        ...push (:kind 'finished :subtype ... :denials ...))
      (t nil))))
```

Five external message shapes collapse into five internal event kinds: `started`, `text`, `tool-use`, `tool-result`, `finished`. The trailing `(t nil)` means an unrecognised message type is silently ignored — forward compatibility with a CLI that adds message types.

Rendering is a `pcase` over `:kind`:

```
(defun claude-client--render-event (event)
  (pcase (plist-get event :kind)
    ('started (format "— claude %s —" (or (plist-get event :model) "")))
    ('text (plist-get event :text))
    ('tool-use (format " [tool: %s]" (plist-get event :name)))
    ('tool-result ...)
    ('finished (format "— %s —" (or (plist-get event :subtype) "done")))
    (_ nil)))
```

Elisp Feature: `pcase`

Elisp’s pattern-matching `cond`. The patterns used in this codebase:

- `'symbol` — matches that exact symbol (quoted, i.e. `(quote symbol)` as a pattern means “equal to this constant”).
- `"string"` — matches an equal string.
- `_` — wildcard, matches anything; conventionally the final clause.
- `(pred functionp)` — matches when the predicate returns non-`nil`.
- ``(a ,x)` — backquote patterns destructure and bind.

`pcase` is the idiomatic dispatcher for small tagged unions like this event model, and reads far better than a `cond` chain of `eq` tests.

11.7 Full re-render on every event

```
(defun claude-client--render ()
  (let ((at-end (eobp))
        (inhibit-read-only t))
    (erase-buffer)
    (dolist (event claude-client--events)
      (when-let ((s (claude-client--render-event event)))
        (insert s)
        (unless (string-suffix-p "\n" s) (insert "\n"))))
      (when at-end (goto-char (point-max))))))
```

Erase and rebuild the whole buffer, every time. That is $O(n^2)$ over a conversation, and for a conversation of a few hundred events it is imperceptible while being trivially correct — no incremental-update bugs, no stale regions.

The `at-end` capture implements sticky scrolling: if the user was at the bottom, stay at the bottom; if they had scrolled up to read something, do not yank them away. `eobp` is “end of buffer predicate”.

`opcode-client.el` uses the identical pattern (line 406), which is a good sign that it was a deliberate choice rather than an accident.

11.8 Turns over one process

The first version spawned a fresh CLI per prompt. It now keeps one subprocess and writes further turns to it, so the session id and the model’s context carry across them:

```
(defun claude-client-send (prompt)
  "Send PROMPT as the next turn of this conversation."
  (cond
   (claude-client--turn-active
    (user-error "A turn is already running; add a note with `n' instead"))
   ((not (process-live-p claude-client--process))
    (user-error "No live Claude process; start one with `g'"))
   (t ...)))
```

The refusal while a turn is in flight is not defensiveness: the CLI reads one turn at a time, so a second write would interleave with the one being answered.

The subtle part is knowing whether a turn *is* in flight. The obvious test — is the process alive? — is wrong, and wrong in a way that passed every unit test:

```
(defvar-local claude-client--turn-active nil
  "Non-nil while a turn is in flight, i.e. between spawn and `result'.
Tracked from the stream rather than from the process: `claude --print'
stays alive after emitting `result' (it waits on stdin for a follow-up
turn), so `process-live-p' reports `run' long after the turn is over
and cannot answer \"is the model working right now?\".")
```

`claude --print` does not exit after answering; it waits on stdin. Reading process liveness as “the model is working” left notes queued forever once a run had finished. Only a live run surfaced it.

11.9 Reopening a past session

```
(when resume-id (list "--resume" resume-id))
```

`--resume`, not `--session-id`. The two look interchangeable and are not: `--session-id` means *create a new session with this id* and fails outright when a transcript already exists. The picker reuses the on-disk store and the labels `mcp-emacs-run-resume.el` already builds, so a session started in either runner continues in the other.

Only the model's context comes back. The rendered log lives in the buffer, not on disk, so the conversation restarts on screen while the model still remembers — an asymmetry worth knowing before it surprises you.

11.10 Notes: the human as a second producer

Until this point the log had one producer and several readers. A note is the human writing to the same log:

```
(defun claude-client-add-note (text) ...) ; bound to `n'
```

A note is recorded the moment it is written, so every subscriber sees it at once and the transcript interleaves human `::` and assistant `::` lines in the order things actually happened. What happens *next* is the interesting part.

11.11 Interruption: the design question, answered by measurement

Issue #39 calls this *the* question — a note arrives mid tool-call: finish, abort, or re-plan? The CLI turns out to answer it for us. Its init event advertises:

```
["interrupt_receipt_v1", "interrupt_cancel_queued_v1", "msg_lifecycle_v1"]
```

and accepts a control request that abandons the turn in flight:

```
(defun claude-client--send-interrupt (proc)
  (process-send-string
   proc
   (concat (json-encode
            `((type . "control_request")
              (request_id . ,(format "int-%s" (float-time)))
              (request . ((subtype . "interrupt"))))
            "\n"))))
```

Two properties were checked against the real CLI before any of this was built on them: the interrupt lands **instantly**, and **the session survives it** — the next turn answers normally. So abandoning a turn costs the in-flight work and nothing else.

That makes the aggressive answer affordable, and it is the default: a note written mid-turn interrupts, and is redelivered as the next turn with framing that says not to resume.

```
(when claude-client--interrupted
  (concat "Your previous turn was interrupted before it "
          "finished, on purpose. Do not resume it; re-plan "
          "around the note below.\n\n"))
```

That framing is load-bearing rather than decorative. Without it the model picks up the abandoned work instead of redirecting; the mutation that drops the text fails two tests. Live, the behaviour is unambiguous — a 2000-word essay cut mid-sentence, then:

```
— error_during_execution —
— 1 note to carry forward —
REDIRECTED
```

The lost partial work is the price of acting immediately. That trade is measured rather than argued, which the earlier queue-until-turn-end behaviour could not establish: with a one-shot runner there was no turn to steer.

11.12 Backpressure: three defects, found by probing

Making a note interrupt exposed three ways rapid notes misbehaved. None were visible by reading the code; all three came from writing a probe that exercised the case.

A second note fired **another interrupt** into a turn already dying, and logged a second interrupted event for one interruption. The interrupt flag now doubles as the guard:

```
((and running claude-client-note-interrupts
      (not claude-client--interrupted)
      (process-live-p claude-client--process))
 ...)
```

An exact repeat **queued twice** and was sent twice, spending context to say nothing new. Repeats are coalesced — though both are still recorded, because the human did write it twice.

The queue was **unbounded**, so a slow model meant notes piling up until they arrived as one enormous prompt with the newest thought behind everything stale. `claude-client-max-pending-notes` bounds it and drops the oldest, because the recent note is the current intent. Drops are logged rather than silent: the record should show that something was said and lost.

11.13 Guarding against a second run

```
(when claude-client--turn-active
  (user-error "A Claude run is already active here; `k' to kill it first"))
```

with the reasoning right above it:

One conversation buffer, one subprocess. Silently restarting would orphan the running CLI — and any ediff it is blocked on, which then waits for a human whose answer no longer goes anywhere.

That is the kind of second-order consequence worth capturing in a comment: the bug is not “two processes”, it is “an orphaned ediff nobody can usefully answer”.

11.14 The agent-backend methods

Everything above is the client’s own machinery. The generic interface from Chapter 8 is satisfied at the bottom of the file, and most methods are one line delegating to a command that already existed — `agent-backend-send` to the `send` path, `agent-backend-interrupt` to `claude-client-interrupt`, and so on.

Three are worth reading:

- **agent-backend-note-policy** returns `:interrupt`, a value the base class does not define. The docstring for the `note-policy` slot explicitly allows this: backends may add policy values beyond `:steer` and `:queue`. See “Interruption: the design question, answered by measurement” above for what that policy actually does.
- **agent-backend-query** shells out to `claude -p`, the CLI’s one-shot mode. No session state is created — which is exactly the contract the base class describes.
- **agent-backend-resume** is implemented here, where the base class default is a `user-error`. This backend can reopen a past session because the CLI writes them to disk; `opcode` cannot.

11.15 The major mode

```
(define-derived-mode claude-client-mode agent-backend-mode "claude"
  "Major mode for the terminal-free Claude conversation buffer.
  Derives from `agent-backend-mode' (issue #41), so the shared keymap is
  inherited and the buffer-local `agent-backend--instance' holds a fresh
  `claude-client-backend' bound to this buffer."
  (setq-local truncate-lines nil)
  (setq-local agent-backend--instance
    (make-instance 'claude-client-backend
                   :buffer (current-buffer))))
```

The mode function is where the backend instance comes from: entering the mode constructs one and stores it buffer-locally. That single `setq-local` is what makes the buffer discoverable by `agent-backend--conversation-buffers` and actionable by every generic in Chapter 8.

11.15.1 Two keymaps, and why

The mode inherits `agent-backend-mode-map`’s C-c-prefixed vocabulary and adds its own single-letter keys. The comment explains the split: the single-letter keys are only reachable in plain Emacs, because under `evil` they are shadowed by the normal-state map, while the C-c vocabulary works either way.

`claude-client--setup-evil` re-registers a *deliberate subset* in `evil`’s normal and motion states. `k` and `g` are left out, and the reasoning is a good example of the codebase’s habit of recording the incident rather than just the fix:

Both are vim motions (`evil-previous-line`, the `gg` prefix) and both had destructive commands behind them here — `k` killed the CLI process and `g` erased the log and started a new session — so moving the cursor up in a conversation silently ended it, and `gg` was dead because `g` consumed the prefix.

Under `evil` they stay motions; the commands remain reachable as C-c C-q and C-c C-r. The

same lesson appears in *agent-session-overview.el* (Chapter 13), where *k* got a confirmation prompt instead.

There is one further layer of shadowing, and it needs its own opt-out:

```
;; evil-snipe binds `s' in a *minor* mode map, which outranks even an
;; evil-registered major-mode map, so `s' would still snipe.
(with-eval-after-load 'evil-snipe
  (when (boundp 'evil-snipe-disabled-modes)
    (add-to-list 'evil-snipe-disabled-modes 'claudio-client-mode)))
```

Minor-mode maps outrank major-mode maps, including evil-registered ones — so registering *s* in evil’s normal state was not enough. *evil-snipe* provides a disable list for exactly this case, and *magit*, *dired*, and *treemacs* are on it by default.

Elisp Feature: **define-derived-mode** and **special-mode**

```
(define-derived-mode CHILD PARENT LIGHTER DOCSTRING &rest BODY)
```

generates a mode function that runs the parent’s setup, then the body, then the mode hook. It also defines *CHILD-map* (inheriting the parent’s keymap), *CHILD-syntax-table*, and *CHILD-hook* if they do not already exist.

special-mode is the conventional parent for read-only, non-file buffers: buffer becomes read-only, *q* buries it, *g* is bound for revert. It is *agent-backend-mode*’s parent, so this buffer inherits that behaviour at one remove — which is why these buffers need `(let ((inhibit-read-only t)) ...)` around every insertion.

Both keymaps here are built with *make-sparse-keymap*: a list-based keymap, appropriate when few keys are bound. The alternative, *make-keymap*, allocates a char-table for the full character range.

A derived mode’s map automatically inherits its parent’s, so *claudio-client-mode-map* need only define what it adds to *agent-backend-mode-map*.

12 opencode-client.el — The Second Backend

915 lines driving opencode’s local HTTP API over plz, consuming its Server-Sent Events stream to render the conversation incrementally.

Architecturally this file earns its place by being *different enough*: it is the reason agent-backend.el exists, and the test of whether that abstraction actually abstracts. Where claude-client owns a subprocess and frames NDJSON off its stdout, this owns an HTTP client and consumes SSE from a server that may outlive Emacs.

It does not reimplement editor tools. Integration goes the other way: opencode is pointed at the mcp-emacs MCP server through opencode.json, so the agent gets the editor surface from Chapter 5 and this file only handles the conversation.

12.1 The shape

```
opencode serve  —HTTP→  requests (send, interrupt, session CRUD)
                —SSE→  event stream → conversation model → render
```

plz is a **soft** dependency, and unusually strict about it: it is loaded only when present *and* required only when a client command is actually invoked, so installing mcp-emacs never hard-requires it.

12.2 One server per project, on parallel ports

The design decision most worth reading is the server registry:

```
(defvar opencode-client--servers nil
  "Alist mapping project directory -> `opencode-client-backend' instance.")
```

Rather than one server on a fixed port, each project gets its own server on its own port, so sessions stay project-scoped and cannot collide. opencode-client--ensure-server looks up default-directory, which means switching projects transparently uses that project’s server and its sessions.

Finding the port is a bind-and-release probe upward from opencode-client-port, with the race written down rather than hidden:

```
;; The probe binds and releases a socket, so a parallel server starting
;; concurrently could race -- acceptable for an interactive tool.
```

That is the honest form of a known limitation: state it, scope it (“an interactive tool”), move on.

Starting the server has a macOS-specific path — when opencode-client-launchd-label is set *and* system-type is darwin, the server is kickstarted as a launchd agent instead of spawned as a child process, so it survives Emacs restarting. Otherwise it is an ordinary start-process.

Either way the client then polls for health against a 20-second deadline and refuses rather than registering a server that never came up.

12.3 What its agent-backend methods reveal

The method list at `elisp/opencode-client.el:632` is the interesting half of the abstraction, because it differs from `claude-client`’s in exactly the places the two agents genuinely differ:

Method	<code>opencode</code>	<code>claude-client</code>
<code>reply-permission</code>	implemented — it has a permission API	inherits the no-op; gating goes through the CLI hook
<code>reply-question</code>	implemented	inherits the no-op
<code>seed-history</code>	implemented — history comes from the server	inherits the no-op
<code>project-root</code>	implemented	—
<code>resume</code>	refuses (inherits)	implemented
<code>query</code>	opens a session, asks, deletes it	shells out to <code>claude -p</code>

Read the query row together with Chapter 8’s account of that generic: the base class promises “answer once, outside any session”, and `opencode` — which is session-only over HTTP — satisfies it by making the session’s lifetime an implementation detail. The effect is one-shot even though the transport is not.

The resume row is the inverse. The base class default is a user-error, and this backend simply keeps it, because there is nothing on disk to reopen. That is the “optional verbs degrade quietly” rule paying off: no stub, no apology, one accurate error message.

12.4 The gap the overview has to work around

Two omissions here matter downstream, and Chapter 13 has to accommodate both:

- `opencode` keeps no per-buffer turn state, and never publishes `finished` (its `step-finish` signal is discarded upstream), so nothing can honestly report whether a turn is running.
- Its buffers are named `*opencode:<title>*`, with no project in the name — unlike the two runners’ parallel `*claude-client:<project>:<n>*` and `*claude:<project>:<n>*` schemes.

Neither is a bug so much as a consequence of the transport. The overview handles them by reporting liveness only, and by deriving the project from `default-directory`.

13 agent-session-overview.el — One Live View

441 lines, and the payoff for the event model.

Each backend answers only for itself, and only when asked: `claude-client-list` and `mcp-emacs-run-list` each print their own sessions into the echo area, which is stale the moment it is printed. This buffer shows all of them at once and keeps itself current, so “is that turn still running?” is answered by looking rather than by switching buffers to check.

It is also the second subscriber to `agent-backend-event-functions`, which is what makes the event log in Chapter 11 a *log* rather than a rendering callback with extra steps.

13.1 Enumeration without a registry, again

The same choice as `agent-backend--conversation-buffers`, made independently and for the same reason — a `buffer-list` scan, no registry:

```
(defconst agent-session-overview--backends
  '( (claude-client
      :label "claude"
      :mode claudes-client-mode
      :regexp "\\`\\*claude-client:\\`(.+\\`):\\`([0-9]+\\`)*\\`"
      :project-in-name t
      :state agent-session-overview--claude-state)
    (eat ...)
    (opcode ...)))
```

One table drives recognition, labelling, and state-reading for all three backends. Two details in it are load-bearing.

Mode is checked before name. A buffer belongs to a backend if its major mode derives from `:mode`, or — when that is nil, or the buffer predates it — if its name matches `:regexp`. The comment gives the case that forces this order: a conversation buffer created before the current naming scheme (plain `*claude-client*`, no project or number) is still a live session and must not be missed just because its name is old. The eat runner has no mode of its own (it is an eat terminal), so it is recognised by name only.

The regexps live here, not in the backends. They are the backends’ own naming schemes, kept in this table rather than reached for through their internals, so this module has no load-order dependency on any of them.

13.2 Reporting only what a backend can actually observe

This is the chapter’s best idea, and it is a refusal:

Backend	States	Why
claude-client	working / idle / finished	publishes turn events, keeps --turn-active and a process handle
eat	live / dead	a TUI with no turn events at all
opencode	live / dead	keeps no per-buffer turn state, never publishes finished

And the reason the other two are not made to look as good:

Inferring working/idle from terminal output activity would misreport a long model think as idle and a spinner as work, so it is not done.

A dashboard that guesses is worse than one that abstains, because a guess is indistinguishable from a measurement once it is rendered in a column. The ? help buffer repeats this to the user, in the buffer itself — the honesty is part of the UI, not just the source.

13.3 The subscriber swallows its own errors

```
(defun agent-session-overview--on-event (_buffer event)
  (when (memq (plist-get event :kind)
    '(started prompt finished interrupted resumed error))
    (condition-case err
      (agent-session-overview--render)
      (error
        (message "agent-session-overview: render failed: %s"
          (error-message-string err))))))
```

Two things worth copying. It filters to the kinds that change what a row would show — a text event per token would re-render the table constantly for no visible difference. And it guards its own errors, because this runs *inside the publishing session's own call stack*: a rendering failure here must not propagate back into the session that published the event. The Org transcript in Chapter 14 follows the same rule, and Chapter 23 lists it as an idiom.

13.4 Two window-management traps

Both are the kind of thing that only shows up under a popup-managing framework, and both have long comments because the symptom is baffling.

k used to end sessions by accident. It is bound to `agent-session-overview-quit-session`, and it is also evil's `evil-previous-line` — so moving the cursor up a list of sessions silently and unrecoverably ended one. The fix is a `yes-or-no-p`, chosen explicitly as *the cheapest fix that keeps the binding where the muscle memory expects it* rather than moving the key.

More generally, `agent-session-overview--setup-evil` re-registers the mode's keys in evil's normal and motion states, because evil's state maps outrank a major mode's own keymap: without it ? searches backward, k moves up, i enters insert state, and g is a prefix map. g and q are deliberately left alone — g r already reverts, and q already quits.

Reading the help used to dismiss the thing it described. Under Doom’s +popup, both the overview and its help buffer match popup rules, with-help-window selects the help window, and q there is +popup/quit-window — which tears down the popup stack rather than the selected window. So ? then q closed the overview too.

The fix is a display-buffer action function that splits the *overview’s own* window, making the help an ordinary window whose q is the plain quit-window:

```
(let ((display-buffer-alist
      (cons `((regexp-quote agent-session-overview--help-buffer-name)
              (agent-session-overview--display-help))
            display-buffer-alist)))
  (with-help-window ...))
```

Prepending to a *local* copy of display-buffer-alist is how this outranks any user or framework rule without mutating the user’s configuration — the same tactic as mcp-emacs-run--display-popup in Chapter 10.

? toggles rather than stacking, and the toggle-off path deletes the window instead of calling quit-window, because the window is one this command split off itself; quit-window would leave the split showing whatever was there before, which is the same window leak from the other side.

13.5 One list drives three surfaces

```
(defconst agent-session-overview--help
  '(("RET" "visit the session")
    ("k"  "quit it (asks first)")
    ...))
```

The docstring states the invariant: *one list drives the header-line hint, the ? buffer, and the keymap, so a new action cannot be documented in one place and missing from another.* --setup-evil iterates the same list. This is the “derive state, do not store it” idiom from Chapter 23 applied to documentation.

The key hint lives in the mode line rather than a header row, and the comment explains the constraint: only tabulated-list-print aligns the column headers with the data, so an in-buffer header row would not be padded to the column widths. The mode line is still always visible and still costs no rows.

14 mcp-emacs-remote.el and mcp-emacs-run-resume.el

Two smaller modules that show off two different integration techniques.

14.1 Remote prompt input

The straightforward half of mcp-emacs-remote.el:

```
(defun mcp-emacs-remote-prompt (&optional initial)
  (interactive
    (list (when (use-region-p)
              (buffer-substring-no-properties (region-beginning) (region-end))))
    (mcp-emacs-remote--send (read-string "Claude prompt: " initial)))
```

Elisp Feature: interactive specifications

(interactive) marks a function as a command. The argument determines how arguments are gathered when invoked interactively:

- (interactive) — no arguments.
- (interactive "sPrompt: ") — a *code string*; s reads a string, p a numeric prefix, r region bounds, f an existing file, and so on. Multiple codes are separated by \n.
- (interactive (list EXPR ...)) — a *form* evaluated to produce the argument list. This is the flexible version, used here to seed from the region.

The key property: when the function is called from Lisp, the interactive form is not evaluated at all. So mcp-emacs-remote-prompt can be called programmatically with an explicit initial, and interactively with a region-derived default.

14.2 The Org transcript via advice

The other half records what the agent is doing into an Org buffer. For the IDE surface it does so with **advice**:

```
(defun mcp-emacs-remote-enable ()
  (interactive)
  (setq mcp-emacs-remote-enabled t)
  (advice-add 'mcp-emacs-ide--call-tool :before #'mcp-emacs-remote--tap-call-tool)
  (advice-add 'mcp-emacs-ide--complete-open-diff :before
    #'mcp-emacs-remote--tap-complete-open-diff)
  (advice-add 'mcp-emacs-ide-start :after #'mcp-emacs-remote--tap-start)
  (when (fboundp 'mcp-emacs-ide-stop)
```

```
(advice-add 'mcp-emacs-ide-stop :before #'mcp-emacs-remote--tap-stop))
(add-hook 'agent-backend-event-functions #'mcp-emacs-remote--tap-runner-event))
```

Elisp Feature: advice

Advice wraps an existing function without editing its definition:

```
(advice-add 'target :before #'my-fn) ; run my-fn first, ignore its value
(advice-add 'target :after #'my-fn) ; run my-fn after
(advice-add 'target :around #'my-fn) ; my-fn receives the original as its
                                     ; first arg and decides when to call it
(advice-add 'target :override #'my-fn) ; replace entirely
(advice-add 'target :filter-return #'my-fn) ; transform the return value
(advice-remove 'target #'my-fn)
```

`:before` and `:after` advice receive the same arguments as the target and their return values are discarded, which makes them safe for pure observation — exactly what a transcript needs.

Advice is the standard Emacs mechanism for extending code you do not control. It is also a maintenance liability: it depends on the target's *name and arity*, both private implementation details. Rename `mcp-emacs-ide--call-tool` or change its signature and the advice silently stops matching or starts erroring.

That is precisely why the newer `claude-client` integration is a hook subscription instead.

Every tap is wrapped in `ignore-errors`, with the reasoning stated:

```
(defun mcp-emacs-remote--tap-runner-event (buffer event)
  "Subscriber for `agent-backend-event-functions'."
  Errors are swallowed: a transcript problem must never take down the
  runner whose events it is observing."
  (ignore buffer)
  (when mcp-emacs-remote-enabled
    (ignore-errors (mcp-emacs-remote--record-runner-event event)))))
```

A logging layer that can crash the thing it logs is worse than no logging.

Note also `(ignore buffer)` — a no-op call whose only purpose is to tell the byte-compiler “yes, I know this parameter is unused”. The alternative is to name it `_buffer`; both appear in this codebase.

14.3 The human's side of the same record

The runner's activity was in the log; the human's Org edits were not. The four `org-task` writes now emit as well, and the transcript subscribes to both, so one record shows both actors in order:

```
- [13:18:34] org session :: STRT [session.org]
- [13:18:34] assistant  :: AI working
- [13:18:34] org item   :: first item → DONE [session.org]
- [13:18:34] result    :: Status: applied
```

The design rule behind this is worth stating plainly, because it is what keeps the log from becoming a liability:

The Org file is the aggregate. The log observes it.

These events are *observations*, not *commands*. They are emitted because Org changed, and nothing may reconstruct state by replaying them. Session and item state lives in the Org subtree the human edits directly — reorder, flip a keyword, rewrite text — and the log is a record beside it, never a competing source of truth.

Two consequences fall out of that. First, the log does not need to be durable: `claude-client--events` being buffer-local and lost on restart is a deliberate property, not a gap, because state does not live there. Second, ordering mostly stops being a problem — with Org as the single writer of record there is nothing to merge, and the only genuine two-writer case left is `FileChanged`, which the ediff gate already handles.

Only *successful* writes emit. A rejected keyword or an item that could not be identified changed nothing, so there is nothing to observe:

```
(if (mcp-emacs-org-task--find-item ref)
    (progn
      (org-todo status)
      (mcp-emacs-org-task--emit 'item-status path :ref ref :status status)
      ...)
    (format "TODO item not found: %s" ref)) ; no emit
```

And the emit site catches subscriber errors itself rather than leaving them to the caller — the org-task functions' outer handler only catches user-error, so an unguarded observer error would escape as a plain error and take the write down with it. The aggregate comes first, the record second.

14.4 Quiet tools

```
(defconst mcp-emacs-remote--quiet-tools '("getDiagnostics" "closeAllDiffTabs")
  "Tool names whose calls are recorded compactly rather than as full entries.
  These fire before every edit and would otherwise dominate the transcript.")
```

Recall from Chapter 9 that Claude Code calls `closeAllDiffTabs` and `getDiagnostics` before *every* native edit. Recording each as a full Org heading with a JSON source block would bury the actual content. They get a one-line note instead.

14.5 The resume picker: reading someone else's data format

`mcp-emacs-run-resume.el` builds a native completing-read over past Claude sessions, which live at:

```
~/ .claude/projects/<slug>/<session-id>.jsonl
```

where the slug is the project path with `/` and `.` replaced by `-`:

```
(defun mcp-emacs-run-resume--slug (root)
  (let ((path (directory-file-name (expand-file-name root))))
    (replace-regexp-in-string "[/\\.]" "-" path)))
```

The interesting part is extracting a *useful label*. A transcript's literal first `type:user` entry is usually not the human's prompt — it is an injected `<command-message>` block, a caveat, or a tool result. So:

```
(defun mcp-emacs-run-resume--noise-p (text)
  "Return non-nil when TEXT is an injected block, not a human prompt."
  (let ((s (string-trim (or text ""))))
    (or (string-empty-p s)
        (string-prefix-p "<command-" s)
        (string-prefix-p "<local-command-" s)
        (string-prefix-p "Caveat:" s))))
```

and the scan is bounded:

```
(with-temp-buffer
  (insert-file-contents file nil 0)
  ...
  (while (and (not prompt)
              (< n mcp-emacs-run-resume-head-lines) ; 40 by default
              (not (eobp))))
  ...))
```

Elisp Feature: bounded `insert-file-contents`

```
(insert-file-contents FILENAME VISIT BEG END REPLACE)
```

`BEG` and `END` are *byte* offsets. `(insert-file-contents file nil 0)` with no `END` reads the whole file — the `0` here is just an explicit start. Passing a real `END` is how you read only a prefix of a large file, which matters when transcripts can be megabytes and you only need the head.

The 40-line window is the actual bound in this code; it is tuned to clear injected caveat blocks without reading whole transcripts.

Content normalisation handles the two shapes a message content field can take:

```
(defun mcp-emacs-run-resume--content-text (content)
  (cond
    ((stringp content) content)
    ((and (or (vectorp content) (listp content)) (> (length content) 0))
     (let ((blocks (append content nil)) text)
       (while (and blocks (not text))
         (let ((b (pop blocks)))
           (when (and (listp b) (stringp (alist-get 'text b)))
             (setq text (alist-get 'text b)))))
       text))
    (t nil)))
```

`(append content nil)` converts a vector to a list — a common Elisp idiom for normalising sequence types before iterating.

15 orgspec — Concept and Data Model

The remaining orgspec chapters describe what is effectively a second project sharing the same host process. It is a port of the load-bearing core of [OpenSpec](#) — a TypeScript tool over Markdown — into Org and Emacs Lisp.

15.1 The workflow

A *change* is a spec-of-intent. You describe what should change, implement it, then **fold** that description into an accumulating source of truth.

```
orgspec/
├── specs/<area>.org          the accumulating truth
                             L1 = requirement, L2 = scenario
└── changes/
    ├── <id>/change.org      one change
                             Intent / Scope / Approach / Tasks + * Delta
                             L2 under Delta = requirement (op-tagged)
                             L3 = scenario
    └── archive/<id>/        folded changes
```

A delta requirement is a level-2 headline under ** Delta*, tagged with exactly one op and carrying an `:AREA:` property naming its target spec:

```
* Delta
** TODO Notify on logout          :ADDED:
:PROPERTIES:
:AREA: auth
:END:
The system SHALL notify the user on account logout.
*** Lockout triggered
- GIVEN three failed logins
- WHEN the fourth fails
- THEN a notification is sent
```

15.2 The marker table

`orgspec.el` (95 lines) exists so that no downstream file hardcodes a marker:

```
(defconst orgspec-op-tags
  '(("ADDED" . added)
    ("MODIFIED" . modified)
    ("REMOVED" . removed)
    ("RENAMED" . renamed)))
```

```
(defconst orgspec-area-property "AREA")
(defconst orgspec-from-property "FROM")
(defconst orgspec-impl-drawer "IMPL")

(defconst orgspec-clarification-regexp "\\[NEEDS CLARIFICATION: [^]]*\\]")
(defconst orgspec-normative-regexp "\\<\\(SHALL\\|MUST\\)\\>")

(defcustom orgspec-todo-blocked "WAIT" ...)
(defcustom orgspec-todo-active "STRT" ...)
(defcustom orgspec-todo-removed "KILL" ...)

(defconst orgspec-fold-order '(renamed removed modified added))
```

Two accessors bridge the string/symbol boundary:

```
(defun orgspec-op-from-tags (tags)
  (seq-some (lambda (cell) (and (member (car cell) tags) (cdr cell)))
    orgspec-op-tags))

(defun orgspec-op-tag-name (op)
  (car (rassq op orgspec-op-tags)))
```

`rassq` is `assq` searching by the *cdr* — a reverse lookup in an alist. `seq-some` returns the first non-`nil` result of applying a function across a sequence.

Note that the TODO keyword mappings are `defcustom`, not `defconst`: `STRT`, `WAIT`, `KILL` are Doom's defaults, and a user with a different workflow changes these three variables rather than patching code.

15.3 The model: three structs

```
(cl-defstruct (orgspec-scenario (:constructor orgspec-scenario-create))
  name      ; headline text – its identity
  body)     ; raw GIVEN/WHEN/THEN text, verbatim

(cl-defstruct (orgspec-requirement (:constructor orgspec-requirement-create))
  name      ; headline text – identity
  op        ; added/modified/removed/renamed, or nil for a spec requirement
  area      ; target spec, from :AREA:
  from      ; previous name for a rename, from :FROM:
  body      ; prose (must carry SHALL/MUST for ADDED/MODIFIED)
  (scenarios nil)
  (impl nil) ; raw lines from the :IMPL: drawer
  source)   ; the raw org subtree text

(cl-defstruct (orgspec-change (:constructor orgspec-change-create))
  id
  (requirements nil))
```

The identity rules are stated in the commentary and are load-bearing:

- **Requirement identity is its exact headline text.** Renames are handled explicitly through

the `:FROM:` property.

- **Scenario identity is its headline text**, used by the MODIFIED drop guard.

The source slot deserves attention. It holds the requirement's *raw org subtree text*, used verbatim by the fold's `org-paste-subtree`. The parsed fields (`body`, `scenarios`) are for validation and inspection; the fold re-pastes the original text so nothing is lost in a parse/serialise round-trip — drawers, formatting, inline markup, all preserved exactly.

15.4 Deliberate omissions

The README lists what *orgspec* drops relative to OpenSpec: conflict-free parallel changes, artifact DAGs, bulk operations. Those serve multi-developer, multi-artifact scale. And the project memory records one more decision: the `:IMPL:` writer was a *won't-do* — automatic requirement-to-code traceability links rot faster than they pay off. The fold still preserves a hand-written `:IMPL:` drawer, because a human who chose to write one meant it.

16 orgspec — Parsing with org-element

orgspec-parse.el, 129 lines. The commentary explains why this file is short:

Going org-native removes the two hardest OpenSpec ports: org-element already sees code blocks/drawers and headline levels correctly, so there is no code-fence mask and no hand-rolled section-tree parser.

The Markdown original had to mask out fenced code blocks (so a # inside a code sample is not read as a heading) and build its own tree from flat heading lines. Org's parser does both.

16.1 The org-element AST

Elisp Feature: org-element

(org-element-parse-buffer) returns an abstract syntax tree of the whole buffer. Each node is a list whose car is the element type:

```
(headline (:raw-value "Notify on lockout" :level 2 :tags ("ADDED") ...) CHILDREN...)
```

The accessors:

- (org-element-type el) — the type symbol: headline, paragraph, drawer, section, plain-list, src-block, ...
- (org-element-property :level el) — a property from the node's plist.
- (org-element-contents el) — the child nodes.
- (org-element-map TREE TYPE FUN) — walk the tree, calling FUN on each node of the given type.

Standard properties include :raw-value (headline text without stars, keyword, or tags), :level, :tags, :todo-keyword, :begin, :end, :contents-begin, :contents-end.

Org properties become element properties with a colon prefix and an upcased name: an :AREA: auth drawer entry is reachable as (org-element-property :AREA headline). This code computes the keyword rather than hardcoding it:

```
(org-element-property (intern (concat ":" orgspec-area-property)) headline)
```

so renaming the property in the marker table is enough.

16.2 The section gotcha

This is the one thing that trips everyone up, and it is documented:

```
(defun orgspec-parse--section-elements (headline)
  "Return the elements in HEADLINE's own section (its direct prose/drawers).
`org-element' nests a headline's non-headline content under a `section'"
```

```
element; descend into it so paragraphs and drawers are visible."
(let ((section (seq-find (lambda (el) (eq (org-element-type el) 'section))
                          (org-element-contents headline))))
  (and section (org-element-contents section))))
```

A headline's contents are **not** its paragraphs and drawers. They are: an optional section element wrapping all the non-headline content, followed by child headlines. Iterating a headline's contents looking for paragraphs finds nothing. You must descend into the section first.

16.3 Extracting body, drawers, scenarios

With that in hand the extractors are direct:

```
(defun orgspec-parse--body (headline)
  (let ((parts '()))
    (dolist (el (orgspec-parse--section-elements headline))
      (when (memq (org-element-type el) '(paragraph plain-list))
        (push (string-trim
                (buffer-substring-no-properties
                 (org-element-property :contents-begin el)
                 (org-element-property :contents-end el)))
              parts)))
    (string-trim (string-join (nreverse parts) "\n"))))
```

Only paragraph and plain-list count as body — drawers, property drawers, and child headlines are excluded by construction.

```
(defun orgspec-parse--impl (headline)
  (let (lines)
    (dolist (el (orgspec-parse--section-elements headline))
      (when (and (eq (org-element-type el) 'drawer)
                  (equal (org-element-property :drawer-name el)
                         orgspec-impl-drawer))
        (setq lines (seq-remove #'string-empty-p
                                (mapcar #'string-trim
                                        (split-string ... "\n" t))))))
    lines))
```

(split-string STR "\n" t) — the trailing t is OMIT-NULLS, dropping empty strings from the result.

Scenarios are simply the child headlines:

```
(defun orgspec-parse--scenarios (headline)
  (let (scenarios)
    (dolist (child (org-element-contents headline))
      (when (eq (org-element-type child) 'headline)
        (push (orgspec-scenario-create
                :name (org-element-property :raw-value child)
                :body (orgspec-parse--body child))
              scenarios)))
    (nreverse scenarios)))
```

16.4 The two entry points

```
(defun orgspec-parse-change (&optional id)
  (let ((tree (org-element-parse-buffer)
          requirements)
        (org-element-map tree 'headline
          (lambda (hl)
            (when (and (= (org-element-property :level hl) 1)
                        (equal (org-element-property :raw-value hl) "Delta"))
              (dolist (child (org-element-contents hl))
                (when (eq (org-element-type child) 'headline)
                  (push (orgspec-parse--requirement child) requirements))))))
    (orgspec-change-create :id id :requirements (nreverse requirements))))
```

Find the L1 headline literally named Delta; its headline children are the delta requirements. `orgspec-parse-spec` is the same shape without the Delta wrapper and without ops: L1 headlines are requirements, L2 are scenarios.

Note that `orgspec-parse--requirement` also captures the raw source region:

```
:source (buffer-substring-no-properties
  (org-element-property :begin headline)
  (org-element-property :end headline))
```

`:begin/ :end` span the whole subtree including children, which is exactly what `org-paste-subtree` needs later.

17 orgspec — The Fold

orgspec-fold.el, 173 lines, and the algorithmic heart of the project.

17.1 The four design commitments

From the commentary:

1. **Apply order is fixed:** RENAMED → REMOVED → MODIFIED → ADDED.
2. **Subtree surgery on org buffers, not text splicing.**
3. **Validate-all-then-write-all:** every target spec is built in a temp buffer and the whole set validated before anything is written, so a late failure leaves the on-disk specs untouched.
4. **MODIFIED must not silently drop a scenario** present in the current spec requirement.

Each is a decision that could plausibly have gone another way, and each has a concrete failure mode behind it.

17.2 Why the order

Consider a change that removes Thing and adds a new Thing with the same name — a common “replace this requirement wholesale” pattern. If ADDED ran first, it would find the existing Thing and signal “already exists”. With REMOVED before ADDED, the old one is gone by the time the add runs.

Similarly, RENAMED runs first so that a subsequent MODIFIED can address the requirement by its *new* name.

The fold test asserts exactly this (test/orgspec-fold-test.el).

17.3 Subtree surgery

The paste operation:

```
(defun orgspec-fold--paste-requirement (req)
  "Paste delta requirement REQ into the current spec buffer at L1, cleaned."
  (org-paste-subtree 1 (orgspec-requirement-source req))
  (save-excursion
    (org-back-to-heading t)
    (orgspec-fold--clean-pasted)))
```

Elisp Feature: Org subtree operations

Org provides a family of structural editors that operate on whole subtrees:

- (org-paste-subtree LEVEL TREE) — insert TREE (a string of org text) at point, **re-leveiling** the whole subtree so its root sits at LEVEL. This is the crucial one:

a delta requirement is stored at L2 with L3 scenarios, and pasting at level 1 shifts everything up by one, giving the L1/L2 shape specs use. Emacs does the arithmetic.

- (org-cut-subtree) — remove the subtree at point into the kill ring.
- (org-end-of-subtree INVISIBLE-OK T0-HEADING) — move past the whole subtree.
- (org-back-to-heading INVISIBLE-OK) — move to the enclosing headline.
- (org-edit-headline TEXT) — replace the headline text, leaving keyword, priority, and tags intact.
- (org-narrow-to-subtree) — restrict the buffer to this subtree.

Using these instead of regexp text-splicing is why the fold does not have to know anything about star counts, and why a scenario body containing a line that *looks* like a heading cannot corrupt the result.

The cleanup after a paste:

```
(defun orgspec-fold--clean-pasted ()
  "Strip the TODO keyword, the op tag, and the `:AREA:' routing property;
keep the `:IMPL:' drawer."
  (org-back-to-heading t)
  (org-todo 'none)
  (org-set-tags nil)
  (org-entry-delete (point) orgspec-area-property))
```

The policy encoded here: workflow state (TODO keyword) and routing metadata (op tag, :AREA:) belong to the change, not to the accumulated spec. Provenance (:IMPL:) is durable and survives.

17.4 The four operations

```
(defun orgspec-fold--apply-renamed (req)
  (let ((from (orgspec-requirement-from req))
        (to (orgspec-requirement-name req)))
    (unless from (signal 'orgspec-fold-error (list ...)))
    (unless (orgspec-fold--find-requirement from)
      (signal 'orgspec-fold-error (list (format "renamed source %S not found" from))))
    (org-edit-headline to)))

(defun orgspec-fold--apply-removed (req)
  (unless (orgspec-fold--find-requirement (orgspec-requirement-name req))
    (signal 'orgspec-fold-error (list ...)))
  (org-cut-subtree))

(defun orgspec-fold--apply-added (req)
  (when (orgspec-fold--find-requirement (orgspec-requirement-name req))
    (signal 'orgspec-fold-error (list (format "added %S already exists in spec" ...))))
  (goto-char (point-max))
  (orgspec-fold--paste-requirement req))
```

Every op *verifies its precondition and signals* rather than silently doing nothing. A REMOVED whose target does not exist means the change file and the spec have diverged, and that is worth

stopping for.

Elisp Feature: custom error types

```
(define-error 'orgspec-fold-error "orgspec fold error")
```

defines a new error symbol with a human-readable message, optionally with a parent (defaulting to error). It is signalled with data:

```
(signal 'orgspec-fold-error (list "renamed source \"X\" not found"))
```

and caught specifically:

```
(condition-case _
  (orgspec-fold-area spec reqs)
  (orgspec-fold-error 'that-specific-failure))
```

Defining a type rather than calling (error "...") lets callers — and, as Chapter 20 shows, tests — distinguish “the fold rejected this” from “something else went wrong”.

17.5 The MODIFIED drop guard

The one operation with real logic:

```
(defun orgspec-fold--apply-modified (req)
  (let ((name (orgspec-requirement-name req)))
    (unless (orgspec-fold--find-requirement name)
      (signal 'orgspec-fold-error (list (format "modified target %S not found" name))))
    (let* ((current (orgspec-fold--scenario-names-at-point))
           (incoming (mapcar #'orgspec-scenario-name
                             (orgspec-requirement-scenarios req)))
           (missing (seq-remove (lambda (s) (member s incoming)) current)))
      (when missing
        (signal 'orgspec-fold-error
          (list (format "modified %S drops scenario(s): %s"
                        name (string-join missing ", "))))))
    (org-cut-subtree)
    (orgspec-fold--paste-requirement req)))
```

MODIFIED *replaces* the spec requirement with the delta’s version. If the author of the change wrote out only the scenarios they cared about, the replacement would silently delete the rest — a real and quiet way to lose requirements. The guard makes it loud: every scenario currently in the spec must appear in the incoming delta.

This is a direct port of OpenSpec’s findMissingCurrentScenarios, and it is the single most valuable rule in the system.

17.6 The pure core

```
(defun orgspec-fold-area (spec-text requirements)
  "Return the new spec text for one area.
Pure: builds the result in a temp buffer and returns the string; writes
```

```
nothing."
(with-temp-buffer
  (let ((org-inhibit-startup t)) (org-mode))
  (insert (or spec-text ""))
  (dolist (op orgspec-fold-order)
    (dolist (req requirements)
      (when (eq (orgspec-requirement-op req) op)
        (goto-char (point-min))
        (orgspec-fold--apply-one req))))
  (buffer-string)))

;; strings in, string out
```

This is the whole fold: nested loops over ops and requirements, in a temp buffer, returning a string. No I/O. That purity is what makes the review feature (§19.3) possible *and* what makes the function trivially testable — Chapter 20 shows the tests exercising it with literal Org strings.

(let ((org-inhibit-startup t)) (org-mode)) skips Org’s per-buffer startup work (agenda file scanning, visibility restoration, hook chains) which is irrelevant to a temp buffer and measurably slow.

17.7 Grouping and the transaction boundary

orgspec-commands-fold-build is where purity meets the filesystem:

```
(defun orgspec-commands-fold-build (id)
  "Fold change ID's delta in memory; return an alist (SPEC-FILE . NEW-TEXT)."
  (orgspec-commands--assert-no-clarifications id)
  (let* ((change (orgspec-commands--read-change id))
        (reqs (orgspec-change-requirements change))
        (groups (orgspec-fold--group-by-area reqs))
        built)
    (orgspec-validate-change-or-signal change)
    (dolist (group groups)
      (let* ((area (car group))
            (file (orgspec-commands--spec-file area))
            (current (when (file-exists-p file)
                        (with-temp-buffer (insert-file-contents file)
                          (buffer-string)))))
        (push (cons file (orgspec-fold-area current (cdr group))) built)))
    (nreverse built)))
```

Gates first (clarifications, then the full validator), then build everything in memory, then return. And orgspec-archive writes:

```
(defun orgspec-archive (id)
  (let ((built (orgspec-commands-fold-build id)))
    (dolist (cell built)
      (make-directory (file-name-directory (car cell)) t)
      (with-temp-file (car cell) (insert (cdr cell))))
    (orgspec-commands--archive-move id)
    (mapcar #'car built)))
```

Every failure mode — a bad delta, a missing rename source, a dropped scenario — happens during `fold-build`, before the first `with-temp-file`. The write loop cannot fail halfway and leave three specs updated and two not. That is the “validate-all-then-write-all” discipline, and it is the reason `orgspec-fold-area` was made pure in the first place.

18 Building a Domain Model on Org

Chapters 13 and 14 describe what orgspec does. This chapter is about *why it is shaped that way*, because the shape is not obvious and the obvious alternative does not work. If you ever want to build something structured on top of Org — a spec workflow, a task system, a knowledge base with invariants — this is the pattern, and this is the reasoning behind it.

18.1 The question

You have Org documents. You want to treat them as domain objects: parse them into a model, apply operations with real invariants, write them back. The natural instinct, coming from XML or JSON or a database, is:

load → object graph → mutate objects → serialise → save

Org appears to offer exactly this. `org-element-parse-buffer` gives you a tree. `org-element-interpret-data` turns a tree back into text. Emacs 29 added mutators — `org-element-set`, `org-element-put-property`, `org-element-insert-before`, `org-element-extract`, `org-element-adopt`, `org-element-create`. All of them exist and all of them work:

<code>org-element-set</code>	YES
<code>org-element-put-property</code>	YES
<code>org-element-insert-before</code>	YES
<code>org-element-extract</code>	YES
<code>org-element-interpret-data</code>	YES
<code>org-element-create</code>	YES
<code>org-element-adopt</code>	YES

(verified on Emacs 30.2 / Org 9.7.11)

So the round-trip is available. **This project uses none of it.** Not one call to `org-element-interpret-data` or any mutator appears in `elisp/`. That is a deliberate choice, and the rest of this chapter is the argument for it.

18.2 Why Org is not a DOM

The decisive difference: **in Org, the buffer text is the model.** There is no canonical in-memory document that the text is a serialisation of. The AST is a *derived, cached view* of the text — the opposite relationship to a DOM, where the tree is authoritative and the text is one rendering of it.

Everything else follows from that inversion:

- The AST goes stale the moment the buffer is edited. `org-element` maintains a cache with invalidation precisely because of this.

- Elements carry `:begin / :end` buffer positions, which are meaningful only against the buffer they came from.
- Org's own editing commands — `org-todo`, `org-set-tags`, `org-promote` — operate on **point in a live buffer**, not on element objects. Org itself does not use the AST as a write path.

And the practical consequence: `interpret-data` is a *formatter*, not a serialiser. It renders an AST according to Org's current formatting preferences. It does not reproduce the input.

18.3 Measuring the round-trip

This is worth doing yourself rather than taking on trust. Take a delta requirement in this project's own format, parse it, interpret it back, and compare:

```
(with-temp-buffer
  (let ((org-inhibit-startup t)) (org-mode))
  (insert src)
  (let* ((tree (org-element-parse-buffer))
        (out (org-element-interpret-data tree)))
    (string= src out)))
```

Result: nil. The input:

```
* TODO Notify on logout                                :ADDED:
:PROPERTIES:
:AREA:  auth
:END:
:IMPL:
[[file:auth.el::x][x]]
:END:
The system SHALL notify.
```

```
** Lockout triggered
- GIVEN three failed logins
```

and what comes back:

```
* TODO Notify on logout                                :ADDED:
:PROPERTIES:
:AREA:      auth
:END:
:IMPL:
[[file:auth.el::x][x]]
:END:
The system SHALL notify.
** Lockout triggered
- GIVEN three failed logins
```

257 characters in, 277 out. Three changes, none of them requested:

1. **The tag column moved.** Org re-aligned `:ADDED:` to its own preferred column.
2. **The property value was re-padded.** `:AREA: auth` became `:AREA: auth`.
3. **A blank line vanished** — the one separating the requirement body from its first scenario.

Semantically identical. Textually different, in a file under version control, touched by a human who chose that formatting.

Now consider what that means for orgspec. The fold rewrites `specs/*.org` on every archive. If each fold reflowed every requirement in the file — not just the ones the change touched — every archive would produce a diff full of whitespace noise, the ediff review from §19.3 would be unreadable, and `git blame` on a spec would point at whoever last archived anything.

The fidelity loss is not a bug in Org. Interpreting an AST *has* to make formatting decisions, because the AST does not record the ones the author made. It is simply the wrong tool for “preserve what the human wrote”.

18.4 The pattern: four rules

What orgspec does instead, stated as rules you can apply elsewhere.

18.4.1 Rule 1 — Parse into your own structs, not the AST

The AST is Org’s vocabulary: headline, section, paragraph, drawer. Your domain has a different one: requirement, scenario, op, area. Build the latter (§15.3):

```
(cl-defstruct (orgspec-requirement (:constructor orgspec-requirement-create))
  name op area from body scenarios impl source)
```

Use `org-element` as an *extraction* tool and then leave it behind. The rest of your system — validator, fold, MCP handlers, tests — never sees an `org-element` node. It sees requirements.

This buys three things. Your invariants are expressible in your own terms (“an ADDED requirement must contain SHALL”). Your code is testable without a buffer. And Org’s parse representation becomes an implementation detail of one file, `orgspec-parse.el`, rather than a dependency threaded through everything.

18.4.2 Rule 2 — Carry the raw source

The slot that makes the whole thing work:

```
:source (buffer-substring-no-properties
         (org-element-property :begin headline)
         (org-element-property :end headline))
```

`:begin` and `:end` span the entire subtree — headline, drawers, body, children. Capturing that text verbatim means you never have to reconstruct it.

The parsed fields (`body`, `scenarios`, `op`) exist for *decisions*: does this validate, does this drop a scenario, which area does it route to. The `source` field exists for *writing*. Two representations, two jobs, and the fidelity problem disappears because the write path never round-trips through the parse path.

The struct docstring says this outright:

```
SOURCE is the requirement’s raw org subtree text, used verbatim by the fold’s org-
paste-subtree (so the pasted content matches the change file exactly, drawers and
all).
```

18.4.3 Rule 3 — Mutate with Org’s buffer API, in a temp buffer

Write by replaying the human’s own editing commands programmatically:

```
(org-paste-subtree 1 (orgspec-requirement-source req))
(org-todo 'none)
(org-set-tags nil)
(org-entry-delete (point) orgspec-area-property)
```

Note what `org-paste-subtree` is doing here. The requirement was stored at level 2 (under *Delta) with level-3 scenarios. Pasting at level 1 **re-levels the whole subtree**, giving the L1/L2 shape specs use. You do not count stars. Org does the arithmetic, correctly, including for content that merely looks like a heading.

The other structural operations follow the same principle:

Instead of	Use
regex for the next <code>^*</code>	<code>org-end-of-subtree</code>
deleting a line range	<code>org-cut-subtree</code>
rewriting the headline line	<code>org-edit-headline</code>
regex on <code>:PROP: value</code>	<code>org-entry-get / -put / -delete</code>
counting stars	<code>org-paste-subtree</code> with a level
scanning for children	<code>org-map-entries</code> with 'tree scope

Every one of these knows about drawers, inline tasks, code blocks, and inherited properties. Your regex does not.

18.4.4 Rule 4 — Keep the transformation pure

Do all of it in a temp buffer and return a string:

```
(defun orgspec-fold-area (spec-text requirements)
  "Return the new spec text for one area.
Pure: builds the result in a temp buffer and returns the string; writes
nothing."
  (with-temp-buffer
    (let ((org-inhibit-startup t)) (org-mode))
    (insert (or spec-text ""))
    (dolist (op orgspec-fold-order)
      (dolist (req requirements)
        (when (eq (orgspec-requirement-op req) op)
          (goto-char (point-min))
          (orgspec-fold--apply-one req))))
    (buffer-string)))
```

Strings in, string out. The buffer is a *computation device*, not a document being edited — which is a genuinely useful way to think about `with-temp-buffer` in Emacs generally.

Three things fall out of this purity, and they are the payoff for the whole pattern:

- **Testability.** A fold test is a literal Org string and a set of assertions on the returned string. No fixtures on disk, no mocks, no cleanup (§20.6).

- **Preview.** Because the transformation produces a value rather than an effect, `orgspec-review-fold` can ediff the result against the on-disk file and write nothing (§19.3). “See the fold, not trust the fold” is free.
- **Transactionality.** All areas are built in memory and only written once every fold has succeeded, so a late failure leaves specs/ untouched (§17.7).

None of those are available if your fold mutates a file buffer in place.

18.5 Where the model actually lives

Worth being explicit, because it is unusual: `orgspec` has **no persistent model**. There is no database, no cache file, no serialised index. The structs exist only for the duration of one operation.

```
change.org  —parse→  orgspec-change  —fold→  new spec text  —write→  spec.org
  (truth)           (transient)           (transient)           (truth)
```

Org files on disk are the only durable state. The model is a lens you look through, briefly.

That is why `orgspec-status` counts checkboxes with a regex over the file rather than consulting an index, and why the agenda dashboard (§19.2) is a *query over the change files themselves*, not over a projection. Nothing can be stale because nothing is stored — the same “derive, do not store” idiom the runner uses for its session registry (§10.2), applied to the domain model.

18.6 What you give up, and when to reconsider

The pattern has real costs.

Re-parsing every time. `orgspec-commands--read-change` re-reads and re-parses the change file for every operation. At this scale — one file, a few dozen requirements — that is microseconds. At thousands of files it would not be.

No incremental update. The fold rewrites a whole spec file. Fine for documents a human reads; wrong for a log that only ever grows.

Cross-file invariants are awkward. “No requirement name may repeat across areas” would mean parsing every spec on every validate. The current validator scopes itself to one change deliberately (§19.1).

If you hit those limits, the escape hatch is a *cache*, not a change of model: keep the Org files authoritative and derive an index from them, with the files’ modification ticks as the invalidation signal. What you should not do is make the model authoritative and the text a rendering — at that point you have stopped using Org and started using a database that happens to write `.org` files, and you lose the property that made Org worth choosing: **a human can open the file and edit it with ordinary commands, and your system picks that up as a matter of course.**

That property is the entire reason this project stores its specs in Org rather than in JSON beside the code. It is the same bet as Chapter 7’s task sessions and Chapter 25’s event log: the human’s editor is the interface, and the file is the API.

18.7 Applying this elsewhere

If you want a structured workflow over Org, the checklist:

1. Define `cl-defstructs` in your domain's vocabulary, not Org's.
2. Write one parse module using `org-element`. Remember the section gotcha (§16.2) — a headline's paragraphs and drawers hide under a section child.
3. Capture `:begin/:end` source text for anything you will write back.
4. Put every marker — tag names, property names, regexes, keyword mappings — in one table, as `orgspec.el` does (§15.2), so nothing downstream hardcodes a string.
5. Write transformations as pure functions over strings, in temp buffers.
6. Mutate with `org-paste-subtree`, `org-cut-subtree`, `org-todo`, `org-entry-*` — never with regexps over headline syntax.
7. Build everything, validate everything, then write everything.
8. Let TODO keywords be your state machine. They are already the user's vocabulary, `org-agenda` already queries them, and you get a dashboard for free (§19.2).

Roughly 300 lines of `orgspec-parse.el` plus `orgspec-fold.el` implement all eight. It is a small pattern with a good ratio.

19 orgspec — Validator, Lifecycle, Review, Tools

The remaining orgspec files, briefly.

19.1 The hard-gate validator

orgspec-validate.el runs the ERROR rules over a parsed change and returns *all* problems, not just the first:

```
(cl-flet ((err (fmt &rest args) (push (apply #'format fmt args) errors)))
  (unless reqs (err "Change must have at least one delta"))
  (dolist (r (append added modified))
    (unless (string-match-p orgspec-normative-regexp (or (orgspec-requirement-body r) ""))
      (err "%s \"%s\" must contain SHALL or MUST" op name))
    (unless (orgspec-requirement-scenarios r)
      (err "%s \"%s\" must include at least one scenario" op name)))
  ...duplicates within an op...
  ...duplicate FROM / TO across RENAMED...
  ...cross-op conflicts...)
(nreverse errors)
```

Elisp Feature: cl-flet and cl-labels

cl-flet defines local functions, lexically scoped to its body:

```
(cl-flet ((err (fmt &rest args) (push (apply #'format fmt args) errors)))
  (err "problem with %s" name))
```

Bindings are *not* recursive and cannot see each other — use cl-labels for that. Both come from cl-lib.

Contrast with a let-bound lambda, which would require funcall:

```
(let ((err (lambda (msg) (push msg errors))))
  (funcall err "..."))
```

cl-flet lets you call it as an ordinary function. This is because Elisp is a *Lisp-2*: a symbol has a separate function cell and value cell, and (foo x) looks up the function cell while foo reads the value cell. cl-flet binds the function cell locally; let binds the value cell.

The two-tier interface:

```
(defun orgspec-validate-change (change) ...) ; returns a list
(defun orgspec-validate-change-or-signal (change) ; signals or returns t
  (let ((errors (orgspec-validate-change change)))
    (when errors
      (user-error "orgspec validation failed:\n- %s" (string-join errors "\n- ")))
    t))
```

The list version backs the `orgspec_validate` MCP tool (report, do not fail); the signalling version backs archive (gate). Same rules, two behaviours.

19.2 The lifecycle and the agenda payoff

```
(defun orgspec-lifecycle--keyword-for-role (role)
  (pcase role
    ('active orgspec-todo-active)
    ('blocked orgspec-todo-blocked)
    ('removed orgspec-todo-removed)
    ('done 'done)
    (_ (error "Unknown orgspec lifecycle role: %S" role))))
```

'done maps to the symbol `done` rather than a keyword string, so `(org-todo 'done)` uses Org's own done handling — logging, CLOSED: timestamps, repeaters — instead of setting a keyword directly.

And then the payoff, which is the entire argument for doing this in Emacs:

```
(defun orgspec-agenda-install (changes-dir)
  (let* ((tag-match (mapconcat #'car orgspec-op-tags "|"))
        (entry `(,orgspec-agenda-key "orgspec in-flight requirements"
          tags-todo ,(format "+{%s}" tag-match)
          ((org-agenda-files ',(orgspec-agenda-files changes-dir))
           (org-agenda-overriding-header
            "orgspec: in-flight delta requirements")))))
    (setq org-agenda-custom-commands
      (cons entry (assoc-delete-all orgspec-agenda-key org-agenda-custom-commands)))
    entry))
```

One `org-agenda-custom-commands` entry, and you have an in-flight-requirements dashboard across every active change — filtered by op tag, showing TODO state, navigable, refileable, clockable. OpenSpec had to build that by hand. Here it is 20 lines because the data was already in the host's native format.

The `(cons entry (assoc-delete-all KEY ...))` pattern makes registration idempotent: remove any prior entry under the same key, then add.

19.3 The review

Because `orgspec-commands-fold-build` returns `(FILE . NEW-TEXT)` pairs without writing, the review is nearly free:

```
(defun orgspec-review-fold (id)
  "Ediff the fold of change ID against the current specs, before writing."
  (interactive "sChange id to review: ")
  (let ((built (orgspec-commands-fold-build id)))
    (dolist (cell built)
      (let* ((file (car cell))
             (new-text (cdr cell))
             (on-disk (when (file-exists-p file)
                          (with-temp-buffer (insert-file-contents file))
```

```

                                (buffer-string))))))
  (orgspec-review--ediff (file-name-nondirectory file) on-disk new-text)))
  (mapcar #'car built)))

```

“See the fold, not trust the fold.” It reuses the window-configuration save/restore convention from Chapter 6, minus the accept/reject bindings — this review writes nothing, so there is nothing to accept.

19.4 Typed MCP tools

orgspec-mcp.el wraps the verbs as MCP tools and registers them:

```

(defun orgspec-mcp-register ()
  "Register the orgspec tools onto `mcp-emacs-server-extra-tools'.
Idempotent: replaces any previously registered orgspec descriptors."
  (let ((names (mapcar (lambda (tl) (plist-get tl :name)) orgspec-mcp--tools)))
    (setq mcp-emacs-server-extra-tools
          (append orgspec-mcp--tools
                  (seq-remove (lambda (tl) (member (plist-get tl :name) names))
                             mcp-emacs-server-extra-tools))))))

```

```
(orgspec-mcp-register)
```

The bare call at the bottom of the file means requiring it registers the tools. Idempotence is achieved by removing same-named descriptors first, so reloading the file during development does not accumulate duplicates.

A schema helper removes repetition, since six of the eight tools take just an id:

```

(defun orgspec-mcp--id-schema (desc)
  (mcp-emacs-server--obj
   "type" "object"
   "properties" (mcp-emacs-server--obj "id" (mcp-emacs-server--prop "string" desc))
   "required" (vector "id")))

```

Handlers are thin adapters that format a value into text:

```

(defun orgspec-mcp--validate (args)
  (let ((errors (orgspec-validate-change
                 (orgspec-commands--read-change (alist-get 'id args)))))
    (if errors
        (format "invalid:\n- %s" (mapconcat #'identity errors "\n- "))
        "valid")))

```

20 The Test Suite

23 test files, roughly 6,600 lines, run on every push and pull request — 22 suites and around 800 assertions. This chapter covers how they work, why they are shaped the way they are, and how to write a new one — because if you are going to modify this code, this is the chapter that lets you do it with confidence.

20.1 Still not ERT, but no longer bare

The suites do not use ERT, Emacs’s built-in test framework, despite the `-test.el` names. A suite is a *batch script*: loading the file runs it, and an assertion is a check call that prints one line.

Why not ERT? ERT gives you selective test running, fixtures, failure diffing, and an interactive UI. What it costs is a layer of indirection between “load this file” and “these assertions ran”. The check approach keeps one real advantage: a suite is *just Elisp you can load*. You can `M-x load-file` it in a running Emacs and watch the output appear, with all your normal debugging tools available.

What has changed is that the two things that were genuinely missing — a shared vocabulary and a runner — now exist.

20.1.1 `test/test-helper.el`: one check, and a description

There used to be three drifted copies of check, one per naming scheme, one of which also printed got/want. They are now one file, and the assertion says what was being tested rather than naming a symbol:

```
(describe "orgspec-agenda-install"
  (it "keeps a single entry when installed twice"
    (check (length org-agenda-custom-commands) 1)))
```

prints

```
DESCRIBE orgspec-agenda-install
PASS keeps a single entry when installed twice
```

The reasoning is in the commentary: `install-idempotent` tells you which call failed; “*keeps a single entry when installed twice*” tells you what the code was supposed to do, which is the part you need when the failure is a surprise.

`test-helper-report` shows got/want on failure but suppresses it on a green run — *on failure the values are the whole point, but on a green run they are noise that buries the description*.

20.1.2 `bin/test-emacs.sh`: an Emacs of its own

434 lines, and the single entry point for both the local loop and CI.

```

bin/test-emacs.sh                # run every suite in batch
bin/test-emacs.sh test/foo-test.el # run just these
bin/test-emacs.sh --quiet [SUITE...] # report only, logs on failure (CI)
bin/test-emacs.sh --compile        # byte-compile elisp/
bin/test-emacs.sh --compile --strict # ... warnings as errors
bin/test-emacs.sh --daemon          # start the test daemon, leave it up
bin/test-emacs.sh --gui             # a windowed test Emacs
bin/test-emacs.sh --eval FORM       # eval in the test daemon
bin/test-emacs.sh --stop            # kill it
bin/test-emacs.sh --deps            # install/refresh deps only

```

The first paragraph of its header explains why it exists, and it is about this project specifically: *fixtures that pop windows, ediffs and *claudio-client* buffers stop landing in the middle of a real editing session, and a suite that wedges or kills its instance costs nothing but that instance.*

Isolation is by state, not just by process. A second Emacs sharing the working one's MCP port or emacsclient socket is not isolated at all, so the test instance gets:

- its own init directory (test/init/ — no Doom, no user config)
- its own emacsclient socket name (mcp-emacs-test)
- the MCP HTTP server on 8775 rather than 8765
- a package dir under .test-emacs/, shared across runs, so dependencies are installed once and ~/.emacs.d is never touched

That port number is the same fact Chapter 11 relies on when it explains why the permission gate's port travels in the hook environment rather than being discovered.

Which checkout gets tested is the current directory's git work tree, not the one the script lives in. The comment records what went wrong before: deriving it from BASH_SOURCE meant running the main clone's copy from inside a worktree silently compiled and tested the main clone — *the worktree contributed nothing, and the run still reported green, so the wrong checkout looked verified.* A worktree run gets its own socket and port, derived from the path so they are stable across runs.

The guard is on test/init/init.el rather than on the script, because that is what a run actually loads — so it also catches a checkout old enough to predate the harness.

20.1.3 The report

The suites do not use ert-run-tests-batch, so the runner builds the report itself: per-suite pass/fail lines, a totals line, and JUnit XML at .test-emacs/report.xml for CI to upload as an artifact. A suite that runs *no* assertions is a failure, not a pass — the XML emits an explicit no assertions ran testcase, which is what catches a suite that silently stopped loading half way through.

The glob is test/*-test.el, so a new suite is picked up automatically. That was not always true — the workflow listed files by hand, the list fell behind, and 203 assertions including the whole terminal-free runner sat outside CI without anyone noticing.

20.2 What batch Emacs does and does not have

Understanding the constraints explains most of the test design:

- **No frames, no windows you can trust.** `--batch` runs with a minimal terminal frame. Anything that manipulates window layout must be stubbed.
- **No user input.** Nothing can wait for a keypress.
- **Packages may be missing.** CI installs `web-server` and `websocket`, but the tests are written to work even without them.
- **Timers and process output do work.** `sleep-for` runs timers; the event loop functions.

So the recurring test problem is: *how do I exercise code whose whole purpose is to open an interactive UI?* The answer is stubbing, and the codebase uses three distinct techniques for it.

20.3 Technique 1: `cl-letf` function stubbing

The workhorse. From `mcp-emacs-apply-diff-test.el`:

```
(cl-letf (((symbol-function 'ediff-buffers)
  (lambda (_a _b setup-hooks &rest _)
    ;; ediff would bind `ediff-control-buffer' and run the
    ;; setup hooks; emulate just enough for the setup lambda.
    (let ((ediff-control-buffer control))
      (dolist (fn setup-hooks) (funcall fn)))
    control))
  ((symbol-function 'set-window-configuration)
  (lambda (&rest _) (setq restored t))))
(let ((ctl (mcp-emacs--ediff-review buffer-a buffer-b "old\n" result)))
  (check "review-returns-control" ctl control)
  (check "quit-hook-installed"
    (and (with-current-buffer control ediff-quit-hook) t) t)
  (with-current-buffer control (run-hooks 'ediff-quit-hook))
  (check "quit-defaults-rejected" (car result) 'rejected)
  (check "layout-restored-on-quit" restored t)))
```

Elisp Feature: `cl-letf` and generalised variables

`let` binds a *variable*. `cl-letf` binds any **place** — anything `setf` can assign to — and restores it on exit:

```
(cl-letf (((symbol-function 'foo) (lambda () 42))
  ((symbol-value 'some-var) 7)
  ((buffer-local-value 'x buf) 1))
  ...)
```

`(symbol-function 'foo)` is the place holding a symbol's function definition. Binding it temporarily replaces the function *globally*, for the dynamic extent of the body, and restores the original afterwards — even on error, because `cl-letf` uses `unwind-protect` internally.

This is the standard Elisp mocking mechanism. Note two properties:

1. It is **not** thread- or reentrancy-safe. The replacement is global while the body runs. In batch tests that is fine.
2. It restores on non-local exit, so a signalling test does not leave a stubbed function behind for the next test in the file.

Compare with `advice-add`, which composes with existing definitions and persists

until removed. Use `cl-letf` for scoped test stubs; use `advice` when you genuinely want to wrap production behaviour.

Look at what the `ediff` stub actually does. It does not simulate `ediff` — it emulates *just the contract the code under test depends on*: bind `ediff-control-buffer`, run the setup hooks, return the control buffer. That is a good stub. It encodes exactly the assumption the production code makes, which means if that assumption ever changes, the test is where you find out.

20.4 Technique 2: `advice-add :override` for capture

From `mcp-emacs-ide-test.el`:

```
(defvar mcp-ide-test--sent nil
  "List of JSON strings the stubbed client \"sent\", newest last.")

(defun mcp-ide-test--capture (_client obj)
  (push (json-encode obj) mcp-ide-test--sent))

(advice-add 'mcp-emacs-ide--send :override #'mcp-ide-test--capture)

;; ... all the tests ...

(advice-remove 'mcp-emacs-ide--send #'mcp-ide-test--capture)
```

Here the stub is installed for the *whole file* rather than per-test, because every test needs it. `:override` replaces the function entirely; the captured messages accumulate in a variable each test resets.

The readback helper parses the most recent message:

```
(defun mcp-ide-test--last ()
  (json-parse-string (car mcp-ide-test--sent) :object-type 'alist))
```

Note it round-trips through JSON — encode on capture, decode on read. That is deliberate: it proves the object the code built is actually encodable, which is a real failure mode given the empty-alist-becomes-null trap from §3.5.

20.5 Technique 3: stub the missing package

`mcp-emacs-server` hard-requires `web-server`, which may not be installed. So:

```
(unless (require 'web-server nil t)
  (defun ws-response-header (&rest _) nil)
  (defun ws-start (&rest _) nil)
  (defun ws-stop (&rest _) nil)
  (provide 'web-server))
(require 'mcp-emacs-server)
```

Define the handful of symbols the module calls, then provide the feature so the subsequent `require` is satisfied. The IDE test does the same for `websocket`.

This works because `require` checks the features list first. `provide` pushes onto it. The stubs never actually run in these tests — `dispatch` is driven directly, no socket is opened — they exist

only to make the module loadable.

20.6 Testing the pure core is easy

orgspec-fold-test.el needs none of the above, because orgspec-fold-area is pure. The whole harness is:

```
(defun fold-test--reqs (change)
  (with-temp-buffer
    (let ((org-inhibit-startup t)) (org-mode))
    (insert change)
    (orgspec-change-requirements (orgspec-parse-change "t"))))
```

and a test is a literal Org string in, assertions on a string out:

```
(let* ((delta "* Delta
** TODO New thing                               :ADDED:
:PROPERTIES:
:AREA: auth
:END:
The system SHALL do the new thing.
*** New scenario
- GIVEN a
- WHEN b
- THEN c
")
      (out (orgspec-fold-area "" (fold-test--reqs delta))))
  (check "added-requirement-L1" (and (string-match-p "^\\* New thing" out) t) t)
  (check "added-scenario-L2" (and (string-match-p "^\\*\\* New scenario" out) t) t)
  (check "added-todo-stripped" (string-match-p "\\* TODO " out) nil)
  (check "added-op-tag-dropped" (string-match-p ":ADDED:" out) nil)
  (check "added-area-stripped" (string-match-p ":AREA:" out) nil)
  (check "added-shall-kept" (and (string-match-p "SHALL" out) t) t))
```

Six assertions covering re-levelling, TODO stripping, tag dropping, property removal, and content preservation. No mocks, no fixtures on disk, no cleanup.

Two idioms worth copying:

- (and (string-match-p ...) t) normalises a match position (an integer) to t, so check's equal comparison works against a literal t. A bare string-match-p returns the match offset, which would fail the comparison.
- A negative assertion compares against nil directly, since string-match-p returns nil on no match.

Testing the error path uses the custom error type from §17.4:

```
(check "modified-drop-guard-signals"
  (condition-case _
    (progn (orgspec-fold-area spec (fold-test--reqs delta)) nil)
    (orgspec-fold-error t))
  t)
```

Run the fold; if it completes, the expression is nil and the test fails; if it signals orgspec-fold-

error specifically, the expression is `t`. Note it catches the *specific* error, so an unrelated failure does not masquerade as a passing test.

20.7 Testing time and timers

The async apply-diff tests need to prove three timing properties: the call returns immediately, the timeout fires, and the timer is cancelled on early resolution. The last one is the clever bit:

```
;; Assert on `timer-list' directly. The single-delivery guard would
;; suppress a second callback even if the timer were never cancelled,
;; so counting calls cannot tell the two mechanisms apart -- only an
;; uncanceled timer shows up here.
(let ((before (length timer-list)))
  (mcp-emacs-apply-diff-async file "new\n" 30 (lambda (out) (push out calls)))
  (check "async-timer-armed" (> (length timer-list) before) t)
  (resolve)
  (check "async-resolve-before-timeout" (length calls) 1)
  (check "async-timer-cancelled" (length timer-list) before))
```

The comment names exactly why the obvious test is insufficient. Counting callbacks cannot distinguish “the timer was cancelled” from “the timer fired but the done guard suppressed it”. Only inspecting `timer-list` — the global list of pending timers — can. That is a test written by someone who thought about what the assertion actually proves.

The timeout test does use real time, kept short:

```
(mcp-emacs-apply-diff-async file "new\n" 1 (lambda (out) (push out calls)))
(check "async-timeout-not-yet" calls nil)
(sleep-for 2)
(check "async-timeout-status" (car calls) "Status: timeout")
(check "async-timeout-single" (length calls) 1)
;; A late human resolve after a timeout must not answer twice.
(resolve)
(check "async-timeout-then-resolve-single" (length calls) 1))
```

`sleep-for` runs the event loop, so timers fire during it (unlike a busy wait). One second of timeout plus two of sleep is a three-second test — the whole suite stays fast enough that nobody is tempted to skip it.

20.8 A reusable test macro

The async tests share setup, factored into a macro:

```
(defmacro mcp--with-async-review (bindings &rest body)
  "Run BODY with `ediff-buffers' stubbed for an async apply-diff test.
  BINDINGS is a list of (VAR . INIT) forms evaluated before the stub is
  installed. Inside BODY, `resolve' calls the captured ON-RESOLVE callback
  and `control' is the fake control buffer."
  (declare (indent 1))
  `(let* ((control (generate-new-buffer " *fake-async-control*"))
         (captured nil))
```

```

,@bindings)
(unwind-protect
  (cl-letf (((symbol-function 'ediff-buffers) ...)
            ((symbol-function 'set-window-configuration) (lambda (&rest _) nil))
            ((symbol-function 'mcp-emacs--ediff-review)
             (lambda (_a _b _entry _result &optional on-resolve _tab)
               (setq captured on-resolve)
               control)))
    (cl-flet ((resolve () (when captured (funcall captured))))
      ,@body))
  (when (buffer-live-p control) (kill-buffer control))))))

```

Elisp Feature: defmacro, declare, and anaphora

A macro receives its arguments *unevaluated* and returns a form, which is then evaluated in the caller's context. The backquote template builds that form; ,@bindings splices the caller's binding list into the generated let*.

(declare (indent 1)) tells Emacs's indentation engine to treat the first argument specially — indent it as a parameter list and the rest as a body. Purely cosmetic, but it makes macro call sites read like built-in forms.

This macro is deliberately *anaphoric*: it introduces the bindings control and resolve that the body can refer to without them appearing in the macro call. That is normally considered poor hygiene — the reader cannot see where resolve comes from — which is why the docstring names both explicitly. In test code, where the macro is defined a few lines above its uses, the trade for brevity is reasonable.

Note also the deliberate exception: one of the async tests does *not* use the macro, because it needs access to the result cell to simulate an accept. The comment says so:

The stub must expose the RESULT cell so the test can record an accept the way the real accept key binding does; otherwise only the reject branch is ever reachable and the applied path goes untested.

A test that silently exercises only half the branches is worse than no test. Noticing that, and writing the more verbose version, is the right call.

20.9 Testing the dispatch layer without a socket

mcp-emacs-server-async-test.el drives the JSON-RPC layer directly:

```

(let* ((sent nil)
      (tool (list :name "fake_async"
                  :async-handler (lambda (_args done) (setq mcp-srv-test--done done))))
      (mcp-emacs-server-extra-tools (list tool)))
  (let ((started (mcp-emacs-server--tools-call-async
                  (list (cons 'name "fake_async") (cons 'arguments nil))
                  42
                  (lambda (response) (push response sent))))))
    (check "async-call-started" started t)
    (check "async-no-immediate-response" sent nil)
    (funcall mcp-srv-test--done "Status: applied\nnew\n")
    (check "async-response-sent" (length sent) 1)

```

```
(check "async-response-id" (mcp-srv--id (car sent)) 42)
(check "async-response-text" (mcp-srv--text (car sent)) "Status: applied\nnew\n"))
```

Two things make this work:

1. **The extra-tools registry is a plain variable**, so a test can let-bind it to a synthetic tool list. No registration API to fight, no cleanup — the binding unwinds automatically. This is a real benefit of the registry design from §3.6.
2. **The dispatch function takes a send callback**, so the test supplies its own instead of a socket.

The registry tests are worth noting for what they assert:

```
(let ((tool (mcp-emacs-server--find-tool "apply_diff")))
  (check "apply-diff-has-async" (functionp (plist-get tool :async-handler)) t)
  (check "apply-diff-keeps-sync" (functionp (plist-get tool :handler)) t))

(let ((tool (mcp-emacs-server--find-tool "project_info")))
  (check "sync-tool-has-no-async" (plist-get tool :async-handler) nil))
```

These are *invariant* tests, not behaviour tests. They encode the design rule “human-answered tools are async, everything else is not” so that a future change violating it fails loudly. If you add an async tool, add a matching assertion here.

20.10 What is and is not covered

Covered well:

- The pure orgspec core: parse, fold, validate — extensively, with literal fixtures.
- The apply-diff decision logic and its async lifecycle.
- IDE protocol dispatch, lockfile round-trip, diff cleanup.
- The async dispatch contract in the server.
- The resume picker’s noise filtering.
- The backend generic interface: that required verbs fail loudly and optional ones inherit their defaults (`agent-backend-test.el`).
- The permission registry’s exactly-once resolution, the rule builders, and the additive settings write (`agent-permission-test.el`).
- Prompt composition: window splitting, history walking, the minibuffer escape hatch (`agent-prompt-test.el`).
- The session overview’s enumeration, mode-before-name matching, and per-backend state honesty (`agent-session-overview-test.el`).

Not covered:

- The live HTTP path. No test starts web-server and issues a real request. `AGENTS.md` documents manual `curl` verification instead.
- Real ediff interaction. Every ediff test stubs it.
- The eat-based runner’s terminal I/O.
- Window placement in a real frame.
- The shell scripts in `bin/`. The permission gate’s fail-closed paths are verified by hand; the Emacs side of that decision is tested.

That is a defensible line: everything below the UI boundary is tested; everything at it is stubbed and verified by hand. If you extend the project, the same line is a reasonable one to hold —

and the check-based harness means adding a test is genuinely cheap, so there is little excuse for leaving new pure logic untested.

20.11 Running the tests yourself

The whole suite, and a single suite:

```
bin/test-emacs.sh
bin/test-emacs.sh test/orgspec-fold-test.el
```

Use the script rather than a hand-rolled emacs `--batch` loop: it owns dependency installation, the isolated init directory, the glob, and the pass/fail contract, so there is exactly one of each. A bare emacs `--batch -l` still works on a suite that needs no dependencies, which is what keeps a suite debuggable in isolation — but it will use your own init unless you point it elsewhere.

For an interactive poke at the code under test, run a *test* Emacs rather than your working one:

```
bin/test-emacs.sh --gui           # a windowed instance, isolated
bin/test-emacs.sh --daemon        # ... or a daemon
bin/test-emacs.sh --eval '(claude-client-open)'
bin/test-emacs.sh --stop
```

This is the loop worth internalising for this project specifically: the code under test pops windows, spawns subprocesses, and binds ports. Doing that in the Emacs you are working in is how a wedged fixture costs you your session.

Loading a suite in your own Emacs still works and is still the fastest loop for pure logic:

```
M-x load-file RET test/orgspec-fold-test.el RET
```

Because the “framework” is `princ`, output goes to standard output in batch and to `*Messages*` interactively.

To debug a failing assertion, the ordinary tools apply: `M-x debug-on-entry` on the function under test, `edebug-defun` (`C-u C-M-x`) on it, or just `M-x ielm` and call it by hand.

21 Modifying the Code

This chapter is procedural: given a change you want to make, here is where to put it and what to be careful of.

21.1 Setting up a development loop

Emacs Lisp's great advantage is that you never need to restart. The loop is:

1. Edit the `.el` file.
2. C-M-x (`eval-defun`) on the function you changed — it is redefined immediately in the live image.
3. Exercise it.

For the MCP server specifically, a tool descriptor lives inside a `defconst`, so re-evaluating one function is not enough. Reload the file and bounce the server:

```
(load-file "elisp/mcp-emacs-server.el")
(mcp-emacs-server-stop)
(mcp-emacs-server-start)
```

or from a shell:

```
emacsclient --eval '(progn (load-file "/path/to/elisp/mcp-emacs-server.el")
                           (mcp-emacs-server-stop)
                           (mcp-emacs-server-start))'
```

Then exercise it directly with `curl`, which is the fastest way to see exactly what the wire format is:

```
curl -s -X POST http://localhost:8765/mcp \
  -H 'Content-Type: application/json' \
  -d '{"jsonrpc":"2.0","id":1,"method":"tools/list"}' | jq .

curl -s -X POST http://localhost:8765/mcp \
  -H 'Content-Type: application/json' \
  -d '{"jsonrpc":"2.0","id":2,"method":"tools/call",
     "params":{"name":"project_info","arguments":{}}}' | jq -r '.result.content[0].text'
```

Note the `defconst` caveat: re-evaluating a `defconst` *does* change its value (unlike `defvar`, which leaves an already-bound variable alone). So `load-file` genuinely refreshes the tool list. If you find a variable stubbornly keeping its old value after a reload, check whether it is a `defvar` — in which case use `setq` or C-M-x with a prefix argument.

21.2 Adding a tool

The rule from `AGENTS.md`: put the work in a helper, keep the descriptor thin.

Step 1 — write the helper in `mcp-emacs.el`. It should return a string, handle its own empty and error cases, and touch live state through `mcp-emacs--current-buffer` if it is buffer-relative:

```
(defun mcp-emacs-count-words ()
  "Return the word count of the current buffer, or a status message."
  (with-current-buffer (mcp-emacs--current-buffer)
    (let ((n (count-words (point-min) (point-max))))
      (if (zerop n)
          "Buffer is empty"
          (format "%d words" n))))))
```

Step 2 — append a descriptor to `mcp-emacs-server--tools`:

```
(list :name "count_words"
      :description "Count the words in the current Emacs buffer"
      :schema (mcp-emacs-server--no-args)
      :handler (lambda (_args) (mcp-emacs-count-words)))
```

With arguments:

```
(list :name "count_words_in"
      :description "Count the words in a named buffer"
      :schema (mcp-emacs-server--obj
                "type" "object"
                "properties" (mcp-emacs-server--obj
                              "name" (mcp-emacs-server--prop "string" "Buffer name")
                              "trim" (mcp-emacs-server--prop "boolean" "Ignore whitespace-only lines"))
                "required" (vector "name"))
      :handler (lambda (args)
                  (mcp-emacs-count-words-in
                   (alist-get 'name args)
                   (eq (alist-get 'trim args) t)))))
```

Checklist for the descriptor:

- `:name` is snake_case, matching the existing set.
- `:description` is one sentence and describes *when to use it*, since that text is what an agent reads to decide.
- Required properties go in a vector, never a list.
- Booleans are extracted with `(eq ... t)`, not truthiness (§5.2).
- No-argument tools use `mcp-emacs-server--no-args`, which produces `{}` not `null` (§3.5).

Step 3 — add the row to the README tool table. The table is the user-facing catalogue; a tool missing from it is effectively undiscoverable.

Step 4 — verify with `curl` before wiring an agent to it.

21.3 Adding a tool that needs a human answer

If your tool must wait for the user — a confirmation, a selection, a review — it **must** be async or it will not work over HTTP. Re-read §3.4.

Write two entry points:

```
(defun mcp-emacs-my-thing (args) ...) ; synchronous, polls
```

```
(defun mcp-emacs-my-thing-async (args on-done)    ; returns nil immediately,
  ...                                           ; calls ON-DONE later
  nil)
```

and register both:

```
(list :name "my_thing"
      :description "... "
      :schema ...
      :async-handler (lambda (args done) (mcp-emacs-my-thing-async args done))
      :handler (lambda (args) (mcp-emacs-my-thing args)))
```

Then reproduce the three guarantees the existing async code provides:

1. **Single delivery.** Use a (list nil) guard cell; the first of {resolution, timeout} to arrive wins and cancels the other.
2. **A timeout.** Never leave an HTTP connection open indefinitely. Cap it, default it, and arm the timer only after setup succeeded.
3. **Cleanup on every path.** Kill temp buffers, cancel timers, restore window configuration.

Copy the structure of `mcp-emacs-apply-diff-async` (§6.8) rather than inventing a new shape; it has already been debugged.

Add an invariant assertion in `mcp-emacs-server-async-test.el` alongside the existing `apply-diff-has-async` check.

21.4 Adding an orgspec op or rule

The marker table (`orgspec.el`) is the single place that names ops. Adding a fifth op means:

1. Add it to `orgspec-op-tags`.
2. Decide where it belongs in `orgspec-fold-order` — the order encodes dependencies between ops (§17.2), so think about which ops must run before yours.
3. Write `orgspec-fold--apply-<op>`, following the existing contract: verify the precondition, signal `'orgspec-fold-error` with a specific message on failure, then mutate the current buffer.
4. Add a `pcase` clause in `orgspec-fold--apply-one`.
5. Add validator rules in `orgspec-validate.el` if the op has structural requirements.
6. Add a fold test with a literal Org fixture.

Adding a *validator rule* is simpler — one `dolist` and one `err` call inside `orgspec-validate-change`, plus a test. Remember that the validator reports *all* problems; do not early-return.

21.5 Adding a subscriber to the event log

If you want something to happen whenever *any* backend sees an event — logging, a modeline indicator, a notification — do not use advice. Subscribe:

```
(defun my-agent-watcher (buffer event)
  (when (eq (plist-get event :kind) 'tool-use)
    (message "agent called %s" (plist-get event :name))))

(add-hook 'agent-backend-event-functions #'my-agent-watcher)
```


The hook is `agent-backend-event-functions`, on the shared layer — not a per-client hook. A subscriber written against it sees `claude-client` and `opencode` alike.

Two rules, both learned the hard way in this codebase:

- **Swallow your own errors.** `condition-case` or `(ignore-errors ...)` around your body. This runs inside the publishing session's own call stack, so a subscriber that signals takes down the session that published the event.
- **Ignore unknown kinds silently.** Guard with a `pcase` and a `(_ nil)` fallback, or filter with `memq` as the overview does, so a new event kind degrades to silence rather than a malformed entry. The hook's docstring makes this a *requirement* on subscribers, which is what lets a publisher add a kind without knowing who is listening.

`mcp-emacs-remote--record-runner-event` and `agent-session-overview--on-event` are the two reference implementations.

21.6 Adding a backend

To add a third agent, subclass `agent-backend` (Chapter 8) and implement the five required verbs: `connect`, `quit`, `send`, `interrupt`, `add-note`. Leave every optional verb alone unless your transport genuinely supports it — the defaults are chosen so that a minimal backend needs no stubs, and a resume that signals a user-error is more useful than one that silently does nothing.

Then:

- Derive your major mode from `agent-backend-mode`, and set `agent-backend--instance-buffer-locally` in the mode body. That single `setq-local` is what makes the buffer discoverable — there is no registry to add yourself to.
- Publish the shared event vocabulary through `agent-backend--publish`. The overview and the Org transcript then work with no changes on their side.
- Add an entry to `agent-session-overview--backends` with a naming regexp and a state function that reports only what you can actually observe.

If you find yourself wanting to report working without a turn-end signal, read Chapter 13 first: reporting less is the established choice here.

21.7 Things that will bite you

A collected list of the traps this codebase has already hit, each with the place it is documented.

Blocking in a process filter. The diff panel renders and its keys are dead. Any human-answered tool must be `async`. (§3.4, `claude-client.el` commentary.)

eat-term-send-string routes by current buffer, not by its argument. Wrap every `send` in `with-current-buffer`. (`mcp-emacs-run.el`:335.)

Empty alist encodes as null. Use `mcp-emacs-server--obj` with no arguments to get `{}`. (§3.5.)

A list where JSON needs an array. Use `vector` or `vconcat`. Applies to required arrays and content arrays.

JSON false is truthy. Compare with `(eq x t)` or configure `:false-object nil`. (§5.2, §11.5.)

Side windows abort ediff. Delete them first, restore the window configuration afterwards.

(§6.3.)

org-element hides a headline's content under a **section**. Descend into it. (§13.2.)

MODIFIED can silently drop scenarios. The guard exists; do not remove it for convenience. (§17.5.)

MultiEdit bypasses the **Edit** gate. Both must be in `claude-client-disallowed-tools`. (§11.3.)

Evil's state maps outrank a major mode's keymap. Every single-letter binding is dead in a Doom-style config unless re-registered with `evil-define-key*` from a `with-eval-after-load 'evil` block. Worse, a *minor* mode map (`evil-snipe's s`) outranks even that, and needs the package's own opt-out list. (`claude-client.el`, `agent-session-overview.el`, `agent-permission.el`.)

Destructive commands under vim motion keys. `k` and `g` were bound to "kill the process" and "erase the log"; under evil they are `evil-previous-line` and the `gg` prefix. Moving the cursor ended sessions. Either leave those keys to evil or add a confirmation. (§11.15, §13.4.)

assq on string keys never matches. The permission registry is keyed by string ids, so it uses `assoc`. Presents as "the gate never fires". (`agent-permission.el:84`.)

json-encode renders nil as null, including for empty arrays. Writing `"deny": null` back into a settings file the CLI also reads corrupts it. Use a vector, and read with `json-array-type 'vector` so untouched arrays survive the round-trip. (§8.7.)

A kept-alive connection leaks its request entry. `ws-filter` only prunes its list on the path that deletes the process, so releasing the socket means doing both. (§5.6.)

A PreToolUse glob runs past &&. `Bash(echo *)` authorises `echo x &&` anything. Never widen a rule on the human's behalf. (§8.7.)

with-help-window under a popup framework can take the parent window with it. Split your own window and supply a `display-buffer` action function. (§13.4.)

Advice depends on private names and arities. If you must advise, expect it to break; prefer a published hook. (§14.2.)

Hand-maintained file lists rot silently. `ci.yml` listed its sources and tests explicitly until seven files had accumulated outside it. Glob instead. (§17.1, §20.4.)

21.8 Style conventions to match

From `.editorconfig` and `AGENTS.md`:

- UTF-8, LF line endings, final newline, no trailing whitespace, spaces not tabs.
- Every file starts with `;;; name.el --- Summary -*- lexical-binding: t; -*-` and ends with `(provide 'name)` plus `;;; name.el ends here`.
- Internal helpers use a double dash: `mcp-emacs-server--obj`. Public functions use a single: `mcp-emacs-get-selection`.
- **Every defun has a docstring**. The byte-compiler warns otherwise, and CI byte-compiles.
- A docstring's first line is a complete sentence. Parameters are referred to in CAPS. Other functions and variables are quoted ``like-this'` so `describe-function` turns them into links.
- Long-form design rationale goes in the `;;; Commentary:` block, not in a separate document. That is why this reference could be written from the source.

Commit messages follow tbagery's conventions: imperative subject around 50 characters, no trailing period, blank line, body wrapped at 72 explaining *why*.

21.9 Before you push

```
# Byte-compile – CI does this, and warnings are the point.  
bin/test-emacs.sh --compile
```

```
# Run every suite.  
bin/test-emacs.sh
```

Both go through `bin/test-emacs.sh`, which is the same entry point CI uses; Chapter 20 explains what it arranges and why running the suites in your working Emacs is a bad idea. `--compile` `--strict` promotes warnings to errors if you want the stricter gate locally.

Byte-compilation catches the things that bite in Emacs: unused variables, undefined functions (hence the `declare-function` declarations), wrong argument counts, missing docstrings, and lexical-binding mistakes. A clean compile is most of a code review.

22 Packaging, CI, and the Plugin Surface

22.1 Distribution as a Claude Code plugin

`.claude-plugin/plugin.json`:

```
{
  "name": "mcp-emacs",
  "description": "Live Emacs integration for Claude Code: buffers, Org, xref, diagnostics, and a c
  "version": "1.0.0",
  "license": "GPL-3.0-or-later",
  "author": { "name": "Gunnar Bastkowski" },
  "homepage": "https://github.com/gbastkowski/mcp-emacs"
}
```

and `.claude-plugin/marketplace.json` declares the repository as a single-plugin marketplace pointing at `./`.

One practical note learned from operating this: the plugin version string is what the update check compares. Pushing new commits without bumping it means the client reports “already at the latest version” and never re-fetches. Bump the version when the plugin content changes.

22.2 MCP client configuration

Both are three lines. For Claude Code (`.mcp.json`):

```
{ "mcpServers": { "emacs": { "type": "http", "url": "http://localhost:8765/mcp" } } }
```

For opencode (`opencode.json`):

```
{ "mcp": { "emacs": { "type": "remote", "url": "http://localhost:8765/mcp", "enabled": true } } }
```

That is the entire client-side integration. No wrapper process, no stdio bridge — the consequence of the server being an HTTP endpoint that is already running.

22.3 Skills

`skills/*/SKILL.md` files are natural-language instructions that tell an agent *when* to use which tools. Five ship with the plugin:

Skill	Routes to
<code>edit-at-point</code>	Use the current buffer/selection when the user says “here” or “this” rather than guessing a path.

Skill	Routes to
review-diagnostics	Pull live Flycheck/Flymake/LSP diagnostics and fix the underlying code.
diagnose-emacs	The setup is broken (no checker, dead LSP, bad exec-path) — fix the environment, not the code.
org-task-loop	Drive the cooperative Org session protocol from Chapter 7.
report-issue	File a bug about the tooling itself via report_tooling_issue.

The review-diagnostics / diagnose-emacs pair is a deliberate split, and each skill’s description explicitly points at the other. Misrouting between them is the difference between “spend an hour hunting a bug that does not exist” and “your exec-path is missing ~/.local/bin”.

commands/emacs-loop.md and .claude/commands/orgspec/*.md are slash commands — thin prompt templates that invoke a skill or a tool with arguments.

22.4 CI

.github/workflows/ci.yml runs on push to main and on every pull request:

1. Check out.
2. Install Emacs 29.4 via purcell/setup-emacs.
3. bin/test-emacs.sh --deps
4. bin/test-emacs.sh --compile
5. bin/test-emacs.sh --quiet
6. Upload .test-emacs/report.xml, with if: always().

The notable thing about that list is how little of it is YAML. Every step after the Emacs install is the same script you run locally, and the workflow says why:

CI and the local loop drifting apart is how a suite ends up green in one and broken in the other; the script owns dependency install, the isolated init directory, the glob and the pass/fail contract, so there is only one of each.

--quiet prints the per-suite report only, because that is the useful artefact in a log you read only when it is red; a failing suite still gets its full output replayed. The artifact upload is if: always() for the same reason — the report is most worth having on the run that failed.

Both the compile and test steps used to list their files by hand. The lists fell behind: three source files and four test suites — 203 assertions, including the entire terminal-free runner — were never compiled or run here. The stated reason for the omission (that those files need packages CI does not install) was not even true; they had simply been forgotten. Globbing removes the failure mode entirely.

One failure condition is worth calling out. The check originally only grepped for FAIL, so a suite that *errored before printing anything* counted as a pass — the most dangerous kind of green. Requiring at least one assertion closes that, and the JUnit output now names such a suite explicitly.

22.5 Cutting a release

`bin/release.sh <version> [--dry-run]` exists because the version is declared in three places that had already drifted apart once:

- the `Version:` header of every `elisp/*.el` file
- the plugin manifest, `.claude-plugin/plugin.json`
- the git tag

Doing it by hand is what let `plugin.json` sit at 1.0.0 while the tags reached v1.7.0, and what left `agent-session-overview.el` claiming a version it was never part of. One script writes all three from one argument, so they cannot disagree.

Two guards are worth copying:

- **Bare semver only.** A leading `v` would produce the tag `vv1.8.0` and a `Version: v1.8.0` header that `package.el` cannot parse, so the script rejects it rather than normalising silently.
- **It refuses to release from a dirty tree or a side branch,** because the tag would point at something other than what was reviewed.

The tag is the part that matters to consumers: Doom's `:pin` resolves a tag, so a release is what makes a version installable by name rather than by SHA.

23 Cross-Cutting Idioms

A collection of patterns that appear repeatedly and are worth adopting generally.

23.1 Degrade, do not fail

Optional functionality never prevents the main path from working:

```
(defun mcp-emacs-run--ide-port ()
  (when mcp-emacs-run-ide-integration
    (condition-case err
      (when (require 'mcp-emacs-ide nil t)
        (or (mcp-emacs-ide-port) (mcp-emacs-ide-start)))
      (error (message "mcp-emacs-run: IDE integration unavailable: %s"
                     (error-message-string err))
             nil))))
```

Missing package, disabled feature, port exhaustion — all produce `nil`, a message, and a session that still launches.

23.2 Guard cells for single delivery

```
(let ((done (list nil)))
  (lambda (outcome)
    (unless (car done)
      (setcar done t)
      ...deliver...)))
```

Whenever two independent paths can complete the same operation — a callback and a timeout, a quit hook and a force-kill — a guard cell makes the first one authoritative. Used in `mcp-emacs-apply-diff-async` and, in a different form, in `mcp-emacs-ide--complete-open-diff` (when `id` plus immediate `remhash`).

23.3 Build in memory, write once

`orgspec` does this at the transaction level (§17.7), but the shape recurs: build the complete result in a temp buffer, validate it, and only then touch the filesystem or the user’s buffers. The payoff is that a failure halfway through leaves nothing partially applied — and the same pure function powers both “apply” and “preview”.

23.4 Derive state, do not store it

The runner has no session registry: buffer names encode identity and default-directory encodes the project (§10.2). Nothing can go stale because nothing is stored.

Prefer this whenever the derivation is cheap. The alternative — a registry with add/remove hooks — has more code and more failure modes.

23.5 Publish events, do not advise

Both approaches appear in `mcp-emacs-remote.el`, and the file's own comment prefers the newer one:

`claude-client` instead *publishes* an append-only event log and announces each event on a hook, so this is a plain subscriber — no advice, and rendering stays one consumer among several.

If you are writing the producer, publish. Advice is for when you do not control the producer.

23.6 Comment the failure, not the code

The comments that earn their space in this codebase all have the same shape: they describe a failure mode, not a mechanism.

```
;; A side or dedicated window cannot be split, which aborts ediff setup
;; before the control buffer exists. Remove side windows first and pin
;; ediff to its plain layout so the control panel lands in this frame.

;; Register the deferred id before opening ediff so a very fast resolve
;; still finds it.

;; Assert on `timer-list' directly. The single-delivery guard would
;; suppress a second callback even if the timer were never cancelled,
;; so counting calls cannot tell the two mechanisms apart.
```

None of these restate what the code does. Each records something that was learned by getting it wrong. That is the standard to hold when you add code here: if you fixed a subtle bug, the fix without the explanation is half the work.

24 The Claude Code CLI Surface

Everything in Chapters 8–11 talks to Claude Code. This chapter collects what the project actually depends on: which invocations, which wire formats, which guarantees, and where the boundaries are.

One framing point first. **This project never calls the Anthropic API.** There is no HTTP request to `api.anthropic.com` anywhere in the codebase, no API key handling, no model-request construction, no token accounting. Every path goes through the `claude` CLI binary, which owns authentication, billing, the agent loop, context management, and model selection. What `mcp-emacs` contributes is the *editor side*: tools the CLI can call, a diff gate it must pass through, and a place to render what it is doing.

That is a deliberate boundary, and it is worth naming because it constrains the design. It means the project cannot, for example, inject a system prompt mid-turn, inspect token usage, or intercept a tool call before the model sees the result — those live inside the CLI. It also means the project inherits the CLI's release cadence, which is why two modules carry a verified-against version number in their commentary.




Figure 24.1: Three CLI invocation surfaces, and the two editor-tool channels they reach

24.1 Three ways the project invokes the CLI

1. As a full-screen TUI (`mcp-emacs-run.el`, Chapter 10):

(`apply` #'eat-make name `mcp-emacs-run-executable` `nil` switches)

Bare `claude` with optional `--continue`, `--resume`, or `--resume <session-id>`, running inside an `eat` terminal buffer. The human reads the CLI's rendered output; Emacs sends keystrokes. Emacs understands nothing about the content — it is a terminal.

2. As a headless one-shot (`mcp-emacs-run--query-headless`):

```
:command (list mcp-emacs-run-executable "-p" prompt "--output-format" "text")
```

`-p` (short for `--print`) means “answer this and exit”. `--output-format text` means plain prose on stdout. Used by `mcp-emacs-explain-selection-in-current-session` when no runner window is visible: fire a question, collect stdout in a sentinel, render it in a markdown popup. No session, no state.

3. As a structured stream (`claude-client.el`, Chapter 11):

```
(list claude-client-executable
      "--print"
      "--output-format" "stream-json"
      "--input-format" "stream-json"
      "--verbose"
      "--mcp-config" (claude-client--mcp-config-file)
      "--strict-mcp-config"
      "--settings" (claude-client--settings-file)
      "--append-system-prompt-file" (claude-client--system-prompt-file))
```

plus `--model` and `--disallowedTools` when configured. This is the one where Emacs actually understands the conversation.

24.2 The flags, and what each one buys

Flag	Effect	Why this project uses it
<code>--print</code> / <code>-p</code>	Non-interactive: answer and exit	No TUI to host
<code>--output-format text</code>	Plain prose on stdout	The one-shot query path
<code>--output-format stream-json</code>	NDJSON events on stdout	The stream runner needs structure
<code>--input-format stream-json</code>	Accept NDJSON on stdin	Multi-turn input without a terminal
<code>--verbose</code>	Emit the full event set	Otherwise intermediate events are suppressed
<code>--mcp-config FILE</code>	Point at MCP servers	Hand the CLI the in-Emacs server
<code>--strict-mcp-config</code>	Use <i>only</i> that config	Ignore the user's global MCP servers
<code>--settings FILE</code>	Permission allowlist	Pre-approve <code>apply_diff</code>

Flag	Effect	Why this project uses it
<code>--append-system-prompt-file FILE</code>	Extra system prompt	Tell the model how to reach the edit path
<code>--disallowedTools A B C</code>	Disable built-in tools	Force edits through the gate
<code>--model NAME</code>	Model override	Optional
<code>--continue</code>	Resume the most recent conversation	TUI runner
<code>--resume [ID]</code>	Resume a specific conversation	TUI runner and the native picker

`--strict-mcp-config` is the quiet one. Without it the CLI would merge the user's global MCP servers into the generated config, and the carefully constructed allowlist would not describe the actual tool set.

24.3 The generated configuration files

Three temp files per stream run, all written with `with-temp-file`.

MCP config — points the CLI at the running in-Emacs server:

```
{ "mcpServers": { "emacs": { "type": "http", "url": "http://localhost:8765/mcp" } } }
```

Settings — the permission allowlist:

```
{ "permissions": { "allow": ["mcp__emacs__apply_diff", "mcp__emacs__open_file", ...],
  "deny": [] } }
```

Note the tool-name convention: an MCP tool is addressed as `mcp__<server-name>__<tool-name>`. The server name comes from the key in the MCP config (`emacs`), so `apply_diff` becomes `mcp__emacs__apply_diff`. Rename the server in the config and every entry in the allowlist must change with it.

System prompt — appended to the CLI's own:

You are running inside Emacs. You have no Write or Edit tools.
To change a file you must call `mcp__emacs__apply_diff` with the absolute path and the complete new file content. The human reviews it in `ediff` and accepts or rejects; the tool result tells you which. Never claim a change was made unless `apply_diff` returned `applied`.

The docstring explains why this is needed: “Without this the model tends to reach for its (disabled) built-in tools first and then report that it cannot edit.” Disabling a tool tells the model *what it cannot do*; the prompt tells it *what to do instead*.

24.4 The stream-json wire format

Output is NDJSON — one complete JSON object per line, framed on newlines (§10.5). Four message types matter.

system / init — the session opens:

```
{"type":"system","subtype":"init","session_id":"abc-123","model":"claude-...", ...}
```

assistant — a model turn, carrying content blocks:

```
{"type":"assistant","message":{"role":"assistant","content":[
  {"type":"text","text":"I'll update the parser."},
  {"type":"tool_use","name":"mcp__emacs__apply_diff","input":{"path":"...", "new_content":"..."}}
]}}
```

user — synthetic, carrying tool results back:

```
{"type":"user","message":{"content":[
  {"type":"tool_result","content":"Status: applied\n..."}
]}}
```

The content of a `tool_result` is either a string or an array of blocks, which is why the handler normalises both:

```
:text (if (stringp content)
          content
          (mapconcat (lambda (c) (or (alist-get 'text c) "")) content "\n"))
```

result — the run ends:

```
{"type":"result","subtype":"success","permission_denials":[...], ...}
```

`permission_denials` is worth a note. It lists tool calls the permission layer refused. The commentary is explicit that this is **not a gate**: it is a post-hoc audit record, arriving only at the end of the run, long after the refusal happened. You cannot use it to approve anything. It is captured into the finished event and currently not rendered.

The mapping is deliberately lossy — five external message shapes become five internal event kinds, and anything unrecognised is dropped by the trailing `(t nil)`. That is forward compatibility: a CLI release that adds a message type does not break the runner.

24.5 The negative finding: `can_use_tool`

The most useful paragraph in `claude-client.el`:

The CLI's `can_use_tool` stdio control request is NOT used: it is never emitted in headless mode (verified against 2.1.220), and the `permission_denials` in the result are a post-hoc audit record rather than a gate.

The obvious design for a permission gate would be: the CLI asks “may I run this tool?”, Emacs prompts the human, Emacs answers. There is a stdio control protocol that looks like it does this. It does not fire in headless mode.

The project's memory records the same finding from the other direction: a `--permission-prompt-tool` flag that would have served is gone from current releases. So the permission model had to be built out of what *does* work.

24.5.1 What does work: the `PreToolUse` hook

The resolution arrived later than the finding, and it is worth stating as plainly as the negative result was. `bin/permission-gate.sh` opens with it:

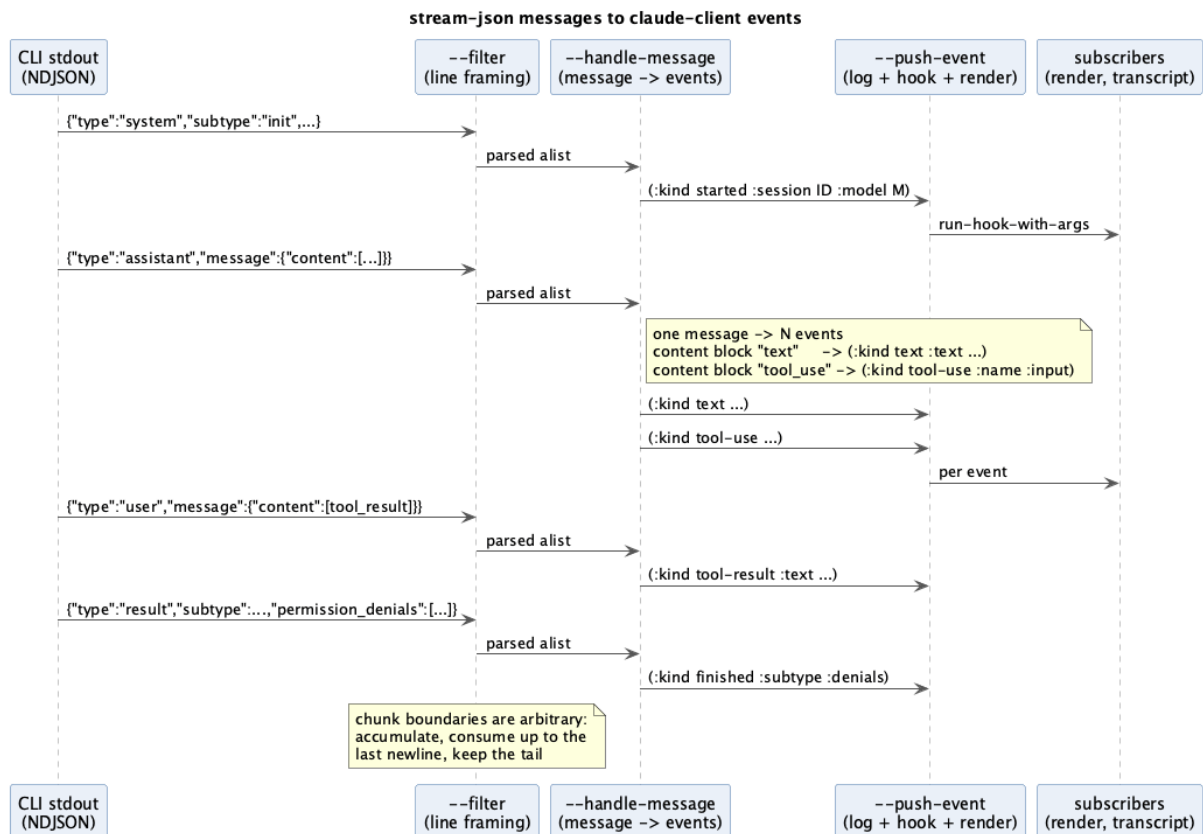


Figure 24.2: Message types collapse into five internal event kinds

The CLI spawns this and blocks on it, which is the whole point: it is the one place a permission decision can still be *made* rather than merely reported. `can_use_tool` is never emitted in headless mode and `permission_denials` only arrives at the end of the turn, long after the refusal — so neither can gate anything.

The hook receives the payload on stdin (`tool_name`, `tool_input`, `tool_use_id`, `session_id`, `cwd`, `permission_mode`), and whatever it prints on stdout decides the call. That blocking spawn is the mechanism the stdio control request failed to provide.

The script stays deliberately dumb — read, POST, print — and the decision logic lives in Emacs, *where it can be tested without a CLI in the loop*. That is why Chapter 20 can list the permission registry as well covered while the script itself is verified by hand.

Everything else about the script is the fail-closed posture:

```
# Fail closed. Every failure path here -- no Emacs, no server, malformed
# reply, curl missing -- prints a deny. A gate that opens up when it breaks
# is not a gate, and this runs unattended by definition.
```

Four paths deny: no curl, an unreachable server, an empty reply, and a reply that does not contain "permissionDecision". The last one refuses to guess: *anything else is a bug on the Emacs side, and guessing at intent here is how a deny silently becomes an allow*.

The timeout is 110 seconds with a stated reason — the CLI kills the hook on its own timeout, so this stays under it so *the human sees a reasoned deny rather than a hook that simply vanished*. Compare `mcp-emacs-server-permission-timeout` on the Emacs side: two deadlines, and the inner one has to be the shorter.

Every `permissionDecisionReason` in the script is written for the model, since it reaches the model verbatim: it says what happened and that retrying will not help.

See Chapter 5 for the `/permission-gate` route this posts to, Chapter 8 for the buffer the human answers, and Chapter 11 for how the hook is installed into a session's settings.

24.6 Three gates, and the hole between them

There are now three gates, covering two different questions. The first two ask *may this file change land?* and the third asks *may this command run at all?*

The stream runner's edit gate is **compositional**, and every part is load-bearing:

1. `--disallowedTools Write Edit MultiEdit NotebookEdit` removes the model's ability to write directly. The docstring flags the trap: "MultiEdit must be disabled alongside Edit: otherwise Claude can batch file changes through it and bypass the ediff gate entirely."
2. MCP tools are denied by default, so the settings allowlist must include `mcp__emacs__apply_diff` — otherwise, as that docstring says, "every edit is refused before the human ever sees the ediff." The gate can fail *closed* as easily as open.
3. `apply_diff` opens an ediff the human answers (Chapter 6).

The TUI runner's gate is **environmental**: the CLI uses its own native Edit/Write, and those route through `openDiff` only when it was launched with `CLAUDE_CODE_SSE_PORT` and `ENABLE_IDE_INTEGRATION` and successfully connected to the WebSocket server.

The third gate is the **PreToolUse hook**, and it is different in kind. The first two constrain *edits*; this one intercepts any tool in `claude-client-gate-tools` — by default Bash, the case where the tool name alone does not tell you whether to allow it. A shell command cannot be routed through an ediff, because there is no diff to show; the only available review is the command itself, before it runs.

Note the division of labour. The mutating built-ins are disabled outright rather than gated, so the hook never fires for them; gating them would raise a menu for a call that cannot happen. Disabling and gating are alternatives, not layers.

And the hole, which the hook narrows without closing: a `claude` the user starts by hand in a shell has neither the edit gate nor the IDE gate. `mcp-emacs-run-ide-integration`'s docstring says so plainly — "IDE integration only works for interactive sessions launched by this runner; a Claude Code started by hand will not connect." Both remain a property of *how the process was started*.

The permission hook is the one part that can outlive its launcher, because A (always allow) writes into `.claude/settings.local.json`, which the terminal CLI also reads — and that cuts both ways. It is why Chapter 8 insists on a confirmation showing the exact rule text: a rule added from a `claude-client` session changes what a hand-started session may do, in a file this project does not own.

24.7 The IDE protocol handshake

The spike finding that shapes `mcp-emacs-ide.el` is that `openDiff` is not enough. Before every native edit the CLI calls `closeAllDiffTabs`, then `getDiagnostics`, then `openDiff`, **blocking on each**. Serve only `openDiff` and the CLI hangs on the first call.

Hence `getDiagnostics` being answered with a stubbed `[]` — it must return something, and




Figure 24.3: The two independent paths a write can take

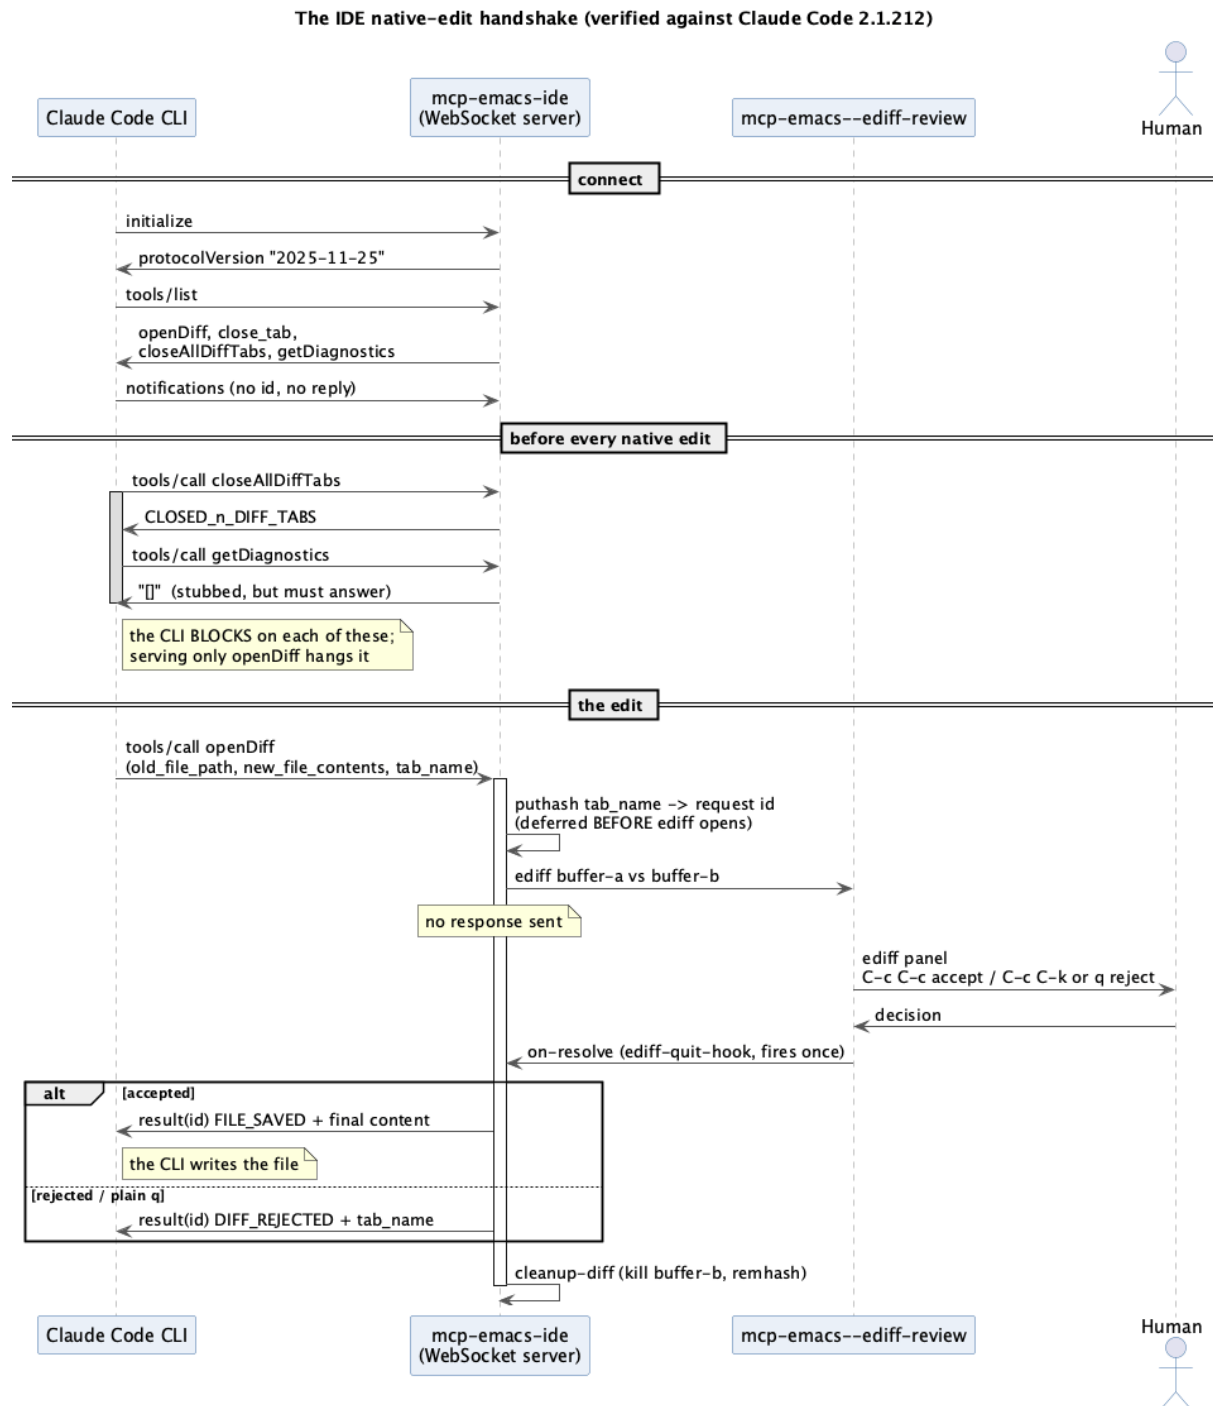


Figure 24.4: The three-call sequence before every native edit

an empty array is accepted.

`openDiff` is the only deferred call. Everything else answers immediately.

24.8 Version fragility, and how the project handles it

Three version numbers are recorded in the source:

Where	Version	What it pins
<code>mcp-emacs-ide.el</code>	Claude Code 2.1.212	The IDE protocol and its call sequence
<code>claude-client.el</code>	Claude Code 2.1.220	<code>can_use_tool</code> not firing in headless mode
<code>mcp-emacs-server.el</code>	MCP 2025-06-18	The protocol version advertised in <code>initialize</code>

The IDE constant carries an instruction for the next person:

```
(defconst mcp-emacs-ide-protocol-version "2025-11-25"
  "MCP protocol version advertised to Claude Code.
  Verified against Claude Code 2.1.212; echo whatever the client offers
  if this ever drifts again.")
```

“Echo whatever the client offers” is the fallback strategy if the version ever moves again — respond with the client’s own version rather than a hardcoded one.

The structural mitigation is that the whole IDE surface is opt-in and isolated: `mcp-emacs-ide-enabled` defaults to `nil`, the module never touches the HTTP server, and a failure to start it degrades to a session without IDE integration (§10.6). An unofficial protocol breaking should cost you a feature, not your editor.

24.9 If you wanted to talk to the API directly

The project does not, but it is worth knowing what would change if it did, since the question comes up.

Going direct would mean owning: authentication and key storage, the request/response loop, the tool-use cycle (the model asks for a tool, you run it, you send the result back, repeat), context window management, streaming response parsing, retry and rate-limit handling, and cost accounting. In exchange you would get full control over the system prompt, genuine pre-tool permission gating, and independence from CLI releases.

The project’s memory records the decision plainly: billing was the blocker. A CLI subscription and API-key billing are different products. As long as the human’s usage is on the former, going direct is not merely more work — it is a different account.

That trade is worth re-examining if the constraint changes, but nothing in the current codebase would need to move: `apply_diff`, the ediff review, the event log, and the render buffer are all transport-agnostic. Only the ~120 lines of `claude-client.el` that spawn a subprocess and parse NDJSON would be replaced.

25 The Event Model

This chapter began as a sketch of a design that did not exist. Most of it now does. It is kept as a chapter rather than folded into the others because the reasoning is what makes the pieces legible: the shape was chosen first, in docs/*VISION.md*, and the features were then chosen against it.

Where something is built, this says so and points at the code. Where a question was open and has since been answered, the answer and the evidence for it are recorded — several were settled by measuring the CLI rather than by argument, and those measurements are the load-bearing part.

25.1 The problem with turns

The vision document states it directly:

The chatbot loop is fundamentally **request/response**: a turn is a closed transaction — you speak, it locks, it answers, it unlocks. “Add a thought at any time” has no place in that model, because there is no *any time*; there is only “your turn” and “its turn.”

That is the whole motivation. Not speed, not capability — the absence of a place to put a thought that arrives while the agent is working.

The proposed alternative:

An **event stream** has no turns. There is an append-only log of events, and both the human and the AI *produce* to it and *react* from it. The log is itself a live artifact you can look at.

Note the last clause. In Emacs, “an append-only log you can look at” is a buffer. The log is not infrastructure hidden behind an API; it is a thing the human reads and writes with ordinary editing commands. That is the same “same live state, concurrently” idea from Chapter 1, made temporal.

25.2 The aggregates already exist

The DDD framing produces a useful surprise, and it is the observation that makes the design plausible rather than speculative:

The aggregates are not records we would invent — **they are the live Emacs artifacts we already have**, and the state lives in them, not in a struct beside them.

Bounded context	Aggregate (its store)	Identity	Transitions today
Session (org-task)	the live Org subtree	SESSION property	org_task_set_ session_status, set_item_status, append_note, append_item
Spec (orgspec)	Change → Requirement → Scenario	headline text	orgspec-lifecycle- advance, orgspec-archive
Runner	the eat buffer or claude-client buffer	*claude:<project>:<n>	run-new, continue, resume, kill, quit
Review	the (list nil) result cell	ediff tab name	accept / reject / timeout, resolved once in ediff-quit-hook

Every one of those is a chapter of this document. The Session aggregate is Chapter 7; the Spec aggregates are Chapters 15–19; the Runner is Chapters 10 and 11; the Review is Chapter 6. Nothing new needs to be built to have aggregates — they are already here, and each already has a well-defined transition vocabulary.

There is a second observation embedded in the table: **Org TODO keywords are already the shared state-transition vocabulary** across the Session and Spec contexts. `STRT`, `WAIT`, `KILL`, `DONE` mean the same thing to a checklist item and to a delta requirement. That is not a coincidence to be exploited later; it is why `orgspec-todo-active` and friends are `defcustoms` reading the user’s own workflow (§15.2).

The domain events these emit, or trivially could: `NoteAdded`, `SessionStatusChanged`, `ItemStatusChanged`, `RequirementAdvanced`, `ChangeProposed`, `FoldArchived`, `RunnerStarted` / `RunnerQuit`, `DiffOpened` / `DiffResolved`, and the two that come from *either* actor — `FileChanged` and `ToolInvoked`.




Figure 25.1: The domain map: contexts, aggregates, and the shared log

25.3 What exists today, and precisely what is missing

Three primitives already exist, each solving a piece of the problem in isolation.

An event feed exists. `mcp-emacs-remote` (Chapter 14) appends tool calls, diff outcomes, and session boundaries into a per-project Org transcript. It is append-only, timestamped, and struc-

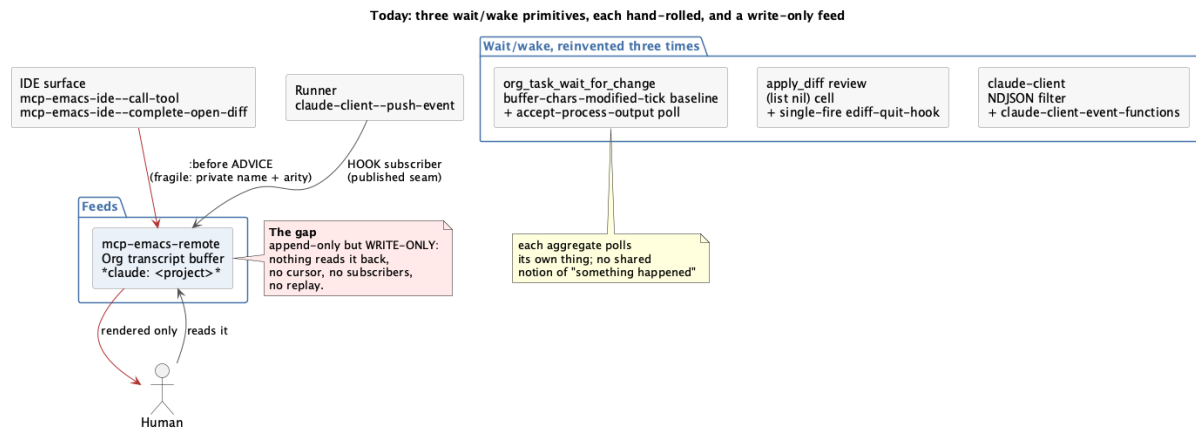


Figure 25.2: Today: three hand-rolled wait/wake primitives and a write-only feed

tured. It was also **write-only** — nothing read it back. That is the piece that has since changed: the transcript is now a *subscriber* on `agent-backend-event-functions`, and the runner publishes to a log rather than writing a record directly.

And there are now two readers, not one. `agent-session-overview` (Chapter 13) subscribes to the same hook and renders a live table. This matters more than a second feature usually would: one subscriber can always be a rendering callback wearing a hook’s clothes. Two independent subscribers, one of which filters the kinds it cares about and neither of which knows about the other, is what makes the claim “one log, many readers” true rather than aspirational.

A wait/wake primitive exists. `org_task_wait_for_change` (§7.6) blocks on a monotonic `buffer-chars-modified-tick` baseline via an `accept-process-output` poll. That is `on(change)` without freezing Emacs — precisely the semantics a reactor needs.

A one-shot resolve primitive exists. The diff review’s `(list nil)` cell plus a single-fire `ediff-quit-hook` (§6.5) is a promise: exactly one resolution, whichever path gets there first.

So the gap was never “we need an event system”. It was narrower, and it has largely closed:

- The feed was **one-directional**. It now has programmatic subscribers; the rendered transcript is one of them, and keeping the conversation window visible across an ediff review is another.
- The producers were **one**. The human is now a second: a note (`claude-client-add-note`) writes to the same log the runner does, and the four `org-task` writes emit observations into it as well, so one record shows both actors in order.
- The two producer mechanisms **disagreed**. The IDE taps still use advice (fragile, §14.2); everything added since publishes a hook (robust, §11.4). The newer one is the pattern; the older one predates it and remains.

What has *not* changed, deliberately: the log is not a store. The Org file remains the aggregate and the log only observes it, so nothing reconstructs state by replaying — which is also why the log’s not surviving a restart costs nothing that matters.

25.4 The target shape

The structural change is small and the consequence is large: the AI stops being *the* consumer and becomes *a* consumer. Once the log has subscribers, a rendered transcript, an agenda refresh, a lint trigger, and the AI reactor are all the same kind of thing. Adding a reader costs




Figure 25.3: Target: one log, many readers

the producers nothing — which is exactly the property §10.4 established with agent-backend-event-functions, generalised.

And the human becomes a first-class *producer*. `NoteAdded` at any time is not a special interrupt path; it is just another event on the same log, which any reactor may or may not care about.

25.5 The cooperative loop, as an event stream

The first band of that diagram is the flow that works today, using the tick token and the poll wait. The second band — “add a thought mid-flight” — is now built; the rest of this chapter is how each question was settled.

25.6 The three hard questions, and their answers

The vision named these as “the actual content of the design”, which was a fair assessment: the plumbing was straightforward and these were not.

1. Interruption semantics — answered: interrupt and re-plan. A `NoteAdded` arrives while the AI is mid tool-call. Finish, abort, or re-plan?

The CLI settled what was affordable. Its init event advertises `interrupt_receipt_v1`, and it accepts a control request that abandons the turn in flight. Two properties were measured before anything was built on them: the interrupt lands instantly, and *the session survives it* — the next turn answers normally. So abandoning a turn costs the in-flight work and nothing else, which makes the aggressive answer cheap enough to be the default.

A note written mid-turn therefore interrupts, and is redelivered with framing that says the previous turn was cut short on purpose and not to resume it. That framing is load-bearing: without it the model picks up the abandoned work instead of redirecting.

Note what this did *not* require: a notion of “safe point”. The earlier worry — is a half-applied edit a safe point? is an open ediff? — dissolved because the edit path is gated at the MCP tool boundary. An interrupt cannot land mid-write; it lands between tool calls, and a review already in flight is answered by the human on its own terms.

2. Ordering and consistency — mostly dissolved by a decision. Two writers on one file makes `FileChanged` ordering a merge problem. This was the same concern that caused the agent backend abstraction to be deferred.

The decision that the Org file is the aggregate removes most of it. For Session and Item state there is a single writer of record — Org — so there is nothing to merge; the log observes rather

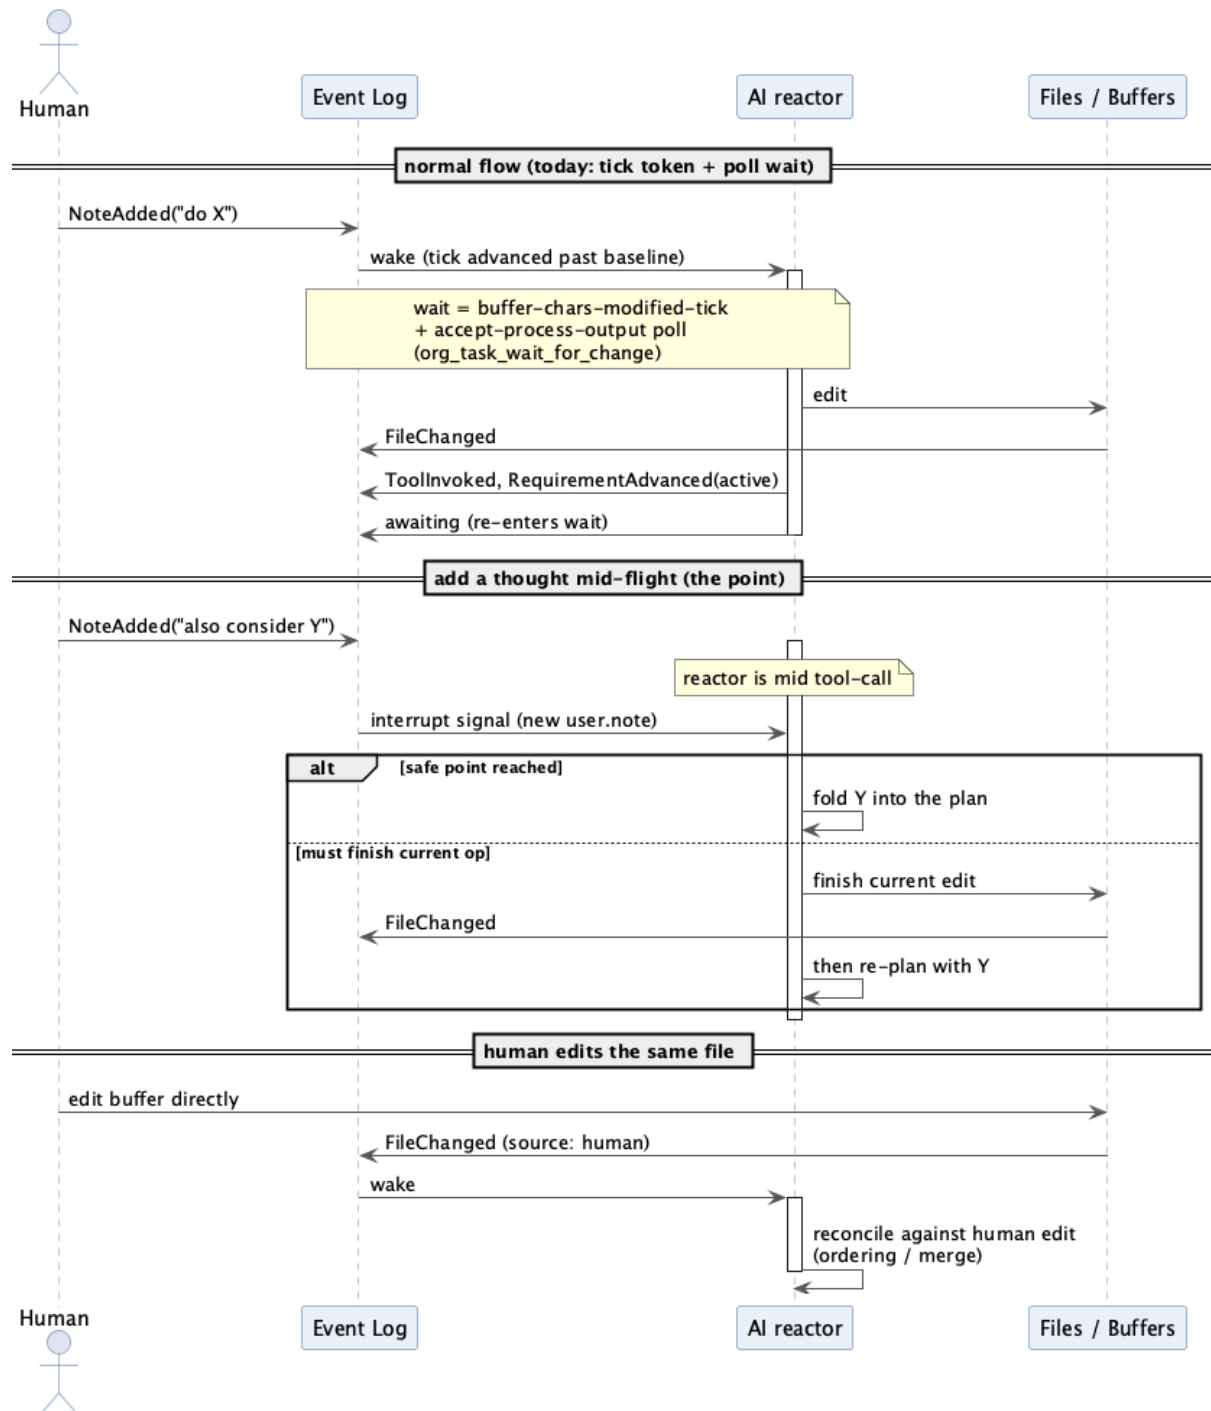


Figure 25.4: The loop, and the mid-flight note it exists for

than competes. What remains is the genuine two-writer case on file contents, and the diff review already answers it: if buffer A was hand-edited during a review, the human's edit wins and the proposal is discarded (§6.4). Last-human-writer-wins, scoped to one file, one review.

3. Backpressure — answered by probing, not by choosing. Several notes arrive in quick succession while the AI is slow. The intended design was “queue them”, and the interesting part is that writing a probe against the implementation found three defects that reading it had not:

- a second note fired *another* interrupt into a turn already dying, and logged a second interrupted event for one interruption;
- an exact repeat queued twice and was sent twice, spending context to say nothing new;
- the queue was unbounded, so a slow model meant notes piling up until they arrived as one enormous prompt with the newest thought behind everything stale.

The answers: one interrupt per turn (the interrupt flag doubles as the guard), exact repeats coalesced but still recorded, and a bounded queue that drops the *oldest* — the recent note is the current intent — with each drop logged rather than silent.

All three now have answers in the repository. Two of them came from measuring rather than deciding, which is the part worth generalising: the interruption rule was chosen because the CLI made a cheap answer available, and the backpressure rules were written because a probe found defects that reading the code had not.

25.7 The smallest first step, and what it decided

The vision was specific about not over-building:

Don't refactor everything onto a bus. Make the `mcp-emacs-remote` transcript a readable log with one subscriber, and wire the single path that most *is* the vision — “add a note at any time” — reusing the tick-token wait. If that cleanly subsumes the hand-rolled `org-task` loop, the abstraction is right. Then settle the interruption rule before going wider.

Both halves were built, and the test was run. **The answer is that it does not subsume the `org-task` loop** — and after the aggregate decision it was never going to.

That is not a failure of the abstraction; it is the abstraction finding its scope. With Org owning Session and Item state, `wait_for_change` plus direct editing is the *right* shape for an owned aggregate: the human edits the file, the AI reacts to what changed. What the log contributes is the push-based wake in place of a poll, and one shared record in place of two side-by-side ones. It does not contribute ownership, and trying to make it would have created a second source of truth for the same checklist.

Concretely: of the `org-task` loop's operations, the log subsumes the note path and improves on the wait; the item and session status operations stay in Org and merely emit observations. That split is stable, not a way-station.

25.8 What this means if you are modifying the code today

The direction has practical consequences right now, independent of whether the log is ever built.

Prefer publishing to advising. If you add a producer, publish an event on a hook (§23.5). The

IDE taps are advice because they predate the pattern, not because advice is right. Every hook-based producer is one that will not need rewriting.

Keep events as plists with a `:kind`. That is the existing shape (`claude-client--events`), it matches the `pcase` dispatch used by consumers, and it degrades well — an unknown `:kind` falls through a `(_ nil)` clause instead of erroring.

Do not build a second wait/wake primitive. There are already three. If you need “wake when X changes”, the tick-token pattern from §7.6 is the one to copy — it is the one the design expects to generalise.

Assume your consumer is one of several. Swallow your own errors (§23.1), never assume you are the only reader, and never mutate the event you were handed.

Following those four while writing ordinary features is what makes the eventual unification a refactor rather than a rewrite.